

Databases and Information Systems

Introduction Conceptual Modeling

Dmitriy Traytel

slides mostly by Marcos Vaz Salles



UNIVERSITY OF
COPENHAGEN

What is a DBMS?

- Database: A very large, integrated collection of data.
- Models real-world **enterprise**.
 - Entities (e.g., students, professors, courses)
 - Relationships (e.g., “attends”, “teaches”, “is a TA for”)
- A **Database Management System (DBMS)** is a software package designed to store and manage databases.

Database Management Systems (DBMS)

A Database Management System (DBMS) provides...

Database Management Systems (DBMS)

A Database Management System (DBMS) provides...

... **efficient, reliable, convenient, and safe multi-user**
storage of and access to **massive** amounts of **persistent** data.

Database Management Systems (DBMS)


A Database Management System (DBMS) provides...

... **efficient, reliable, convenient, and safe multi-user**
storage of and access to **massive** amounts of **persistent** data.

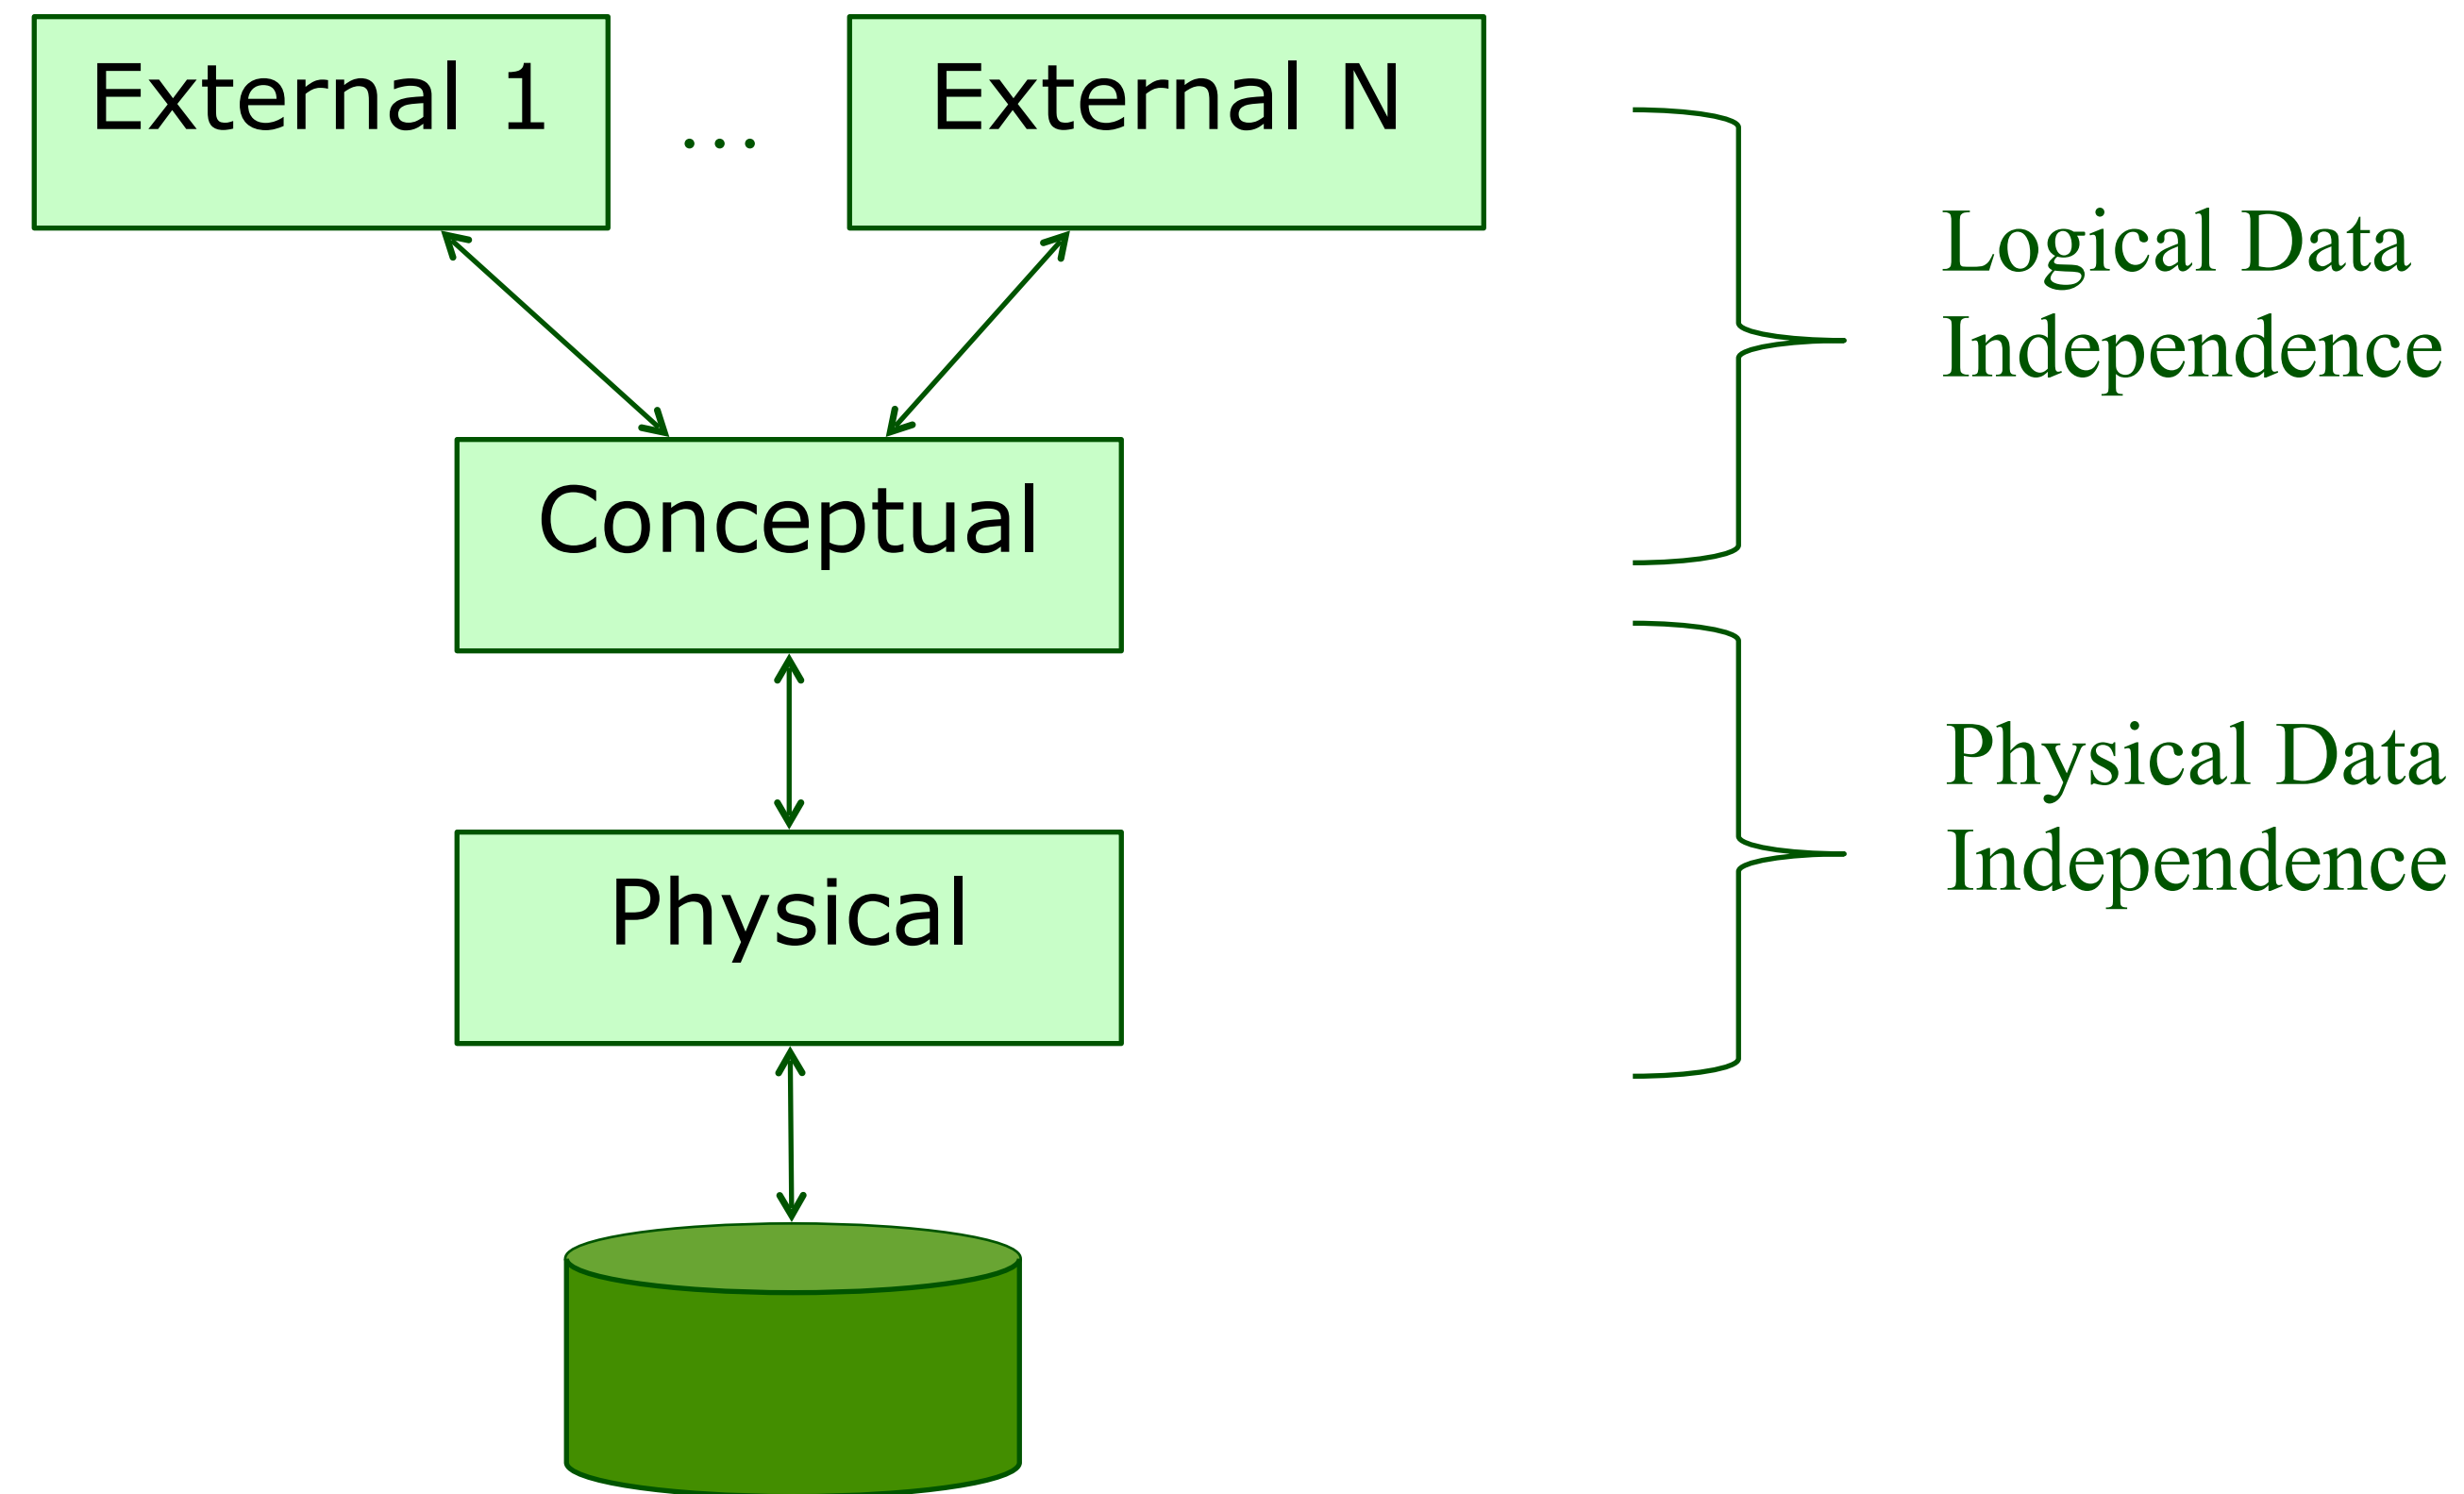
We will focus on Relational DBMS (RDBMS)

Why Use a DBMS?

- **Data independence** and efficient access.
- Reduced application development time.
- Data integrity and security.
- Uniform data administration.
- Concurrent access and recovery from crashes.



Levels of Abstraction



Databases: Key Concepts and People

- **Key Concepts**

- Data model
- Schema versus data
- Data definition language (DDL)
- Data manipulation (or query) language (DML)

- **Key People**

- DBMS implementer
- Database designer
- Database application developer
- Database administrator

Databases: Key Concepts and People

- **Key Concepts**

- Data model
- Schema versus data
- Data definition language (DDL)
- Data manipulation (or query) language (DML)

What are examples of the key concepts in the case of a real application, e.g., a social network?

- **Key People**

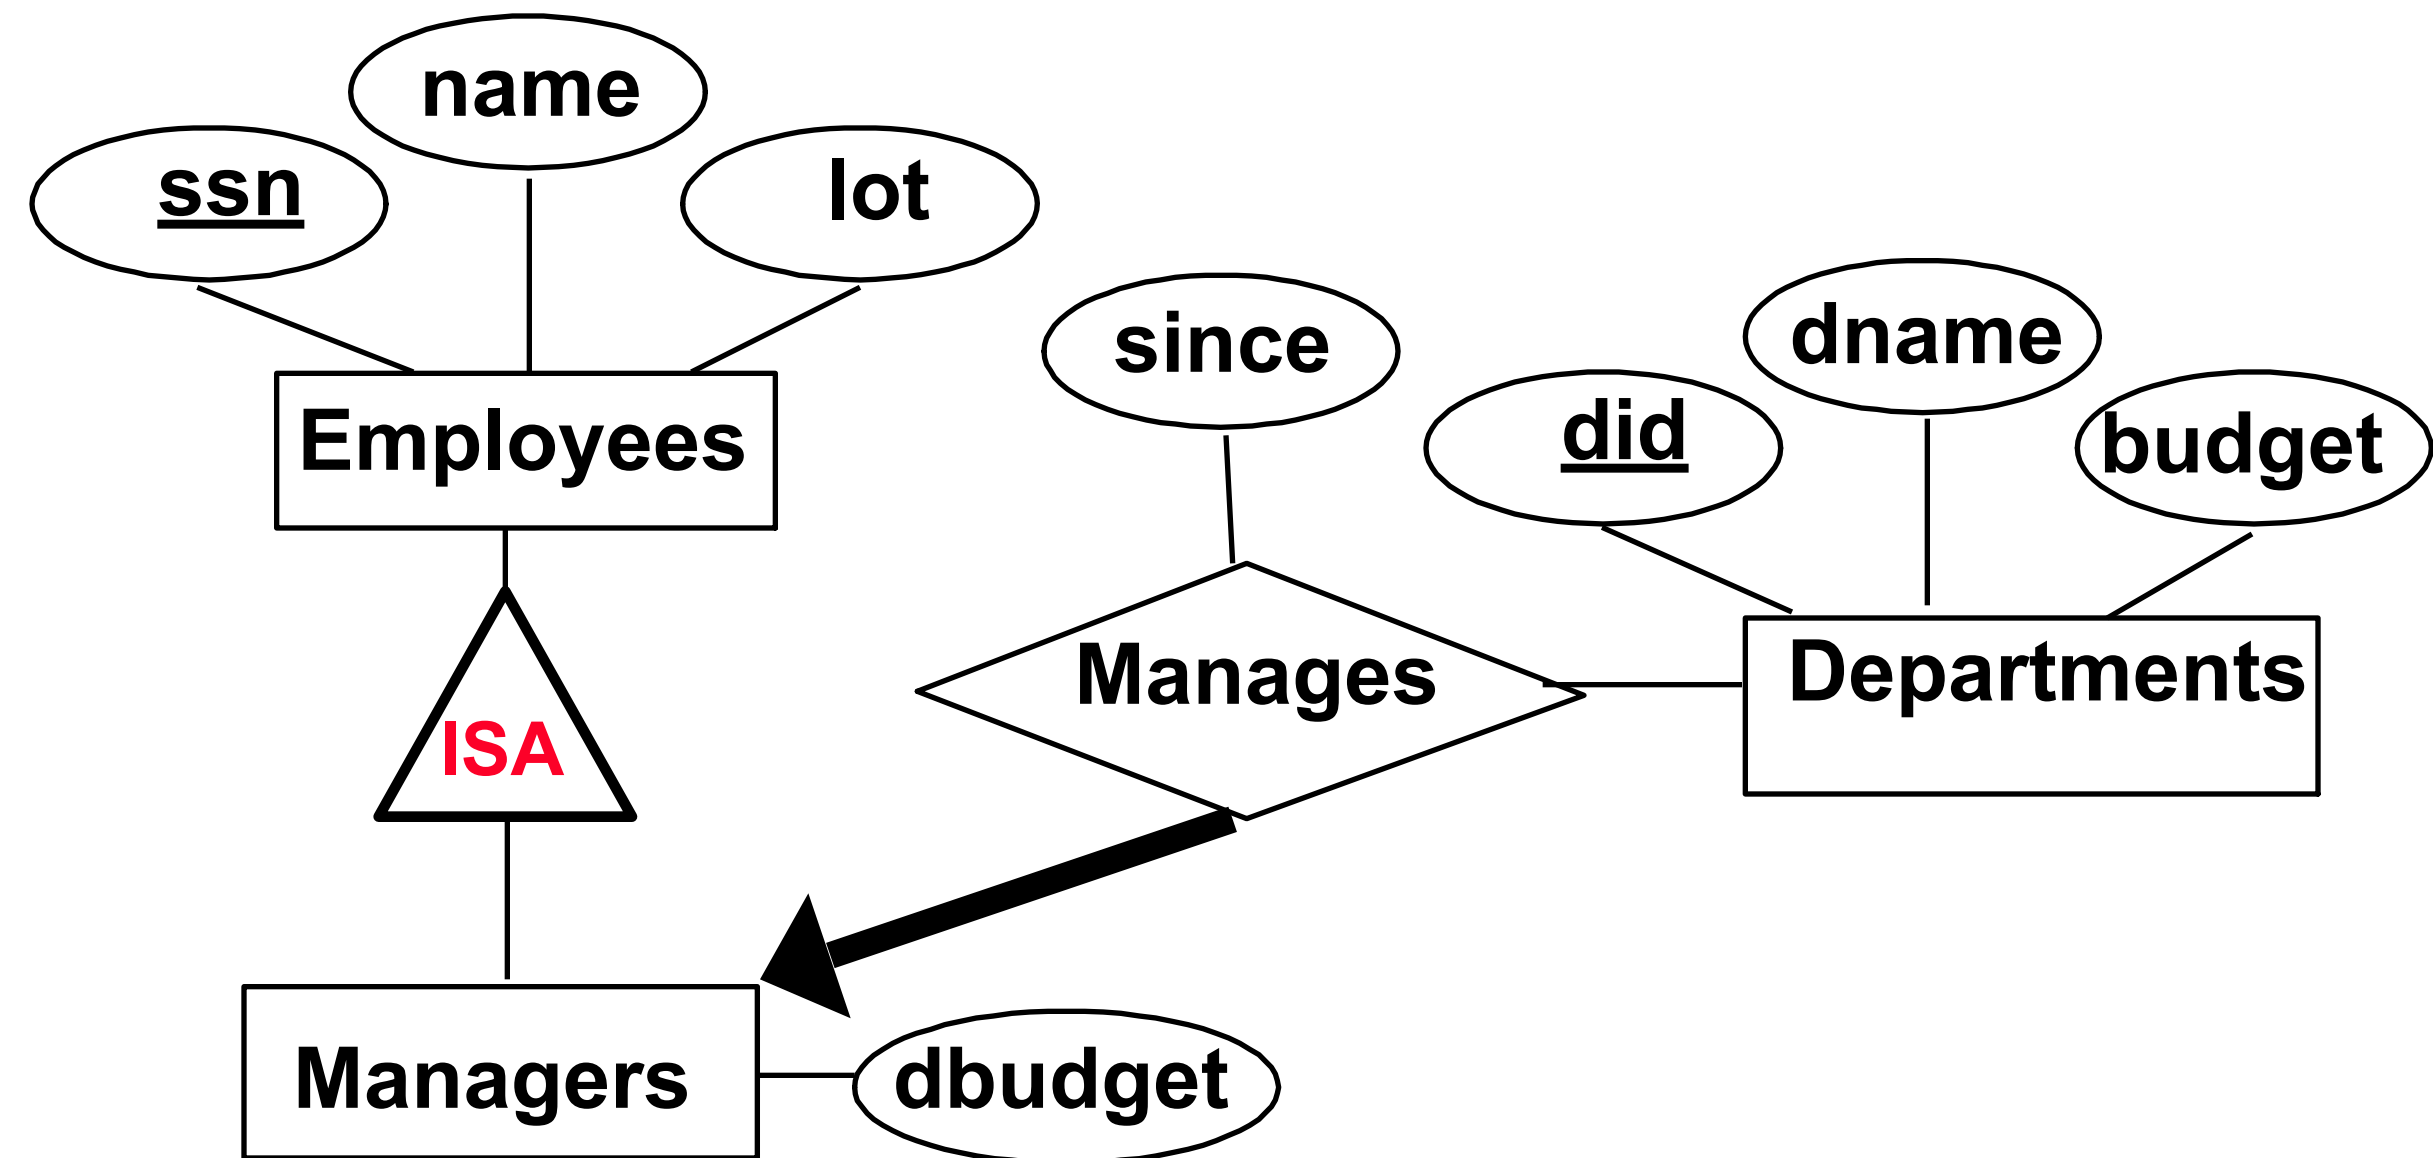
- DBMS implementer
- Database designer
- Database application developer
- Database administrator

What should we learn today?

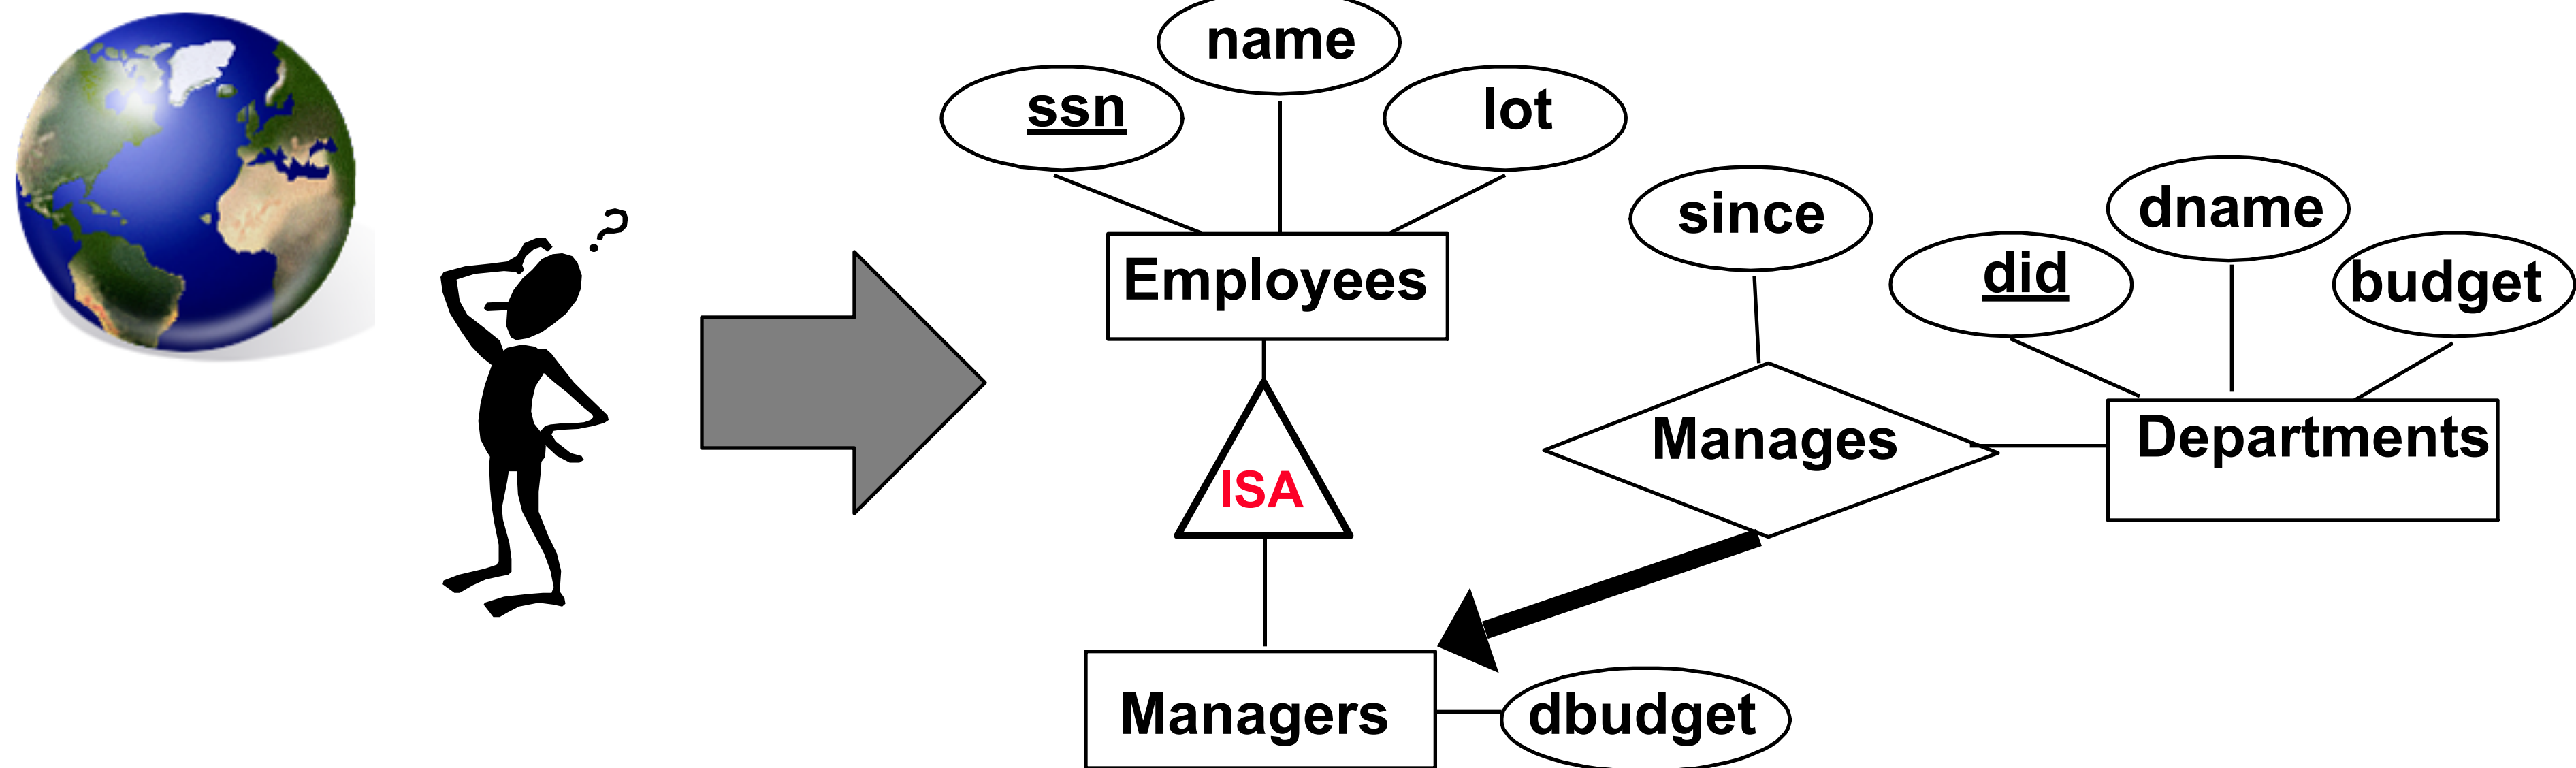


- Explain the concepts of entity (set), relationship (set) and express these concepts in entity-relationship (E/R) diagrams
- Explain and express constraints (key, uniqueness, and ref. integrity) in E/R diagrams
- Explain and capture in E/R diagrams weak entity sets and ISA hierarchies
- Analyze trade-offs and argue for the advantages and disadvantages of a E/R design
- Explain what a data model is, along with the distinction between schemas and instances
- Define and explain the relational data model
- Define and explain the main classes of integrity constraints in the relational model, in particular key and foreign key constraints
- Model relational schemas and express them in SQL

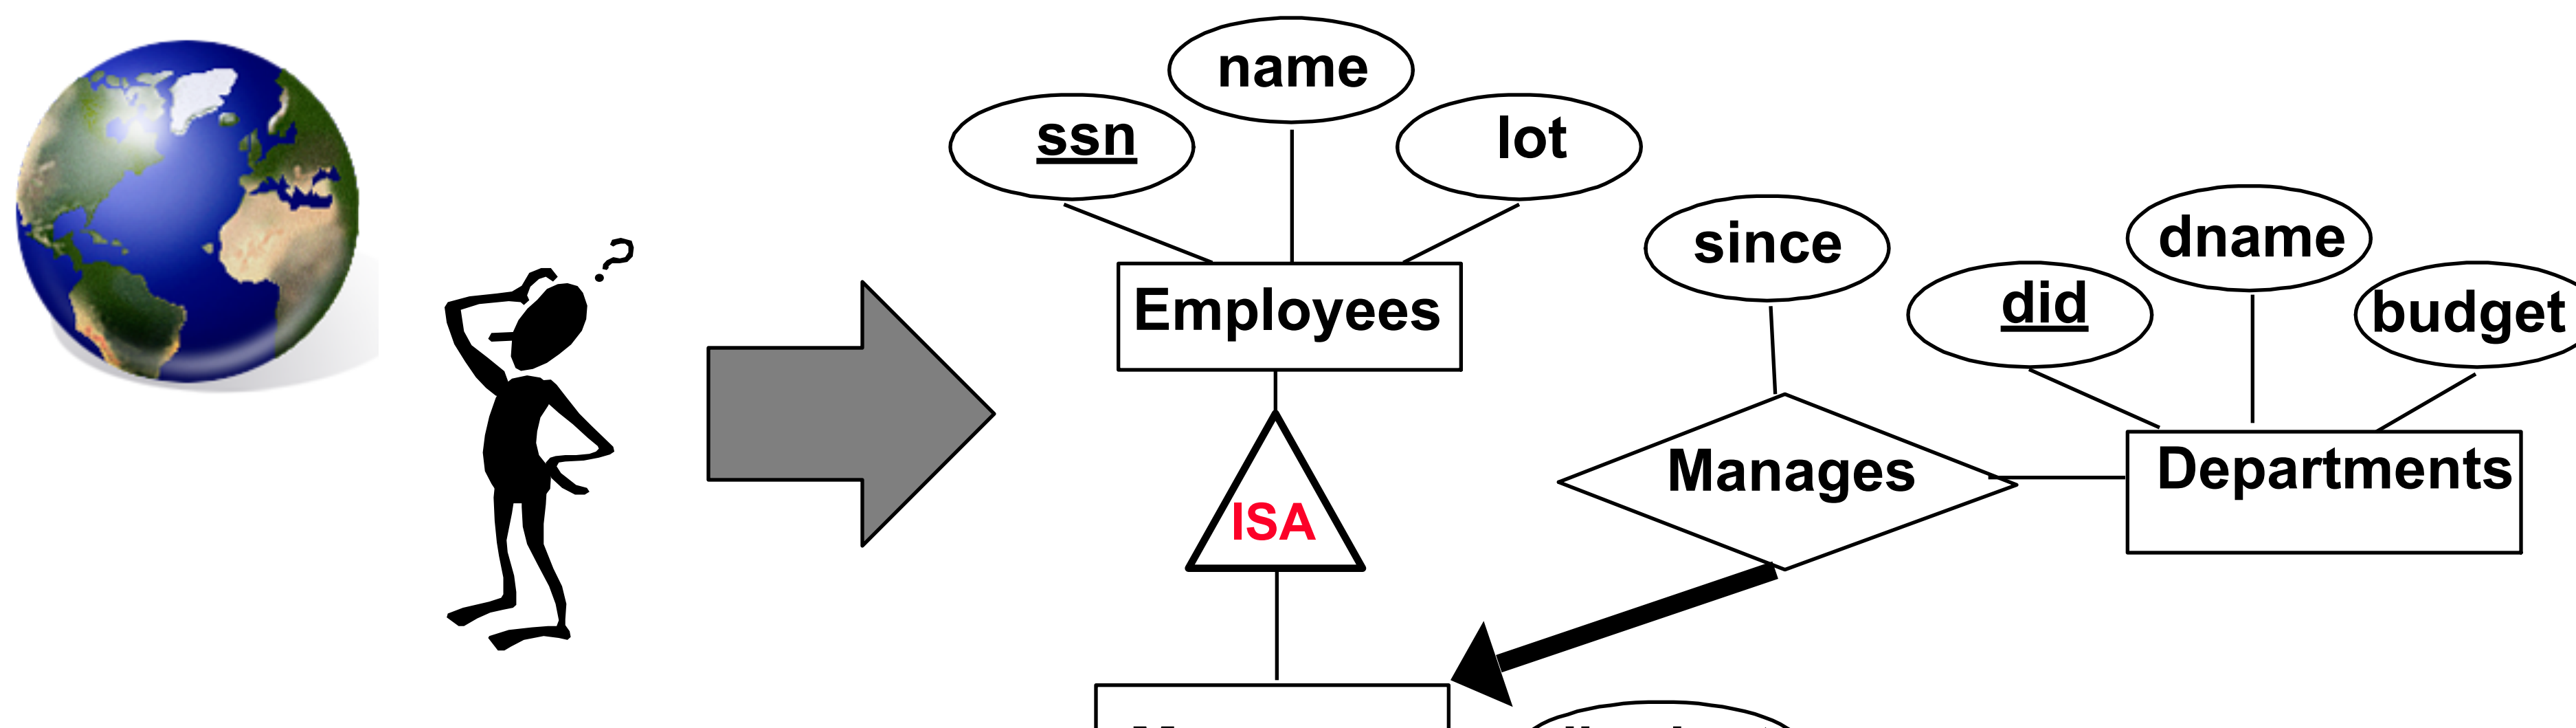
Conceptual Models



Conceptual Models




Conceptual Models



How do we create a conceptual model?



E/R Model Basics

- **Entity**: Real-world object distinguishable from other objects. An entity is described (in DB) using a set of **attributes**
- **Entity Set**: A collection of similar entities. E.g., all employees
 - All entities in an entity set have the same set of attributes
 - Each entity set has a **key**
 - Each attribute has a **domain**




E/R Model Basics

- **Relationship**: Association among two or more entities.
 - E.g., Attishoo works in Pharmacy department.
- **Relationship Set**: Collection of similar relationships.
 - An n-ary relationship set R relates n entity sets E1 ... En
 - Each relationship in R involves entities e1 in E1, ..., en in En




E/R Model Basics

- **Relationship**: Association among two or more entities.
 - E.g., Attishoo works in Pharmacy department.
- **Relationship Set**: Collection of similar relationships.
 - An n-ary relationship set R relates n entity sets E1 ... En
 - Each relationship in R involves entities e1 in E1, ..., en in En



Relationships (Contd.)


- Relationships can have roles
- For example, if we want to capture supervisor-subordinate relationship:



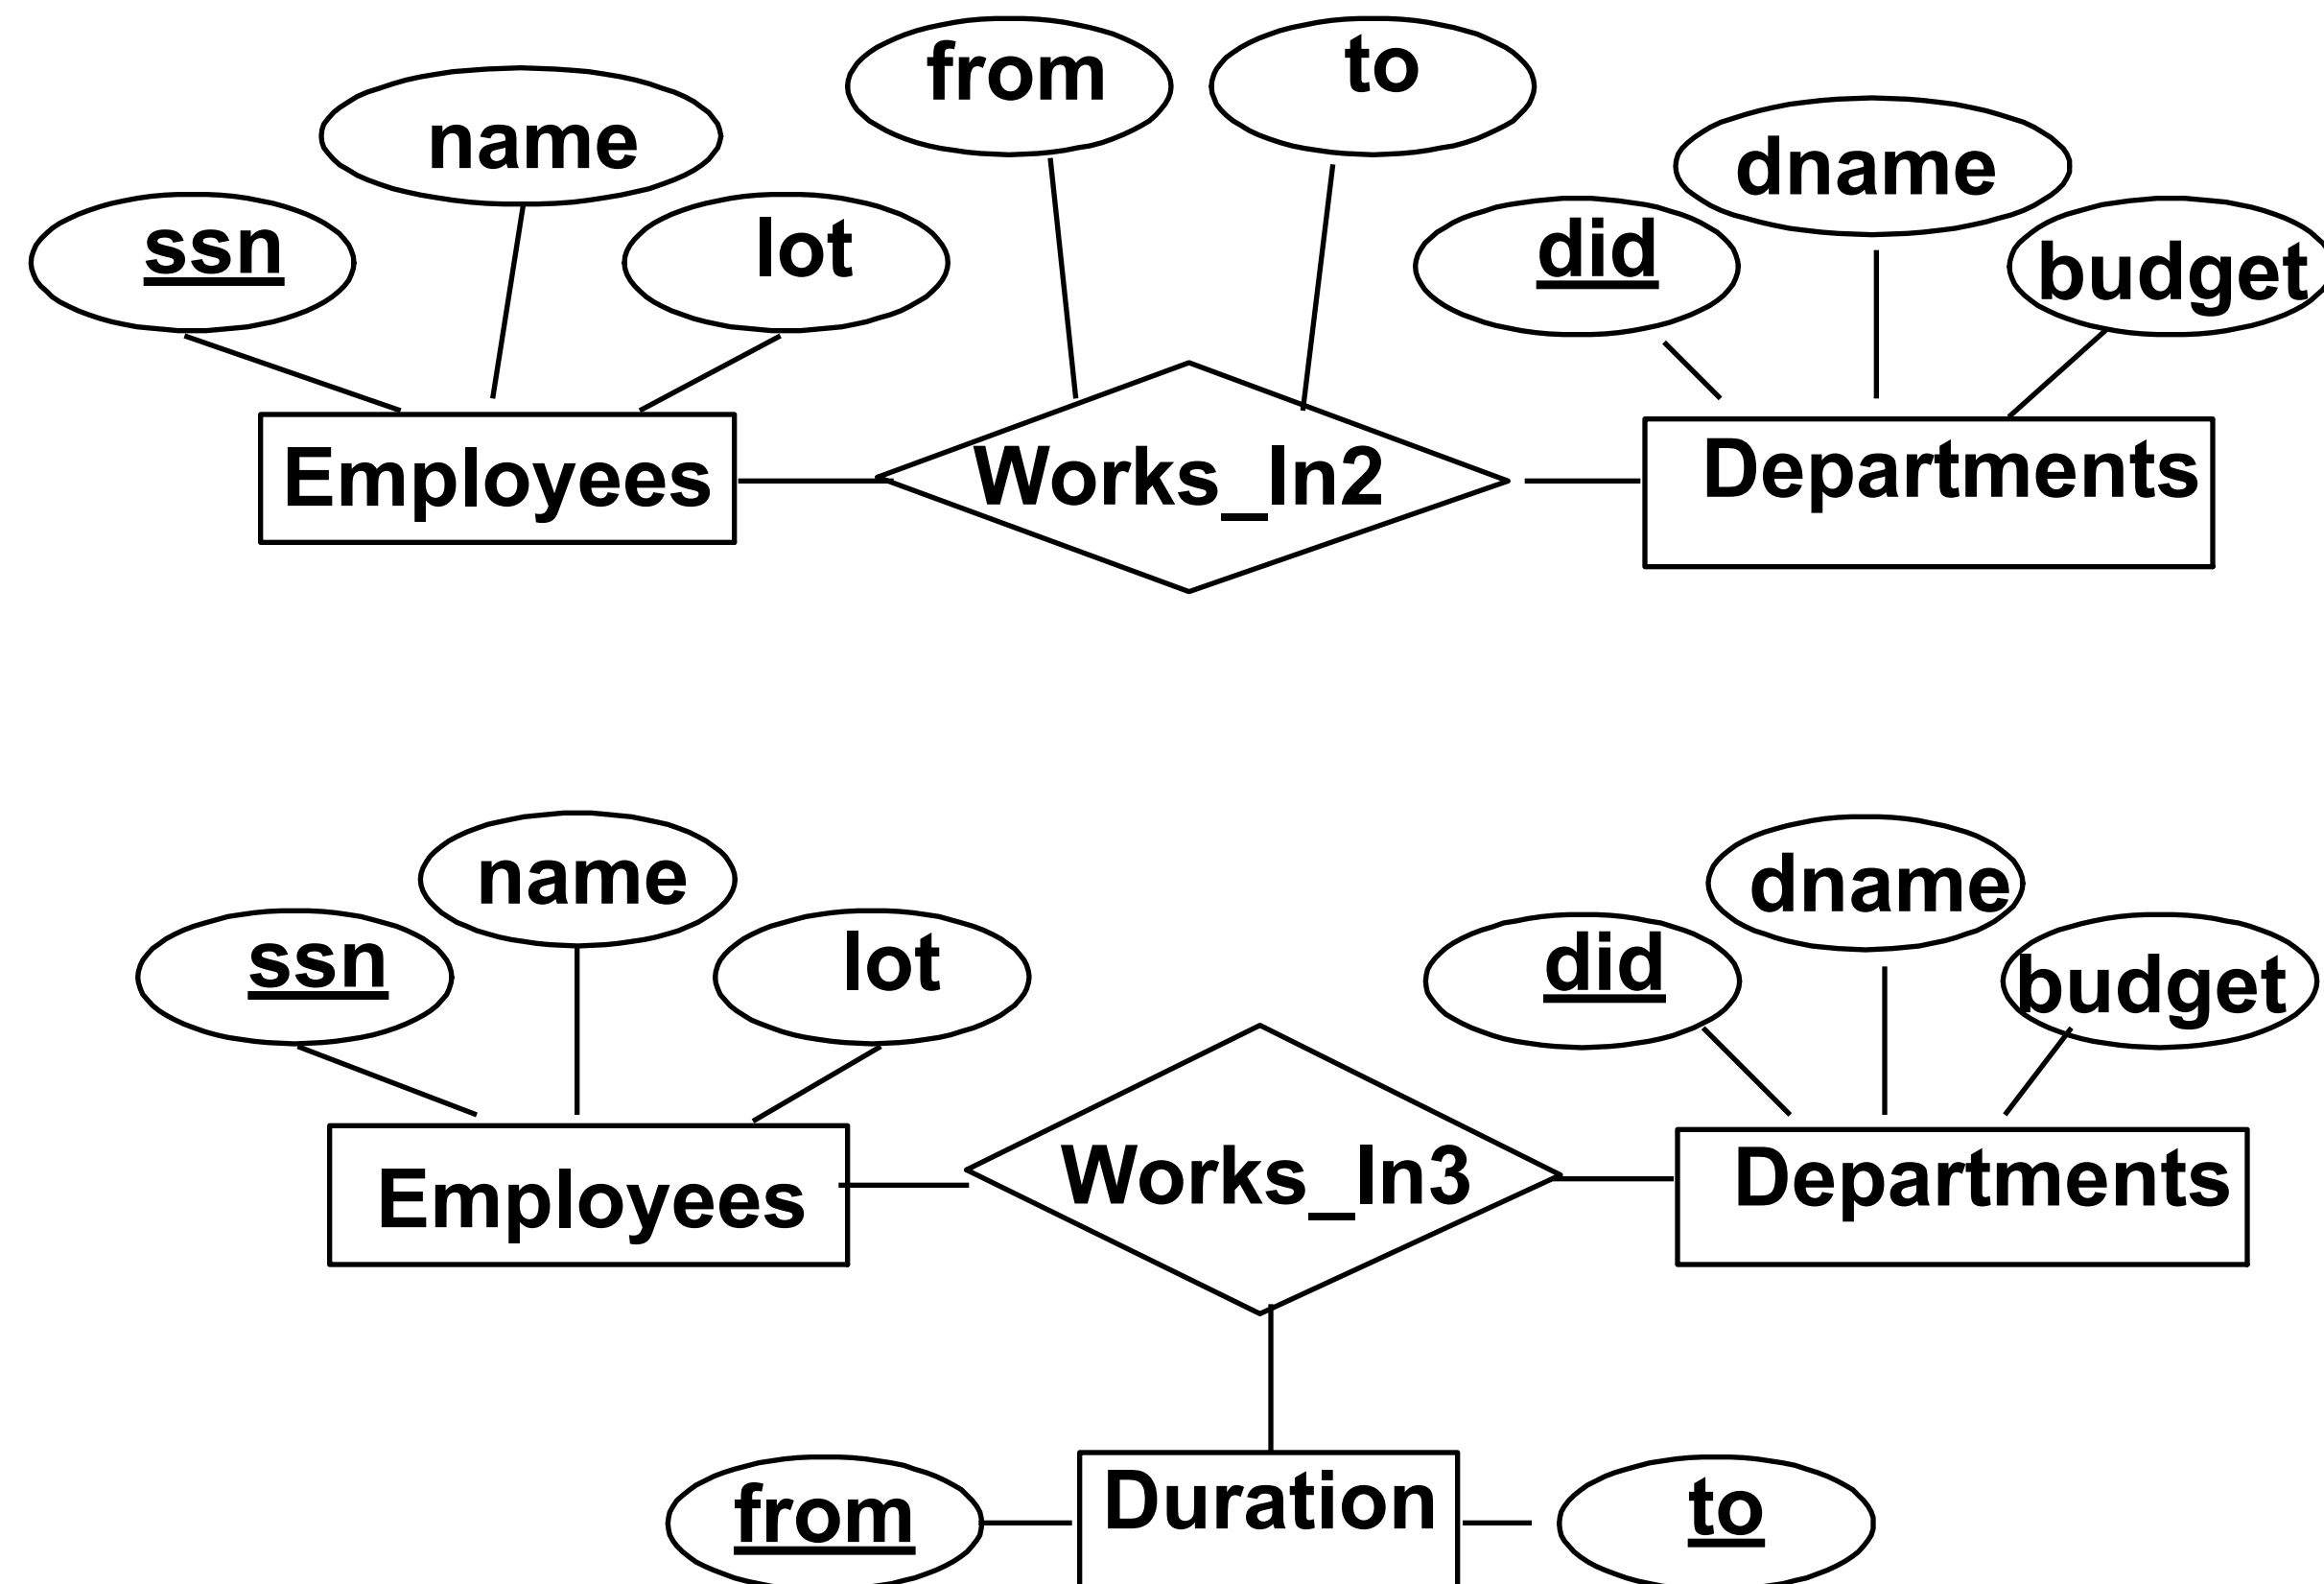
-

Relationships (Contd.)

- Relationships can have roles
- For example, if we want to capture supervisor-subordinate relationship:




How are these different?




- Attributes on relationship vs. ternary relationship

Uniqueness Constraints




- An employee can work in many departments; a department can have many employees (many-many)
- Each department has **at most one** manager, according to the *uniqueness constraint* on Manages (many-one)



Referential Integrity

- A department has **exactly one** manager, but an employee may manage **at most one** department:



- Referential-integrity constraint says that employee is required to exist in this one-one relationship

E/R Modeling Exercise


Drugwarehouse.com has hired you to design its database schema.

Here is the information that you gathered:

- Patients are identified by their SSN, and we also store their names and age.
- Doctors are identified by their SSN, and we also store their names and specialty.
- Each patient has exactly one primary care physician, and we want to know since when the patient has been with her primary care physician.
- A doctor can be a primary care physician for multiple patients.

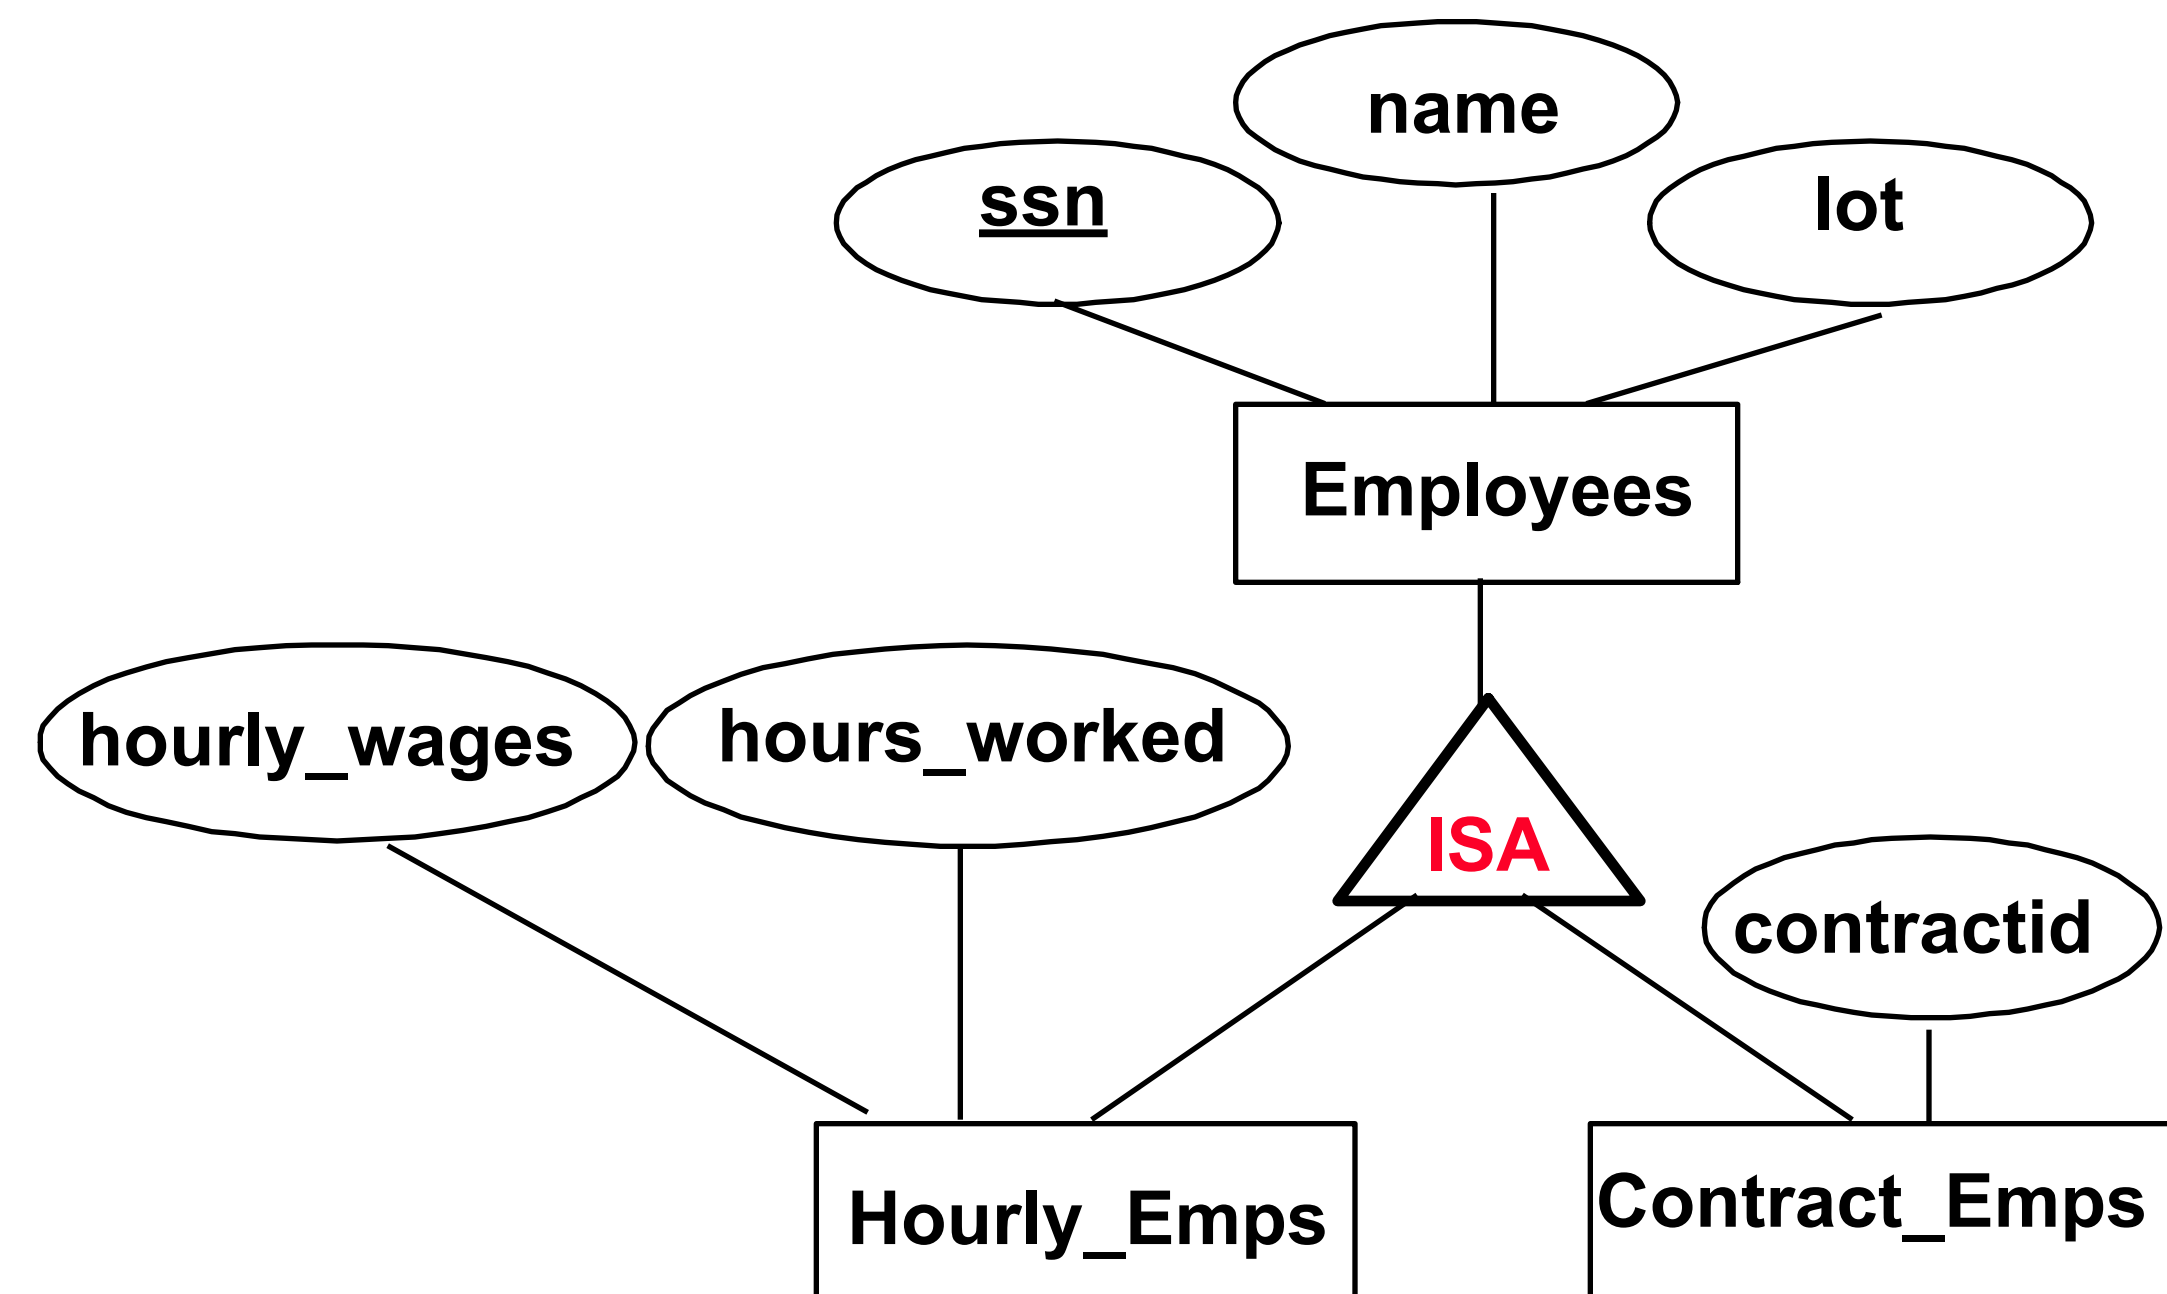
Weak Entities

- A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (one owner, many weak entities).
 - Keys of weak entity set (partially) composed of attributes from *supporting* entity sets.



ISA ('is a') Hierarchies

- Similar to object-oriented modeling
- In OO, objects are in one class only.
 - Subclasses inherit from superclasses.
- In contrast, E/R entities have representatives in all subclasses to which they belong.
 - Rule: if entity e is represented in a subclass, then e is represented in the superclass. (and recursively up the tree).



Conceptual Design Using the E/R Model


- Design principles:
 - Be faithful to the application.
 - Avoid redundancy.
 - Strive for simplicity.
- Design choices:
 - Should a concept be modeled as an entity or an attribute?
 - Should a concept be modeled as an entity or a relationship?
 - Identifying relationships: Binary or ternary?
- Constraints in the E/R Model:
 - A lot of data semantics can (and should) be captured.
 - But some constraints cannot be captured in E/R diagrams. (Which for example?)

Entity vs Attribute

- Should *address* be an attribute of Employees or an entity (connected to Employees by a relationship)?
- Depends upon the use we want to make of address information, and the semantics of the data:
 - If we have several addresses per employee, *address* must be an entity (since attributes cannot be set-valued).
 - If the structure (city, street, etc.) is important, e.g., we want to retrieve employees in a given city, *address* must be modeled as an entity (since attribute values are atomic).

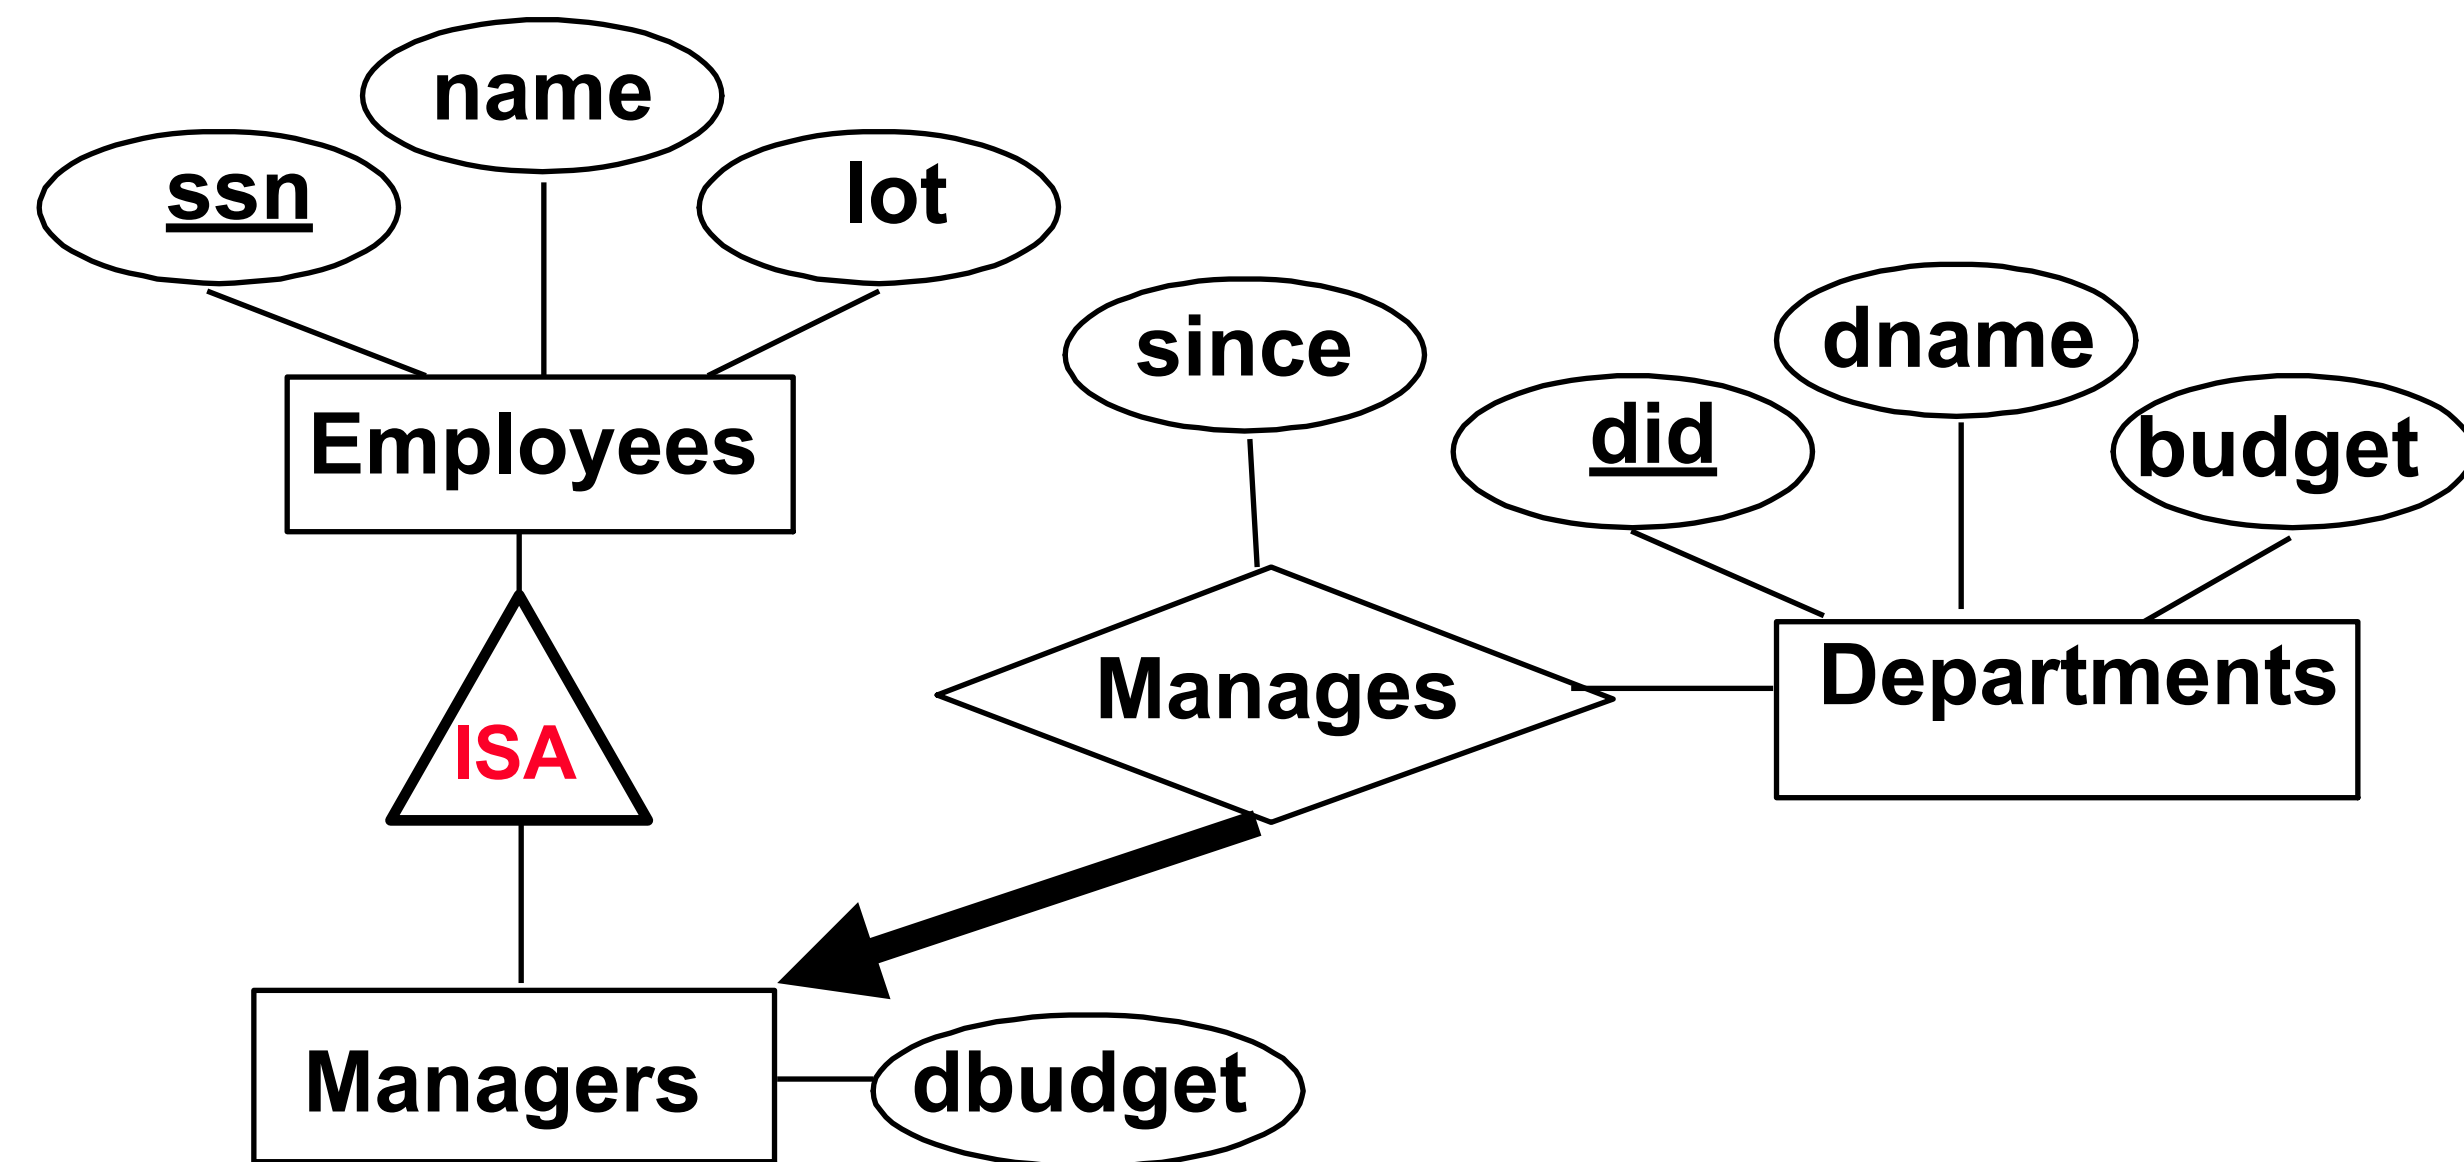
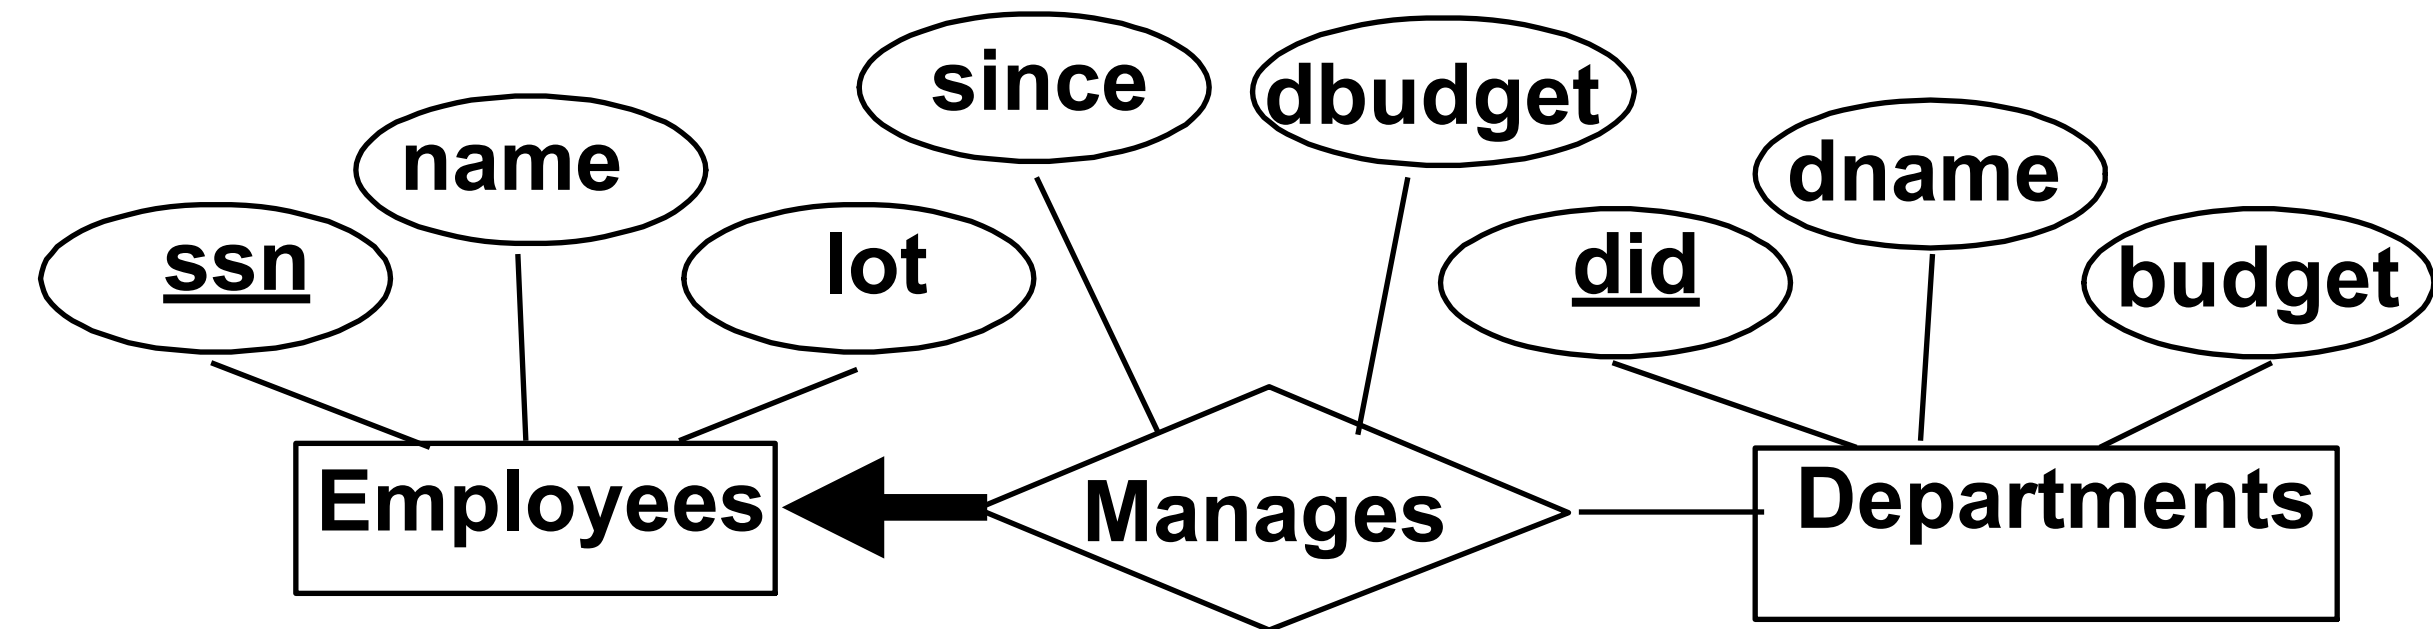
Entity vs Relationship

- E/R model OK if a manager gets a separate discretionary budget for each department.
- What if a manager gets a discretionary budget that covers *all* managed departments?
 - **Redundancy:** *dbudget* stored for each department managed by manager.
 - **Misleading:** Suggests *dbudget* associated with department-manager combination.



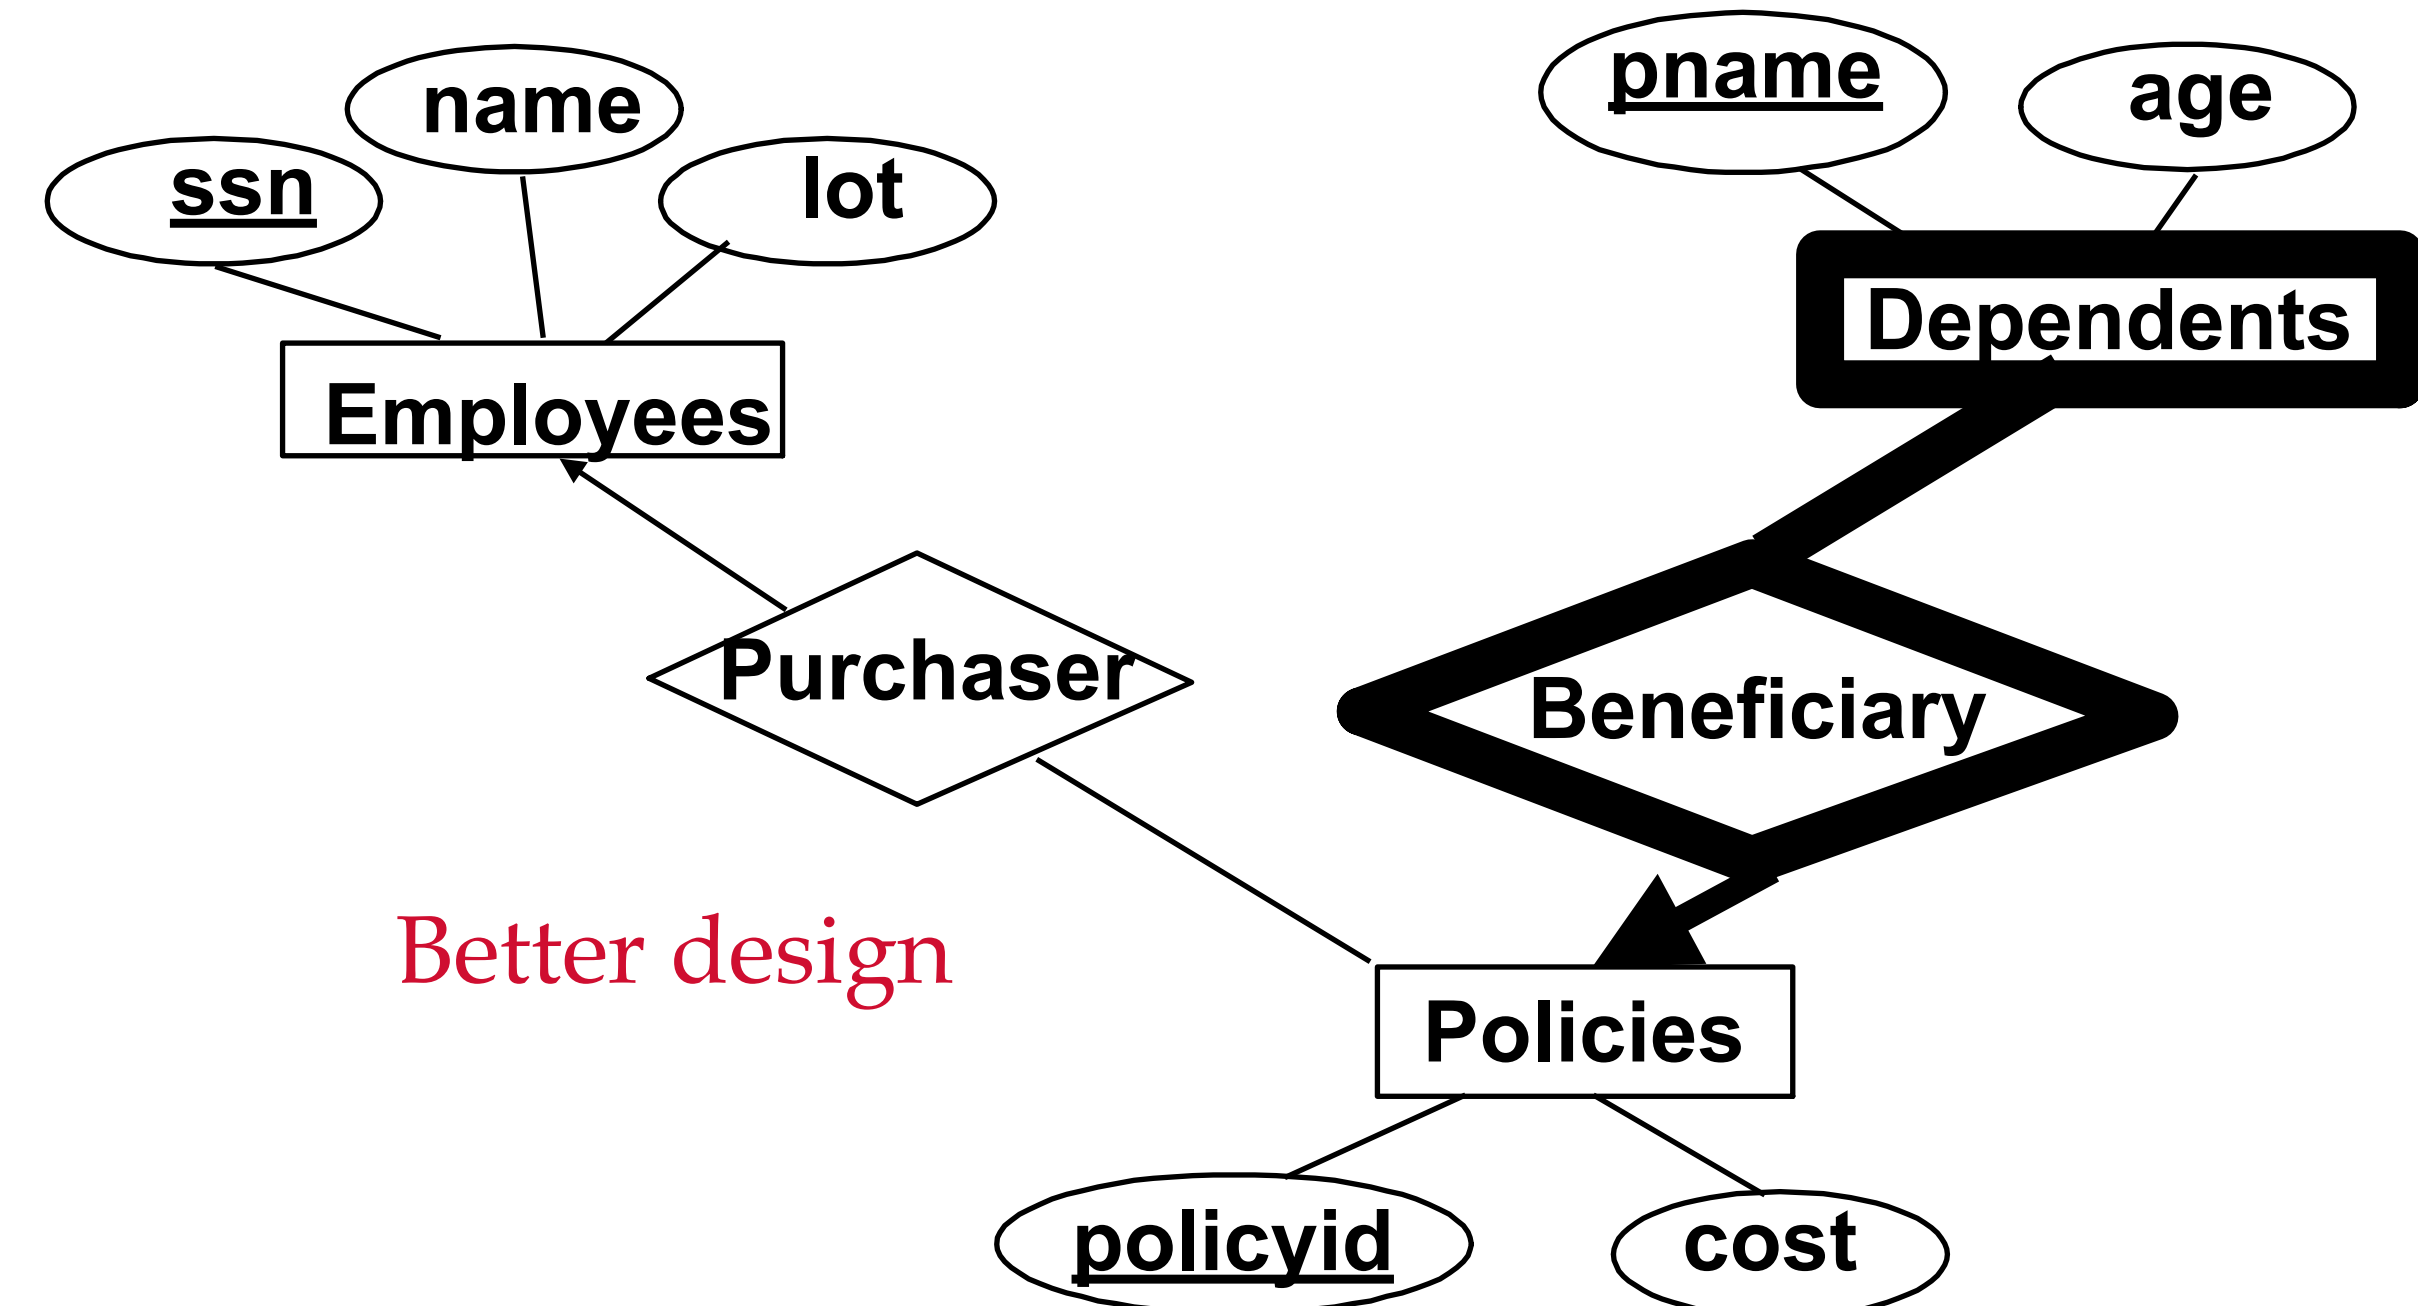
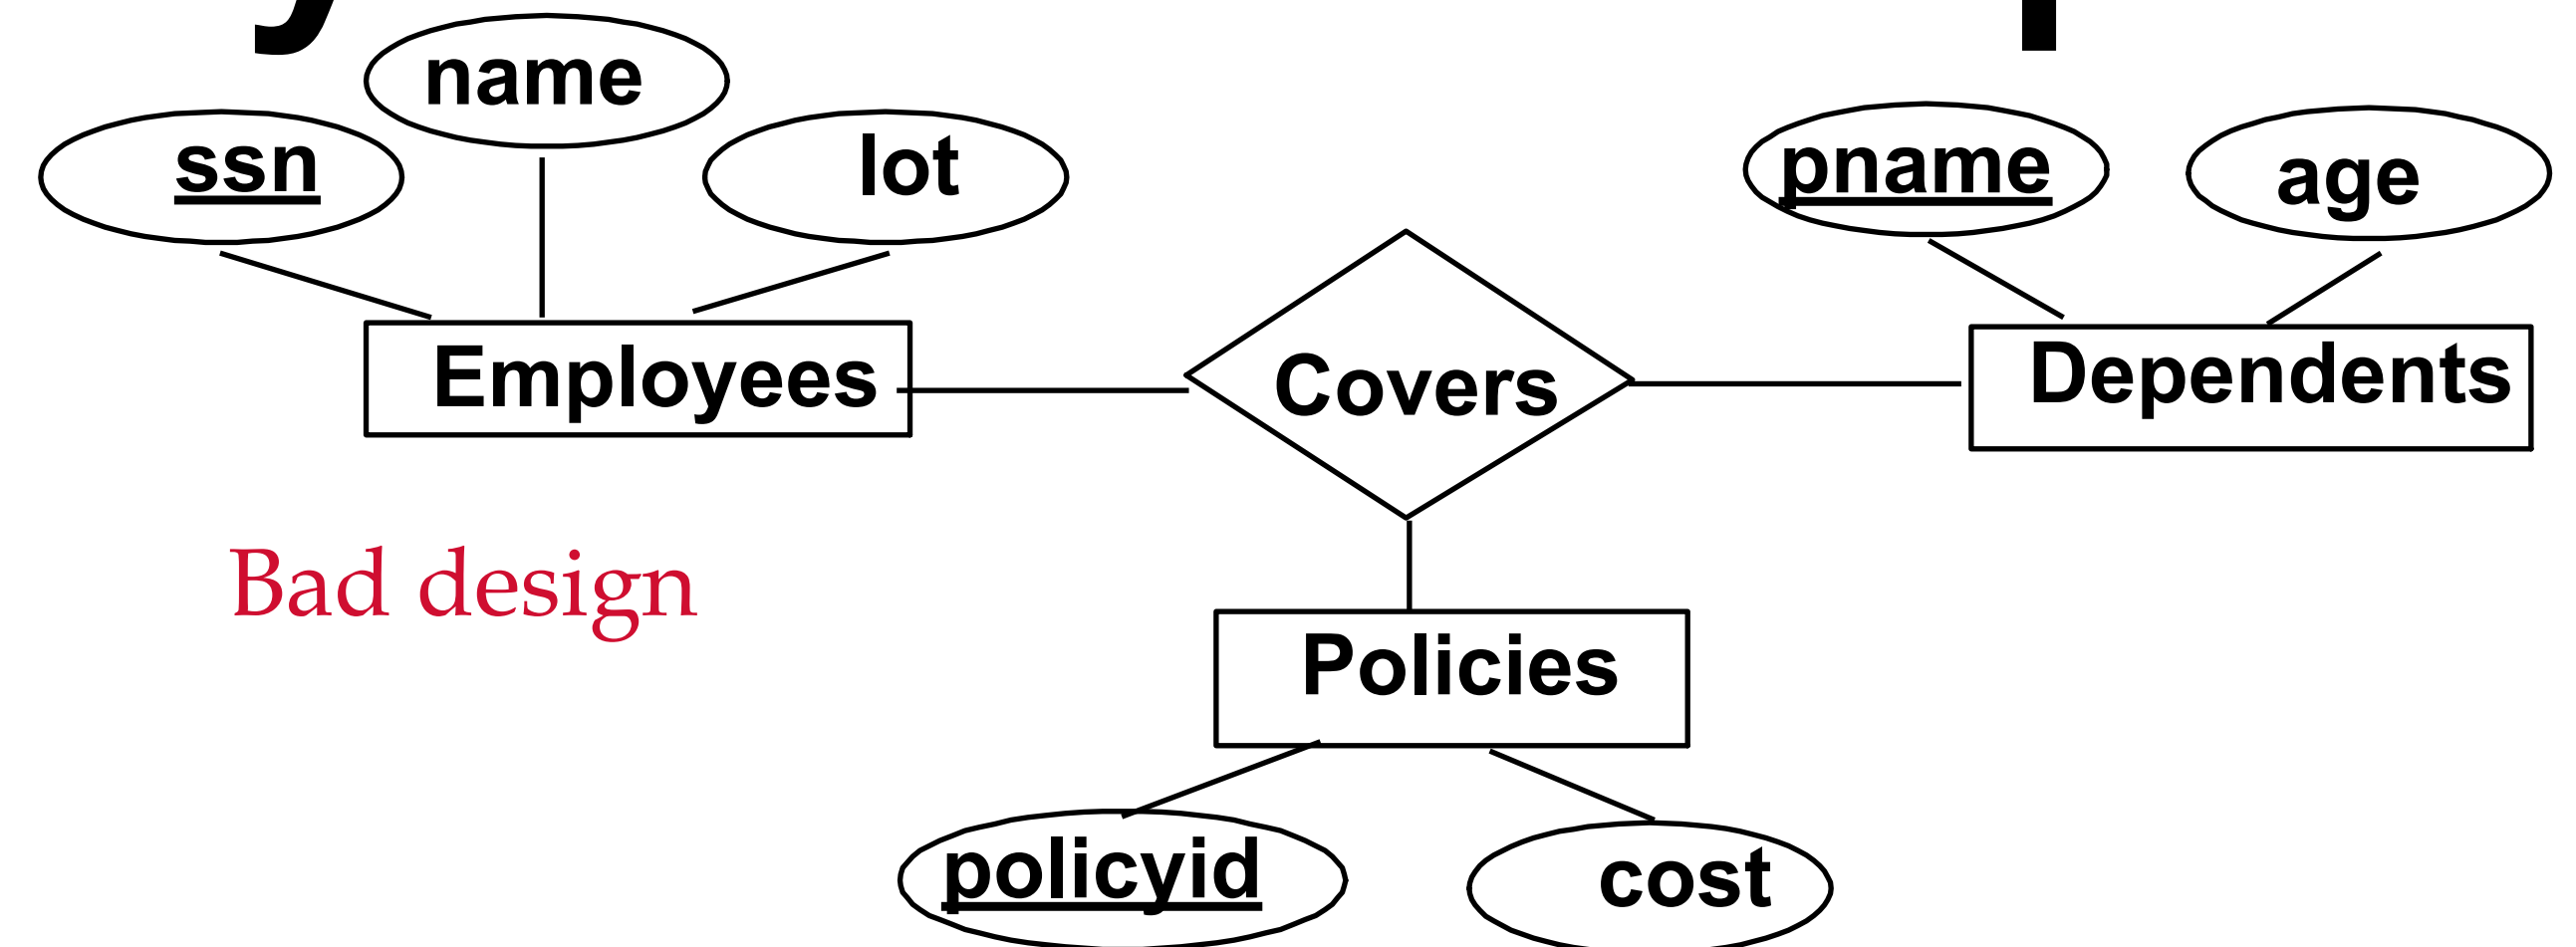
Entity vs Relationship

- E/R model OK if a manager gets a separate discretionary budget for each department.
- What if a manager gets a discretionary budget that covers *all* managed departments?
 - **Redundancy:** *dbudget* stored for each department managed by manager.
 - **Misleading:** Suggests *dbudget* associated with department-manager combination.



Binary vs Ternary Relationships

- If each policy is owned by just one employee, and each dependent is tied to the covering policy, first diagram is inaccurate.
- What are the additional constraints in the second diagram?



Binary vs Ternary Relationships (Contd.)

- Previous example illustrated a case when two binary relationships were better than one ternary relationship.
- An example in the other direction: a ternary relation **Contracts** relates entity sets **Parts**, **Departments** and **Suppliers**, and has descriptive attribute *qty*. No combination of binary relationships is an adequate substitute:
 - S “can-supply” P, D “needs” P, and D “deals-with” S does not imply that D has agreed to buy P from S.
 - How do we record *qty*?
 - However, we can model Contracts as an *entity* instead, and then use binary relationships

Summary of Conceptual Design

- Conceptual design follows requirements analysis
 - Yields a high-level description of data to be stored
- E/R model popular for conceptual design
 - Constructs are expressive, close to the way people think about their applications
- Basic constructs: *entities*, *relationships*, and *attributes* (of entities and relationships)
- Some additional constructs: *weak entities*, *ISA hierarchies*.
- Note: There are many variations on E/R model.

Summary of ER (Contd.)

- Several kinds of integrity constraints can be expressed in the ER model: *key constraints*, *uniqueness constraints*, and *referential-integrity constraints*. Some *foreign key constraints* are also implicit in the definition of a relationship set.
 - Some constraints (notably, *functional dependencies*) cannot be expressed in the ER model.
 - Constraints play an important role in determining the best database design for a use case.

Summary of E/R (Contd.)

- E/R design is **subjective**. There are often many ways to model a given scenario! Analyzing alternatives can be tricky. Common choices include:
 - Entity vs. attribute, entity vs. relationship, binary or n-ary relationship, whether or not to use ISA hierarchies, how to define constraints.
- Ensuring good database design: resulting relational schema should be analyzed and refined further.
- Functional dependencies and normalization techniques can be useful but cannot be captured in a E/R model. More on this later in the course 😊

Questions so far?

Why Study the Relational Model?

- Codd, 1969
- Most widely used model
 - Vendors: IBM, Microsoft, Oracle, Sybase/SAP, etc.
- Query languages
 - Relational calculus
 - Relational algebra
 - SQL
- “Legacy systems” in older models
 - E.G., IBM’s IMS
- Competitors
 - In the early 90s: *object-oriented model*
 - A synthesis: *object-relational model* (Oracle, DB2)
 - *Document-oriented*: XML, JSON
 - Often collected under the term *NoSQL*

Relational Database: Definitions

- **Relational database:** a set of *relations*
- **Relation:** made up of two parts:
 - **Schema:** specifies name of relation, plus name and type of each column.
 - e.g., Students(sid:string, name:string, login:string, age:integer, gpa:real)
 - **Instance:** a *table*, with rows and columns.
#Rows = *cardinality*, #fields = *degree* or *arity*.
- Can think of a relation as a **set** of rows or **tuples** (i.e., all rows are distinct).


Example Instance of Students Relation

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8


- Cardinality = 3, degree = 5, all rows distinct
- Do all columns in a relation instance have to be distinct?

Logical Design: E/R to Relational

- Entity sets to tables.



```
CREATE TABLE Employees
(ssn CHAR(11),
 name CHAR(20),
 lot INTEGER,
PRIMARY KEY (ssn))
```



```
Employees
(ssn : string,
 name : string,
 lot : int)


https://traytel.bitbucket.io/rc-eval/
```

Example Instance

Employees

<u>ssn</u>	name	lot
0983763423	John	10
9384392483	Jane	10
3743923483	Jill	20

```
INSERT INTO Employees(ssn,name,lot)
VALUES (0983763423, 'John', 10), (9384392483, 'Jane', 10), (3743923483, 'Jill', 20);
```



```
Employees
("0983763423", "John", 10) ("9384392483", "Jane", 10) ("3743923483", "Jill", 20)
```

<https://traytel.bitbucket.io/rc-eval/>

Integrity Constraints (ICs)

- **IC:** condition that must be true for *any* instance of the database
 - Domain constraints
 - Key constraints
 - Foreign key constraints (later)
- A *legal* instance of a relation is one that satisfies all specified ICs.
 - DBMS should not allow illegal instances
 - Avoids data entry errors too!

Primary Key Constraints

- A set of fields is a superkey for a relation if:
 - No two distinct tuples can have same values in all fields belonging to the superkey
- A set of fields is a key if:
 - The set of fields is a superkey
 - No proper subset of the set of fields is a superkey
- If there is >1 key for a relation, one of the keys is chosen (by DBA) to be the *primary key*.
- E.g., {ssn} is a key for Employees. (What about {name}?) The set {ssn, name} is a superkey.

What does this mean?

```
CREATE TABLE Enrolled  
  (sid CHAR(20)  
   cid CHAR(20),  
   grade CHAR(2),  
   PRIMARY KEY (sid,cid))
```

Candidate Keys

- Possibly many *candidate keys* (specified using **UNIQUE**), one of which is chosen as the *primary key*.



```
CREATE TABLE Enrolled  
  (sid CHAR(20)  
   cid CHAR(20),  
   grade CHAR(2),  
   PRIMARY KEY (sid),  
   UNIQUE (cid, grade))
```

- Each student is enrolled in at most one course
- No two students in a course get the same grade


Where do ICs Come From?

- ICs are based upon the semantics of the real-world enterprise that is being described in the database relations.
- We can check a database instance to see if an IC is violated, but we can **NEVER** infer that an IC is true by looking at an instance.
 - An IC is a statement about *all possible* instances!
 - From example, we know *name* is not a key, but the assertion that *sid* is a key is given to us.
- Key and foreign key ICs are the most common; more general ICs supported too.

E/R to Relational (contd.)



E/R to Relational (contd.)



Relationship Sets to Tables

```
CREATE TABLE Employees
    (ssn CHAR(11),
     name CHAR(20),
     lot  INTEGER,
     PRIMARY KEY (ssn))
```

```
CREATE TABLE Departments
    (did INTEGER,
     dname CHAR(20),
     budget FLOAT,
     PRIMARY KEY (did))
```

Relationship Sets to Tables

```
CREATE TABLE Employees
    (ssn CHAR(11),
     name CHAR(20),
     lot  INTEGER,
     PRIMARY KEY (ssn))
```

```
CREATE TABLE Departments
    (did INTEGER,
     dname CHAR(20),
     budget FLOAT,
     PRIMARY KEY (did))
```

```
CREATE TABLE Works_In(
    ssn  CHAR(11),
    did  INTEGER,
    since DATE,
    PRIMARY KEY (ssn, did),
    FOREIGN KEY (ssn) REFERENCES Employees,
    FOREIGN KEY (did) REFERENCES Departments)
```


Example Instance

Employees

<u>ssn</u>	name	lot
0983763423	John	10
9384392483	Jane	10
3743923483	Jill	20

Departments

<u>did</u>	dname	budget
101	Sales	10K
105	Purchasing	20K
108	Databases	1000K

Works_In

<u>ssn</u>	<u>did</u>	since
0983763423	101	1 Jan 2003
0983763423	108	2 Jan 2003
9384392483	108	1 Jun 2002

Foreign Keys, Referential Integrity

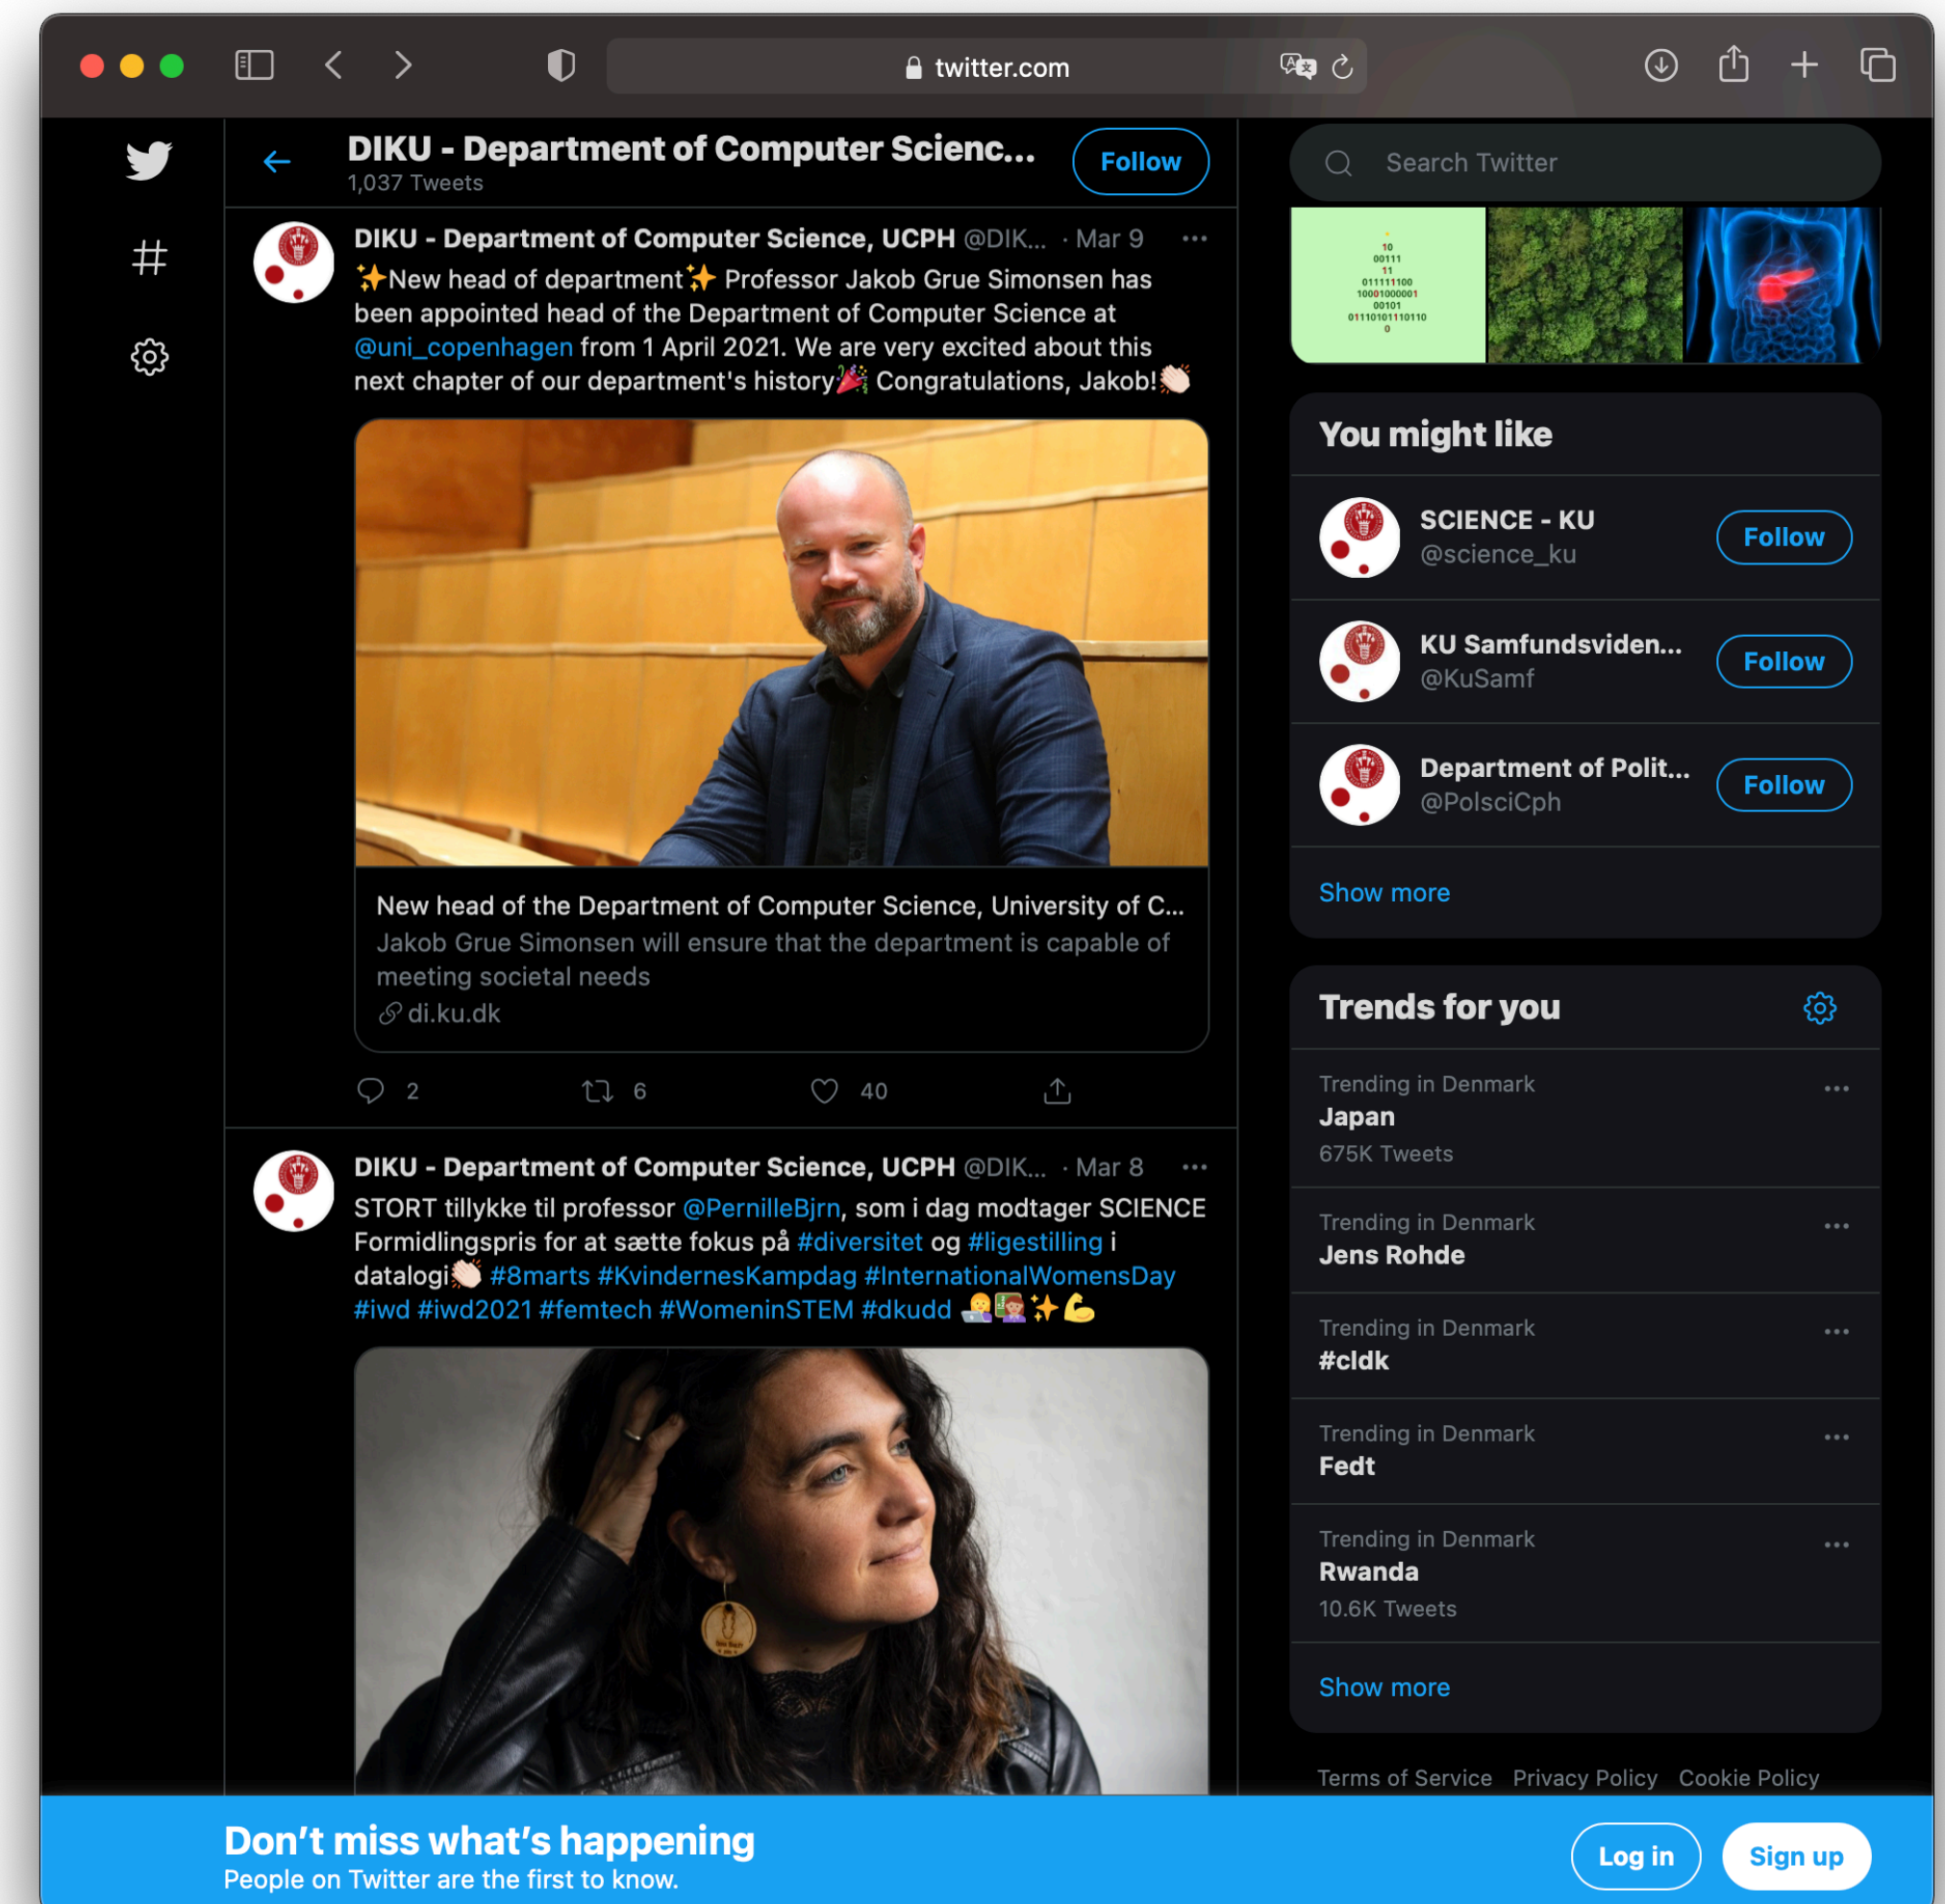
- Foreign key: Set of fields in one relation that is used to `refer' to a tuple in another relation
 - Must correspond to primary key of the second relation
 - Like a `logical pointer'.
- If all foreign key constraints enforced, referential integrity is achieved, i.e., no dangling references.
 - Not like HTML links!

Enforcing Referential Integrity

- What if a new “Works_In” tuple is added that references a non-existent employee?
 - Reject it!
- What if an Employee tuple is deleted?
 - Also delete all Works_In tuples that refer to it.
 - Disallow deletion of Employee tuple that is referred to.
 - Set *ssn* to some default value
 - Set *ssn* in Works_In to *null*, denoting ‘unknown’
- Similar if primary key of Employee tuple is updated

Discussion: Relational Twitter

- Let us say you would like to model two tables from Twitter:
 - Users
 - Tweets
- Discuss
 - Which attributes should each table have?
 - What are the keys in these tables? Would you also need foreign keys? If so, which?
 - Do not model Followers for now! But you may think about it. 😊
 - If you have time, try to write the SQL



What should we learn today?

- Explain the concepts of entity (set), relationship (set) and express these concepts in entity-relationship (E/R) diagrams
- Explain and express constraints (key, uniqueness, and ref. integrity) in E/R diagrams
- Explain and capture in E/R diagrams weak entity sets and ISA hierarchies
- Analyze trade-offs and argue for the advantages and disadvantages of a E/R design
- Explain what a data model is, along with the distinction between schemas and instances
- Define and explain the relational data model
- Define and explain the main classes of integrity constraints in the relational model, in particular key and foreign key constraints
- Model relational schemas and express them in SQL

