

Databases and Information Systems

Views, Triggers, Transactions, Indices

Dmitriy Traytel

slides partly by Marcos Vaz Salles



UNIVERSITY OF
COPENHAGEN

Boyce-Codd Normal Form (BCNF)

- Relation R with FDs F is in **BCNF** if, for all $X \rightarrow Y$ in F^+
 - $Y \subseteq X$ (called a **trivial** FD), or
 - X contains a key for R .
- In other words, R is in BCNF if the only non-trivial FDs that hold over R are key constraints.
 - No dependency in R that can be predicted using FDs alone.
 - If we are shown two tuples that agree upon the A value, we cannot infer the C value in one tuple from the C value in the other.
 - If example relation was in BCNF and $A \rightarrow C$, the 2 tuples would have to be identical (A is a key).

A	B	C
10	1	3
10	2	?


Decomposition into BCNF

- Consider relation R with FDs F . If $X \rightarrow Y$ violates BCNF, decompose R into $R - Y$ and XY .
 - Repeated application of this idea will give us a collection of relations that are in BCNF; lossless-join decomposition, and guaranteed to terminate.
 - e.g., CSJDPQV, key C , $JP \rightarrow C$, $SD \rightarrow P$, $J \rightarrow S$
 - To deal with $SD \rightarrow P$, decompose into SDP, CSJDQV.
 - To deal with $J \rightarrow S$, decompose CSJDQV into JS and CJDQV
- In general, several dependencies may cause violation of BCNF. The order in which we “deal with” them could lead to different sets of relations!

Exercise

- Courses(course,teacher,hour,room,student,grade)
- Abbreviated CTHRSG
- FDs: $C \rightarrow T$, $HR \rightarrow C$, $HT \rightarrow R$, $HS \rightarrow R$, $CS \rightarrow G$
- Determine all the keys
- Decompose into BCNF

What should we learn today?



- Explain the concept of **(materialized) views** and create them in SQL
- Explain the concept of **triggers** and create them in SQL
- Argue for when to use **(materialized) views** vs. **triggers**
- Explain the **ACID** properties of transactions
- Formulate **transaction programs** in SQL
- Identify the main types of representations and access methods for relations in DBMS, namely **heap files** and **indexes**
- Explain the core algorithms for insertion, deletion, and search in a **B⁺-tree**

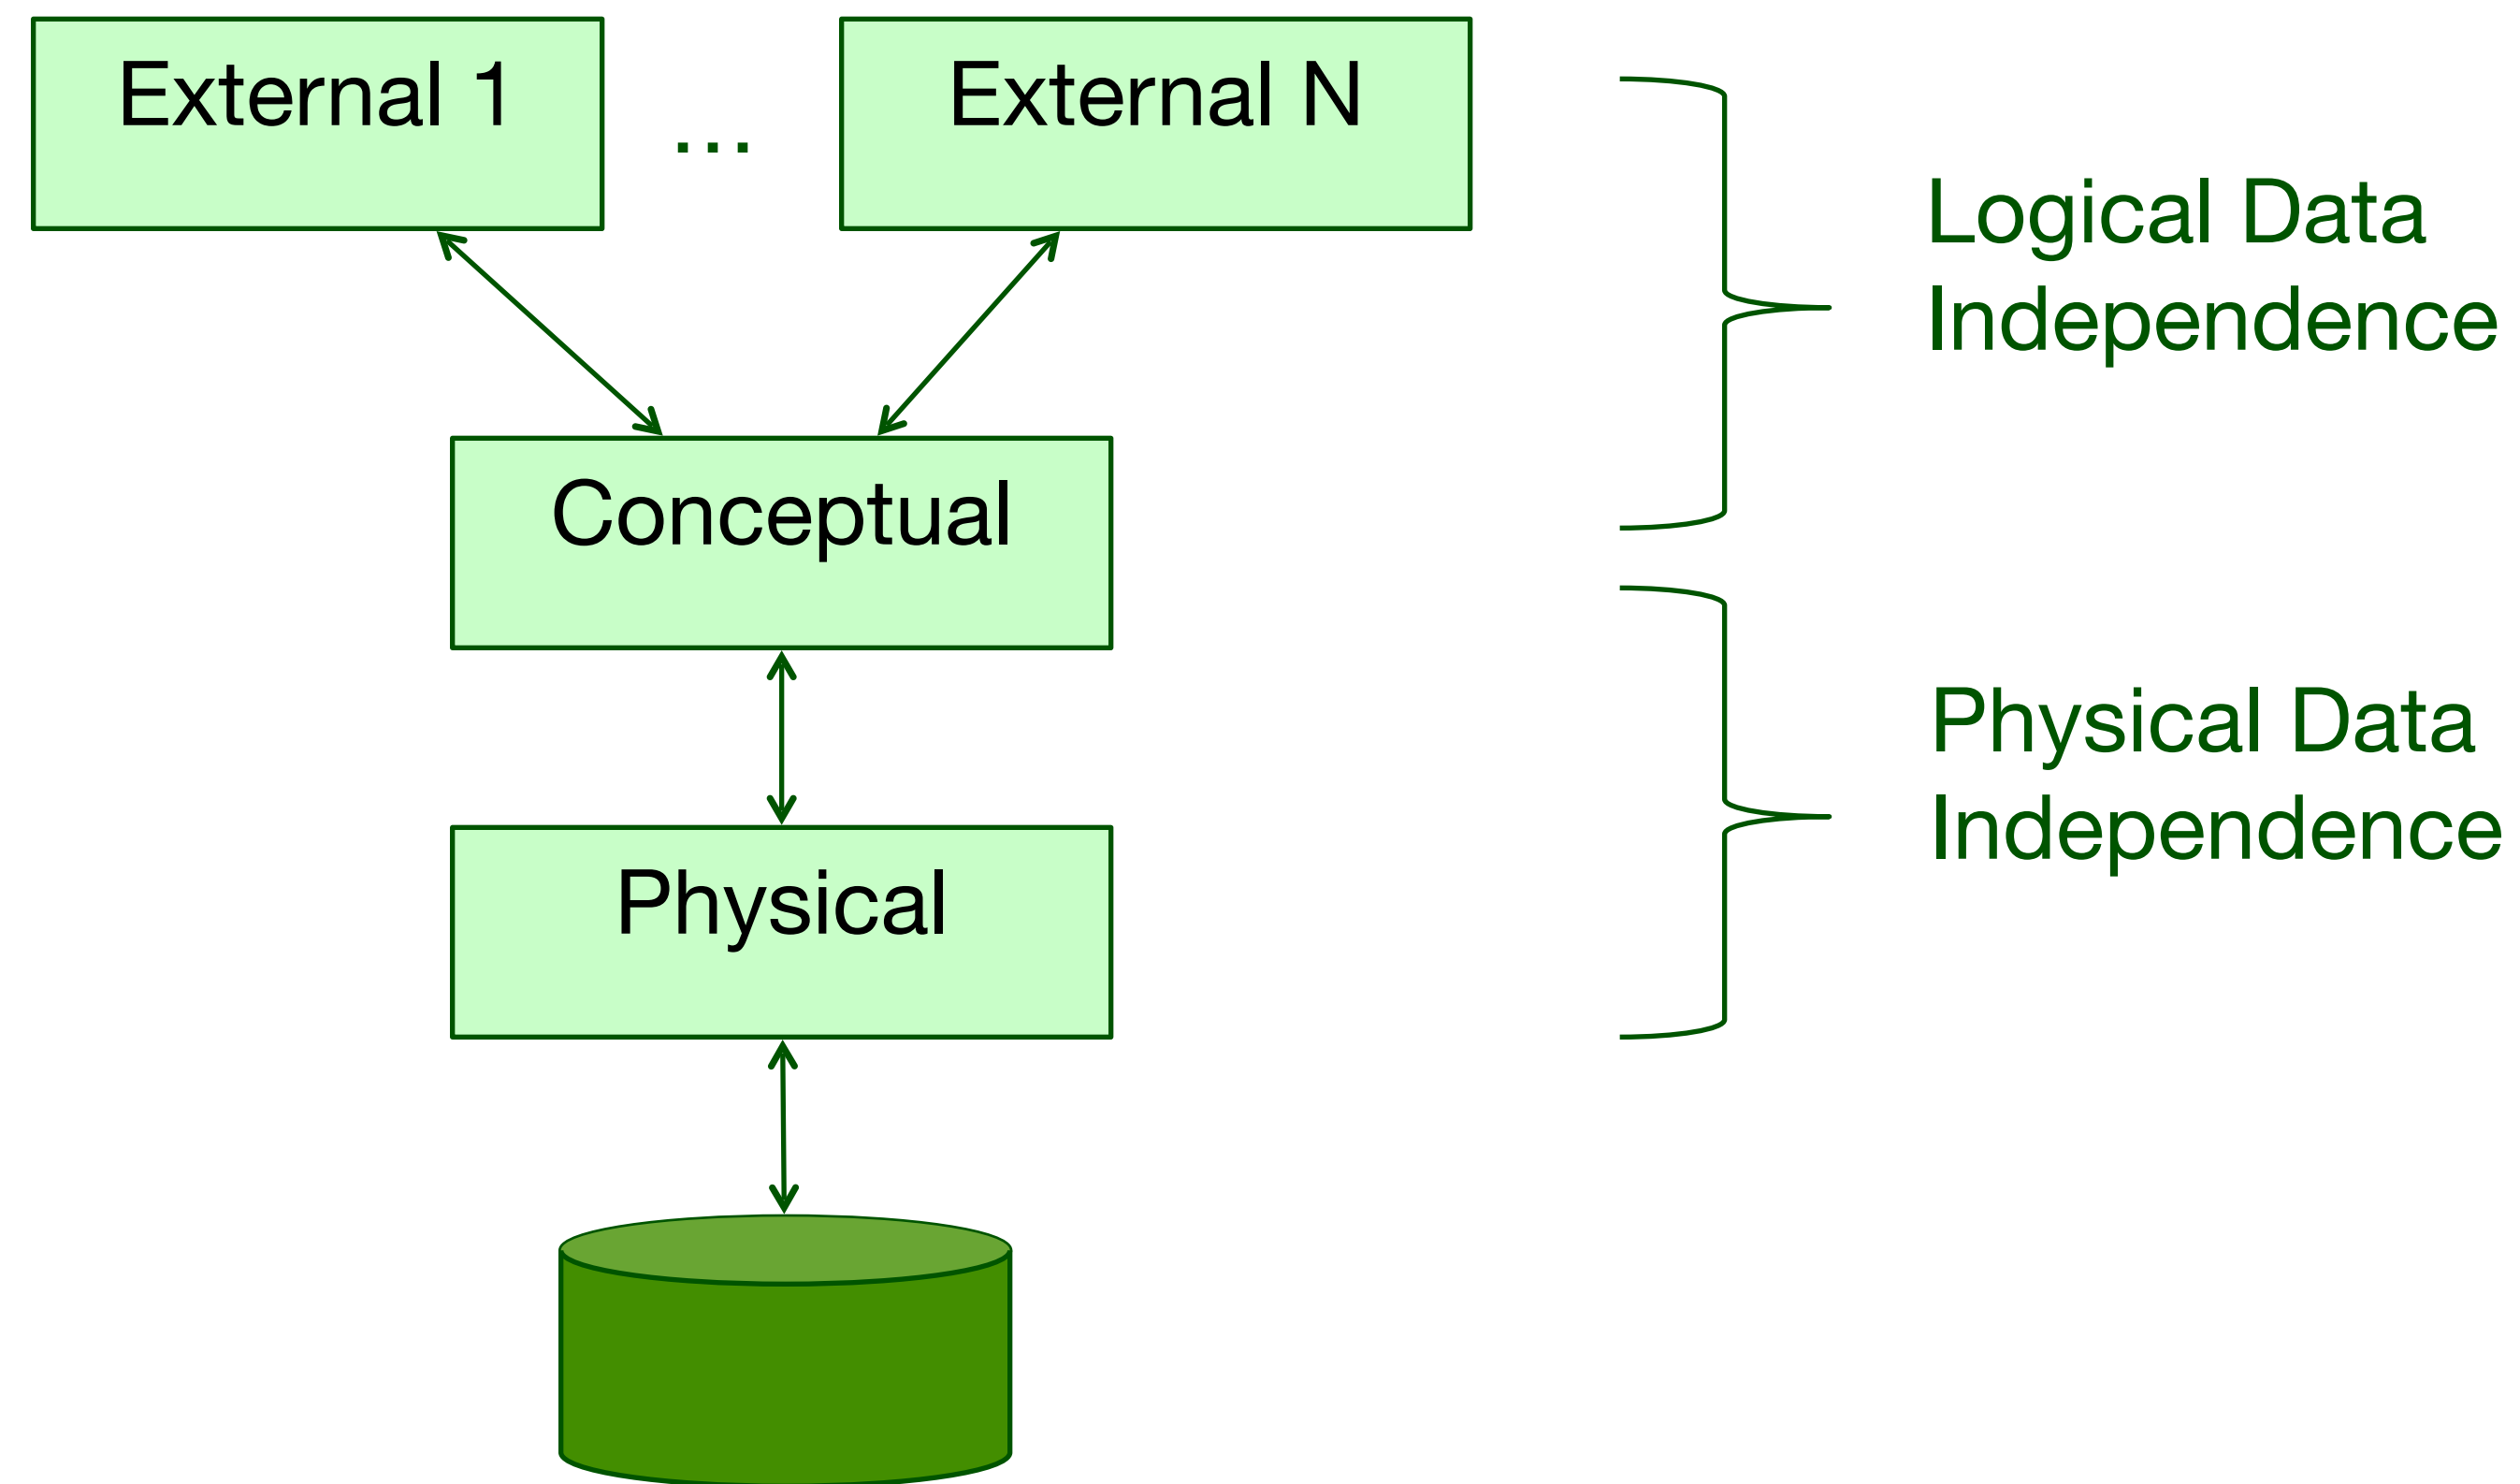
Views

- Special construct to add logical table to **schema** of the database

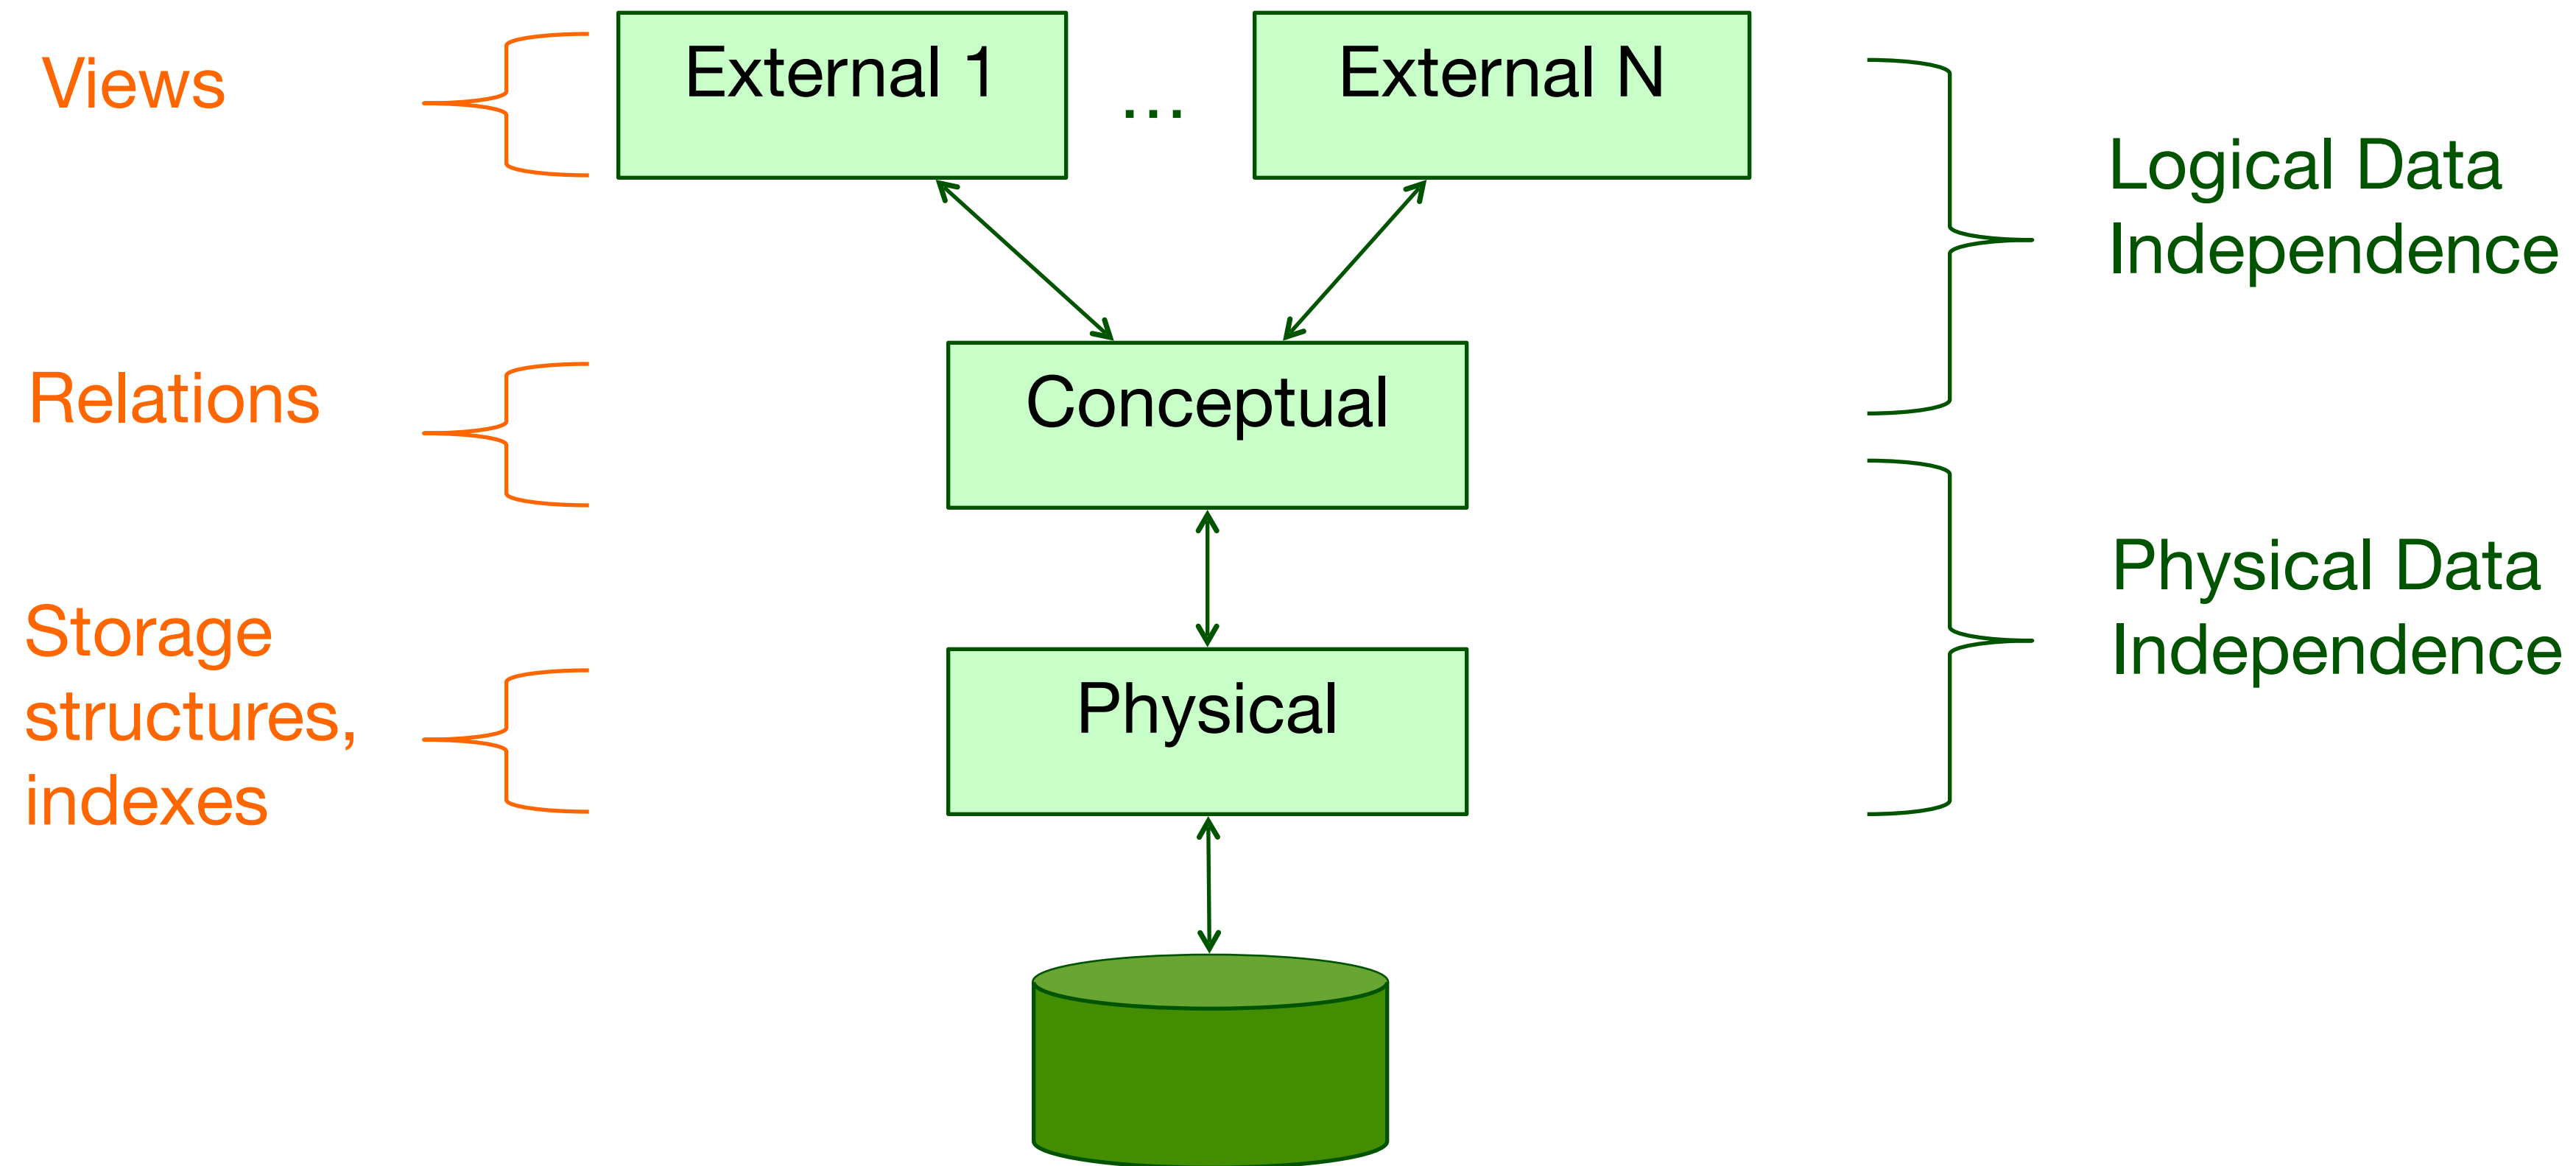
```
CREATE VIEW Red_Green_Sailors (sid, sname) AS
  SELECT DISTINCT S.sid, S.sname
  FROM   Sailors S, Boats B, Reserves R
  WHERE  S.sid=R.sid AND R.bid=B.bid
        AND (B.color='red' OR B.color='green')
```

```
DROP VIEW Red_Green_Sailors
```

Remember this picture?



Remember this picture?



Why use views?

- Hide some data from some users
- Make some queries easier / more natural
- Modularity of database access

Real applications tend to use lots and lots (and lots and lots!) of views

Why use views?

- Hide some data from some users
- Make some queries easier / more natural
- Modularity of database access

Real applications tend to use lots and lots (and lots and lots!) of views

- Virtual vs. Materialized Views
 - + Improve query performance
 - + Transparency to applications
 - Materialized view could be LARGE!
 - View maintenance under updates

Materialized View Example

```
CREATE MATERIALIZED VIEW CA-CS AS  
SELECT C.cName, S.sName  
FROM College C, Student S, Apply A  
WHERE C.cName = A.cName AND S.sID = A.sID  
AND C.state = 'CA' AND A.major = 'CS'
```

- Can use CA-CS as if it's a table (it is!)
 - DBMS will store a copy of result of view query

Materialized View Example

```
CREATE MATERIALIZED VIEW CA-CS AS
SELECT C.cName, S.sName
FROM College C, Student S, Apply A
WHERE C.cName = A.cName AND S.sID = A.sID
AND C.state = 'CA' AND A.major = 'CS'
```

- Can use CA-CS as if it's a table (it is!)
 - DBMS will store a copy of result of view query
- But: Modifications to base data invalidate view
 - DBMS can be given a policy to refresh view, e.g., daily
 - Incremental maintenance possible in some cases

Materialized View Example

```
CREATE MATERIALIZED VIEW CA-CS AS
SELECT C.cName, S.sName
FROM College C, Student S, Apply A
WHERE C.cName = A.cName AND S.sID = A.sID
AND C.state = 'CA' AND A.major = 'CS'
```

- Can use CA-CS as if it's a table (it is!)
 - DBMS will store a copy of result of view query
- But: Modifications to base data invalidate view
 - DBMS can be given a policy to refresh view, e.g., daily
 - Incremental maintenance possible in some cases
- Also: If you can update the view, base tables must stay in sync
 - DBMS must propagate updates through view (not always possible...)

View Maintenance

- Two steps:
 - **Propagate**: Compute changes to view when data changes.
 - **Refresh**: Apply changes to the materialized view table.
- **Maintenance policy**: Controls when we do refresh.
 - **Immediate**: As part of the operation that modifies the underlying data tables.
 - + Materialized view is always consistent
 - Updates are slowed down
 - **Deferred**: Some time later, as a separate operation.
 - View becomes inconsistent
 - + Can scale to maintain many views without slowing down updates

View Maintenance

- Two steps:
 - **Propagate**: Compute changes to view when data changes.
 - **Refresh**: Apply changes to the materialized view table.
- **Maintenance policy**: Controls when we do refresh.
 - **Immediate**: As part of the operation that modifies the underlying data tables.
 - + Materialized view is always consistent
 - Updates are slowed down
 - **Deferred**: Some time later, as a separate operation.
 - View becomes inconsistent
 - + Can scale to maintain many views without slowing down updates
 - **Lazy**: Delay refresh until next query on view; then refresh before answering the query.
 - **Periodic (Snapshot)**: Refresh periodically. Queries possibly answered using outdated version of view tuples. Widely used, especially for asynchronous replication in distributed databases, and for warehouse applications.
 - **Event-based**: E.g., refresh after a fixed number of updates to underlying data tables.

Triggers

- “Event-Condition-Action Rules”
 - When event occurs, check condition; if true, do action
- Why would we need triggers?
 - 1) Move monitoring logic from application into DBMS
 - 2) Enforce constraints
 - Beyond what constraint system supports
 - Automatic constraint “repair”
- Implementations vary significantly
 - Different DBMS may support different kinds of events and actions
 - Different DBMS may allow triggers to be written in different languages (e.g., through stored procedures)

Triggers in SQL

```
CREATE TRIGGER name
BEFORE|AFTER|INSTEAD OF events

[ referencing-variables ]

[ FOR EACH ROW ]

WHEN ( condition )

action
```

A Simple Example of a Trigger

```

CREATE TRIGGER AUR_NetWorthTrigger
    AFTER UPDATE OF netWorth ON MovieExec
    REFERENCING
        OLD ROW AS OldTuple
        NEW ROW AS NewTuple
    FOR EACH ROW
    WHEN (OldTuple.netWorth > NewTuple.netWorth)
        UPDATE MovieExec
        SET netWorth = OldTuple.netWorth
        WHERE cert# = NewTuple.cert#;
    
```

Trigger name in database schema

Event: After update of attribute

Access to tuple values

Row level trigger

Condition

Action

Prevent lower values

Reset to old value

NOTE: Not PostgreSQL syntax

- **Granularity**

- **Row-level:** Event = Change of a single row (a single UPDATE statement might result in multiple events)
- **Statement-level:** Event = Statement (a single UPDATE statement that changes multiple rows is a single event).

Aggregate Maintenance with Triggers

Orders(order_num, item_num, quantity, store_id, vendor_id)
Items(item_num, price)
VendorOutstanding(vendor_id, amount)

- Insertions

```
INSERT INTO orders VALUES (1000350,7825,100,'xxxxxx6944','vendor4');
```

- Queries (first without, then with redundant tables)

```
SELECT vendor_id, SUM(quantity * price)
FROM orders NATURAL JOIN items
GROUP BY vendor_id;
```

vs.

```
SELECT * FROM vendorOutstanding;
```

Aggregate Maintenance with Triggers (contd.)

```
CREATE TRIGGER
trUpsertVendorOutstanding
AFTER INSERT ON orders
FOR EACH ROW
EXECUTE PROCEDURE
pyUpsertVendorOutstanding();
```

```
CREATE OR REPLACE FUNCTION pyUpsertVendorOutstanding()
RETURNS trigger
AS $$
    # new values inserted
    new_quantity = TD["new"]["quantity"]
    new_item_num = TD["new"]["item_num"]
    new_vendor_id = TD["new"]["vendor_id"]

    # Python code that accesses DB follows
    ...
$$ LANGUAGE plpython3u;
```

See script [trigger-example.sql](#) in Absalon

Aggregate Maintenance with Triggers (contd.)

```
CREATE OR REPLACE FUNCTION pyUpsertVendorOutstanding()
RETURNS trigger
AS $$
    # new values inserted
    new_quantity = TD["new"]["quantity"]
    new_item_num = TD["new"]["item_num"]
    new_vendor_id = TD["new"]["vendor_id"]

    # prepare upsert SQL statement
    upsert_stmt = plpy.prepare(
        ("INSERT INTO vendor_outstanding "
         "as v(vendor_id, amount) VALUES "
         "($1, $2 * (SELECT items.price "
         "      FROM items WHERE items.item_num = $3)) "
         "ON CONFLICT (vendor_id) DO "
         "UPDATE SET amount = v.amount + EXCLUDED.amount"),
        ["text", "int", "int"])

    # execute upsert on vendor_outstanding
    plpy.execute(upsert_stmt, [new_vendor_id, new_quantity, new_item_num])
$$ LANGUAGE plpython3u;
```

Other Uses of Triggers

- Maintaining an audit trail or history of modifications
 - In an auction, record audit trail of bids per user
- Automatically populating attributes
 - Whenever a bid is updated by a user, record the current time (ignoring any other time value given by application)
- Making updates conditional
 - Disallow users to revise bids down, only allow bids to be revised to be higher than current value

Tricky Issues in Triggers

- Row-level vs. Statement-level
 - **New/Old Row** and **New/Old Table**
 - **Before, Instead Of, After**
- Multiple triggers activated at same time
- Trigger actions activating other triggers (chaining)
 - Also self-triggering, cycles, nested invocations
- Conditions in **When** vs. as part of action

DBMS vs. Multiple Users

- Support for concurrent access necessary
 - Lower response time, users do not have to queue behind large jobs
- Make maximal use of CPU / disk resources for performance
 - With single user at a time, access to disk wastes huge amount of CPU cycles
 - Processing user requests concurrently increases throughput

Transactions

- Reliable unit of work against database
- **ACID** Properties
 - **Atomicity**: transactions are all-or-nothing
 - **Consistency**: transaction takes database from one consistent state to another
 - **Isolation**: transaction executes as if it was the only one in the system
 - **Durability**: once transaction is “committed”, results are persistent in the database

Transactions in SQL

Transaction T1: TRANSFER

```
BEGIN;  
UPDATE accounts  
SET balance = balance - 10  
WHERE name = 'Dmitriy';  
--  
UPDATE accounts  
SET balance = balance + 10  
WHERE name = 'Magnus';  
COMMIT;
```

Transaction T2: INTEREST

```
BEGIN;  
UPDATE accounts  
SET balance = balance * 1.01;  
COMMIT;
```

- **BEGIN** and **COMMIT** keywords delimit transaction
- **COMMIT** confirms that work should be made durable in the database

Transactions in SQL

Transaction T1: TRANSFER

```
BEGIN;  
UPDATE accounts  
SET balance = balance - 10  
WHERE name = 'Dmitriy';  
--  
UPDATE accounts  
SET balance = balance + 10  
WHERE name = 'Magnus';  
ROLLBACK;
```

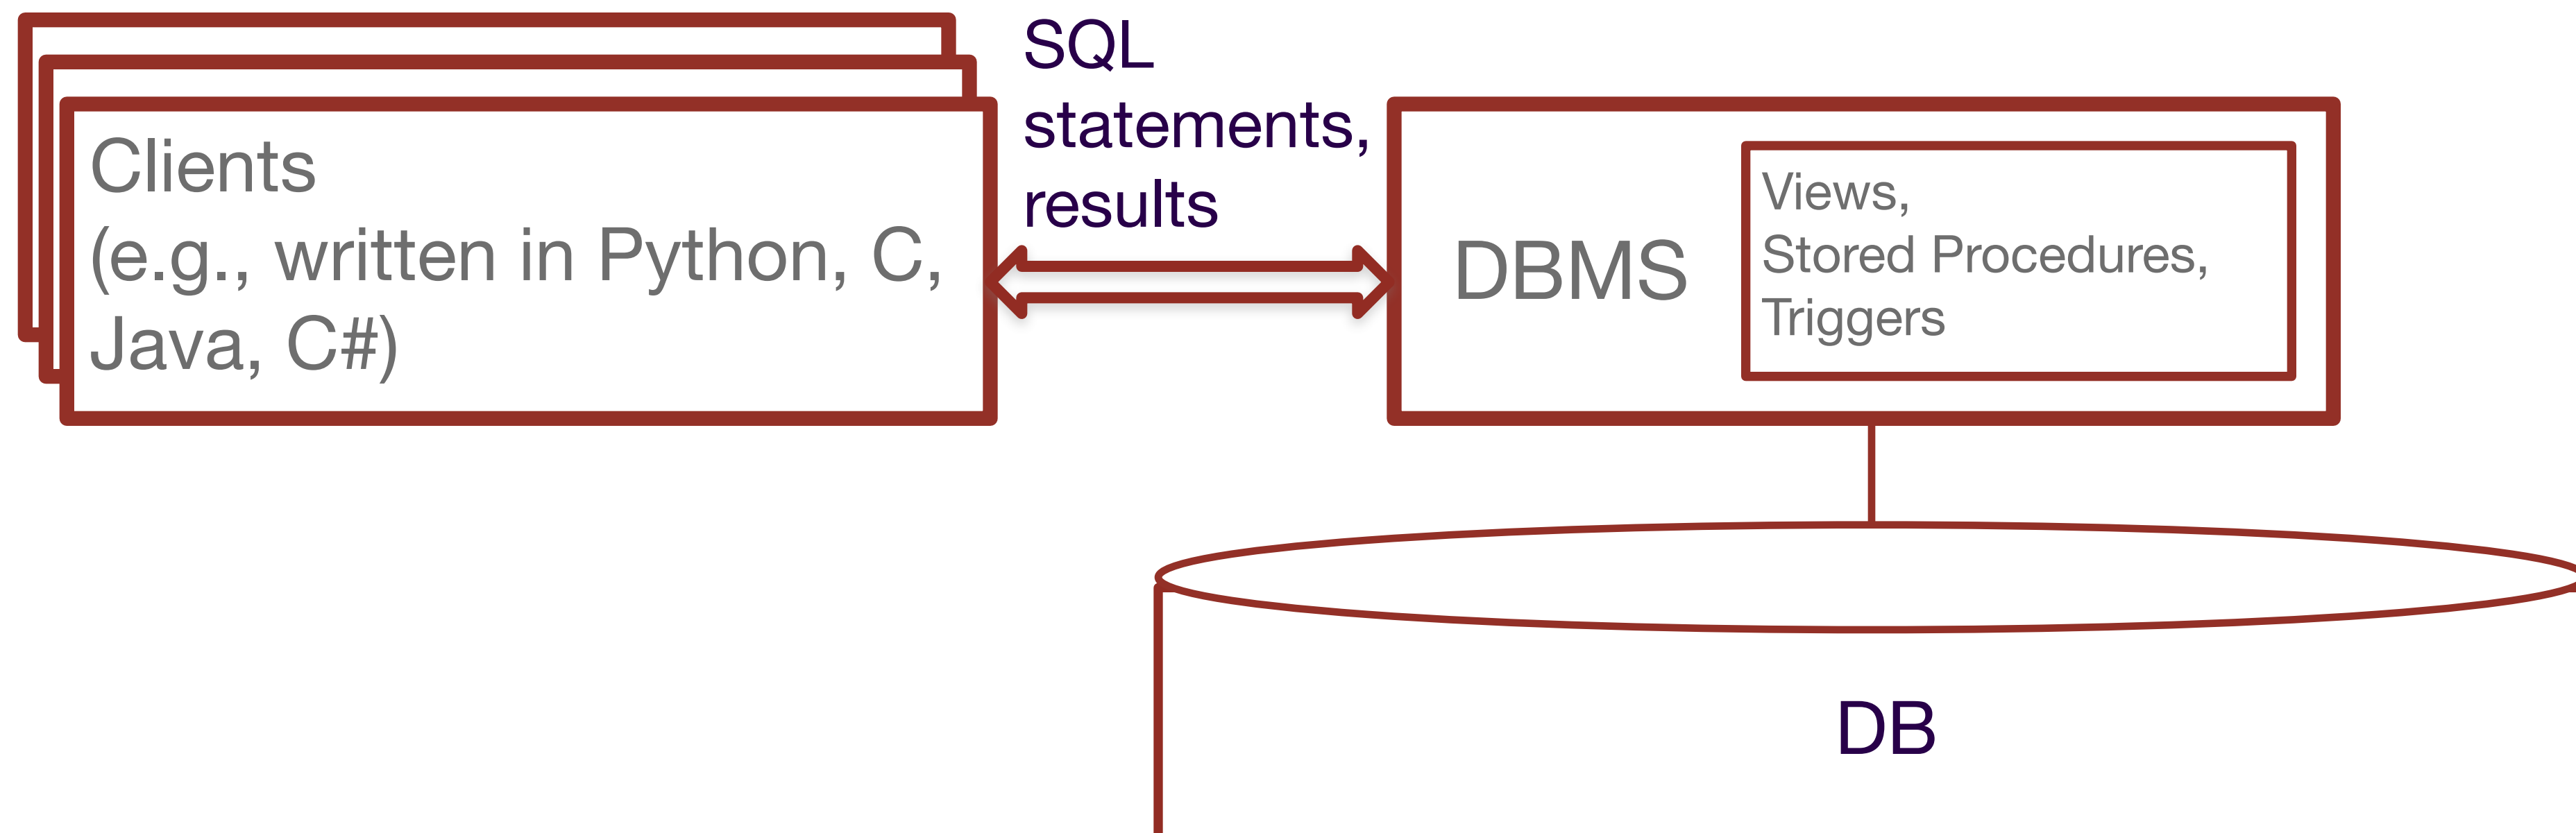
Transaction T2: INTEREST

```
BEGIN;  
UPDATE accounts  
SET balance = balance * 1.01;  
ROLLBACK;
```

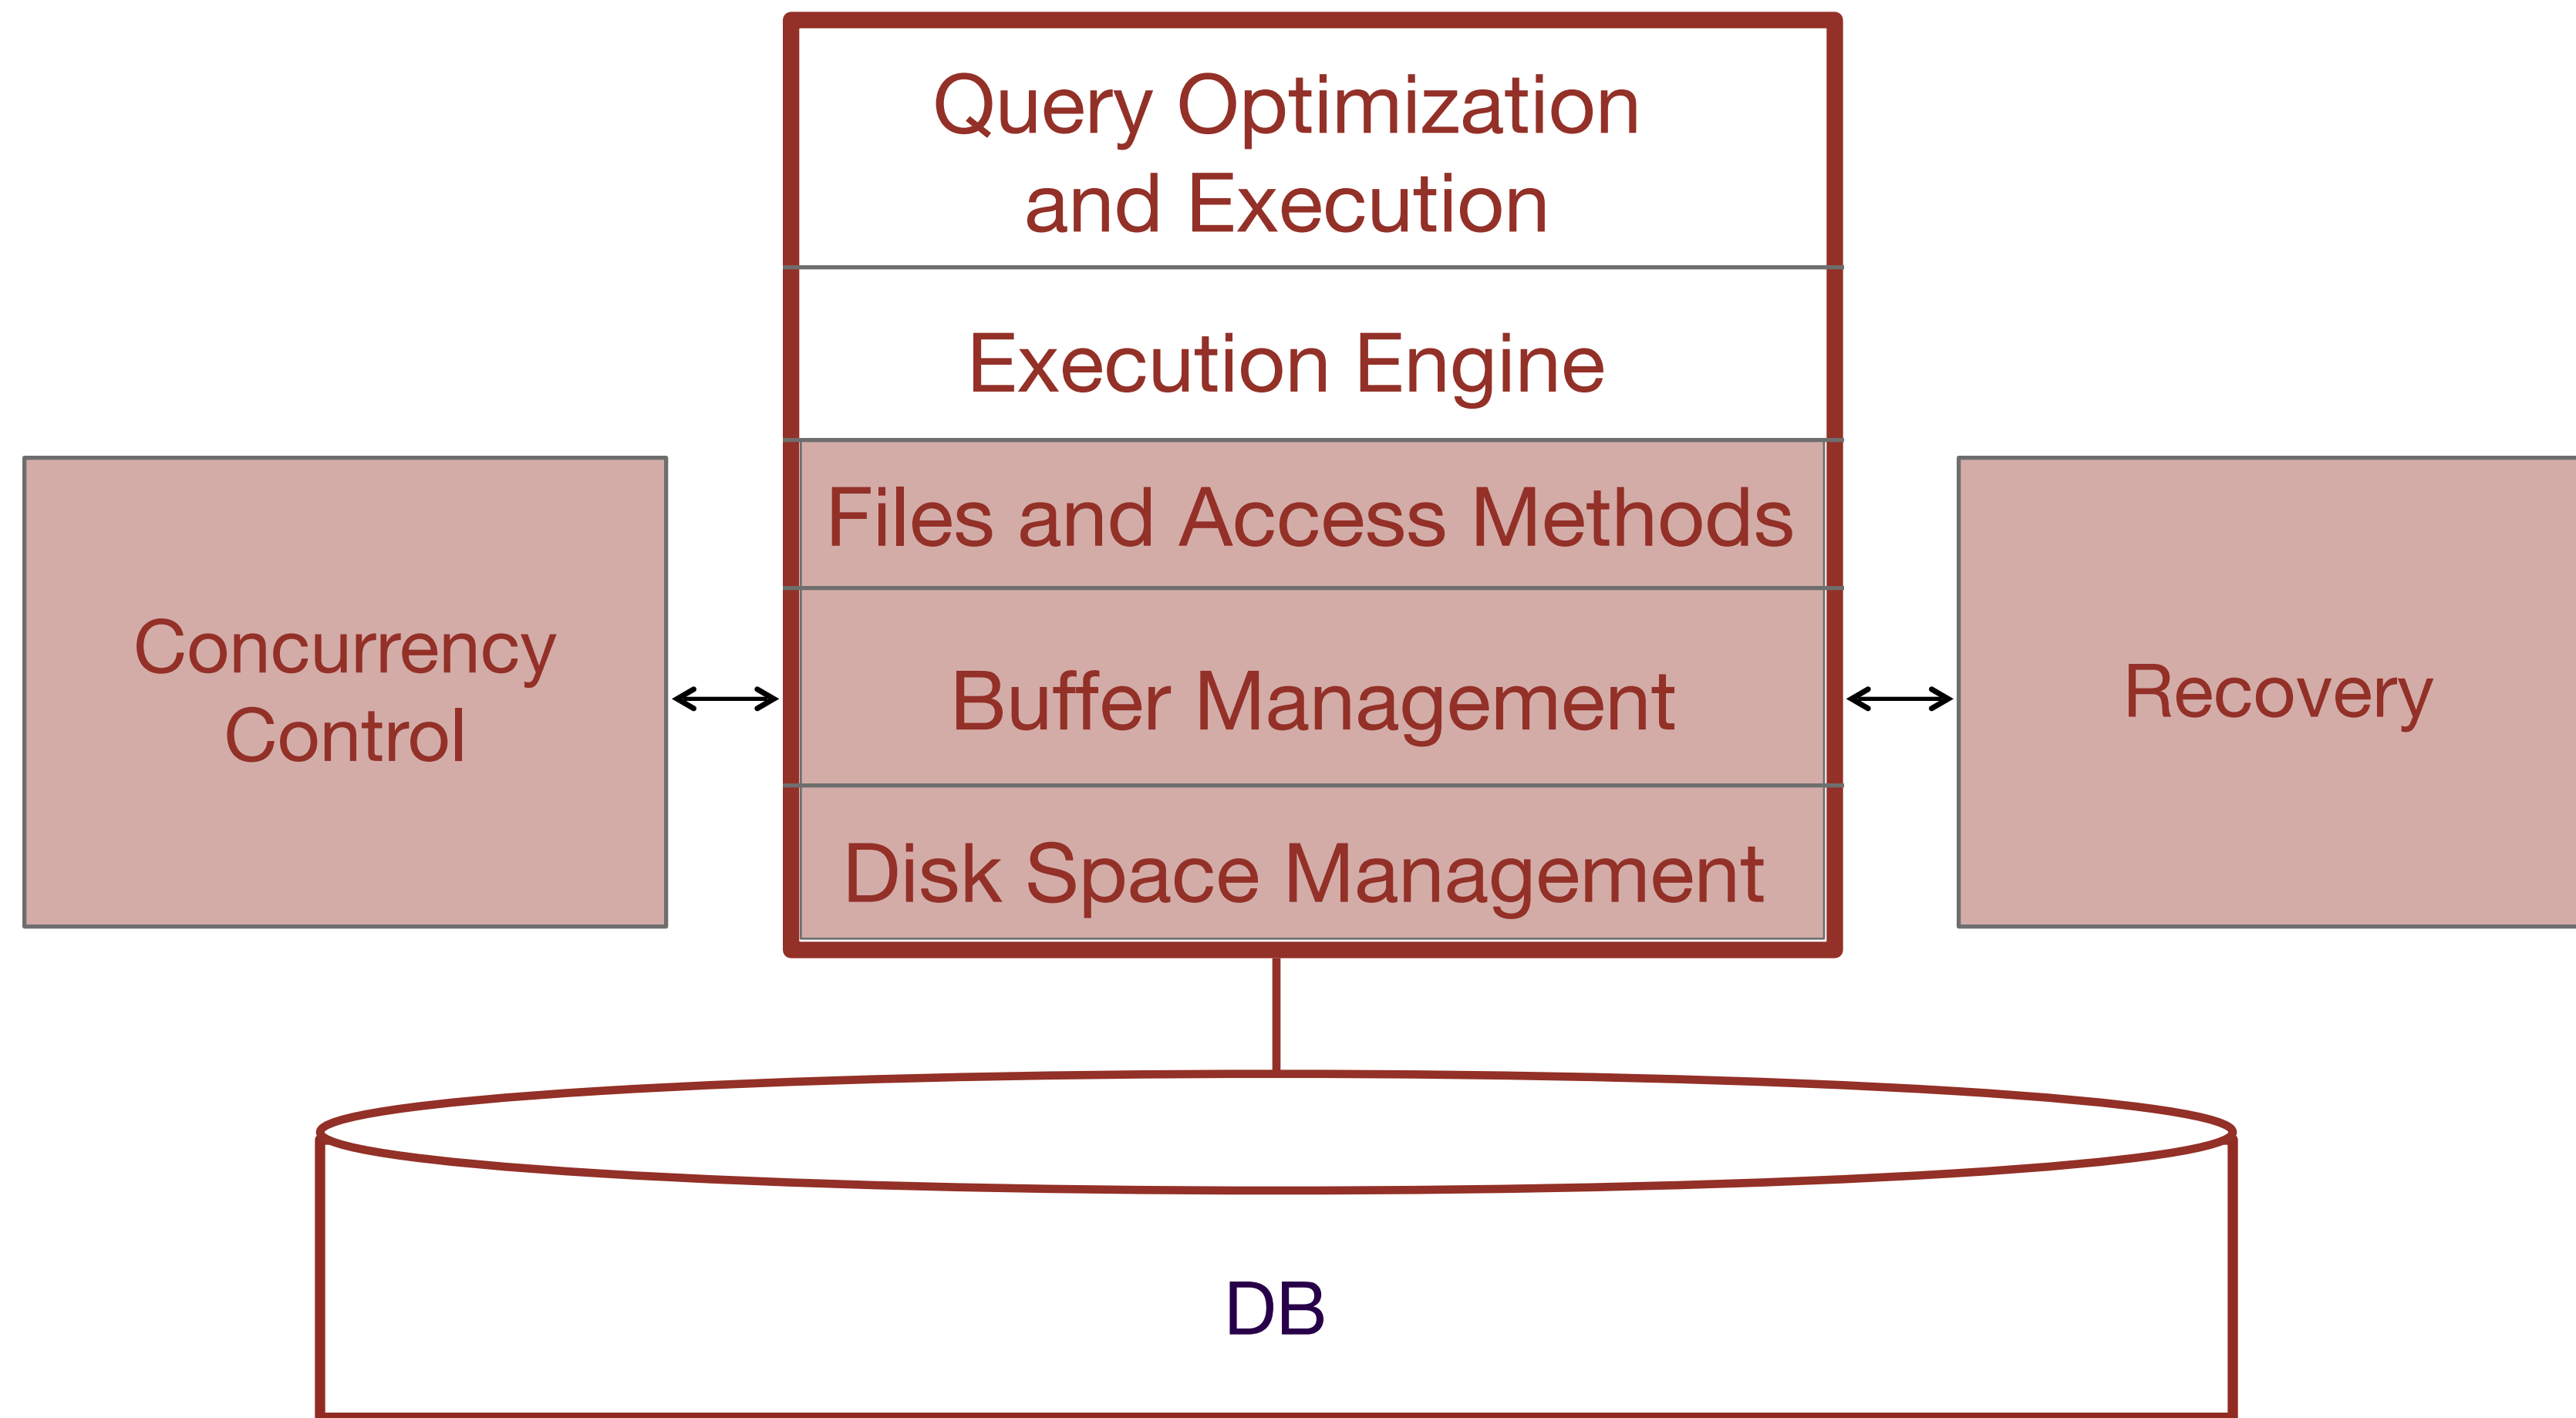
- **ROLLBACK** implies that effects of transactions should be undone
- In this example, database back to original state

Where does the code run?

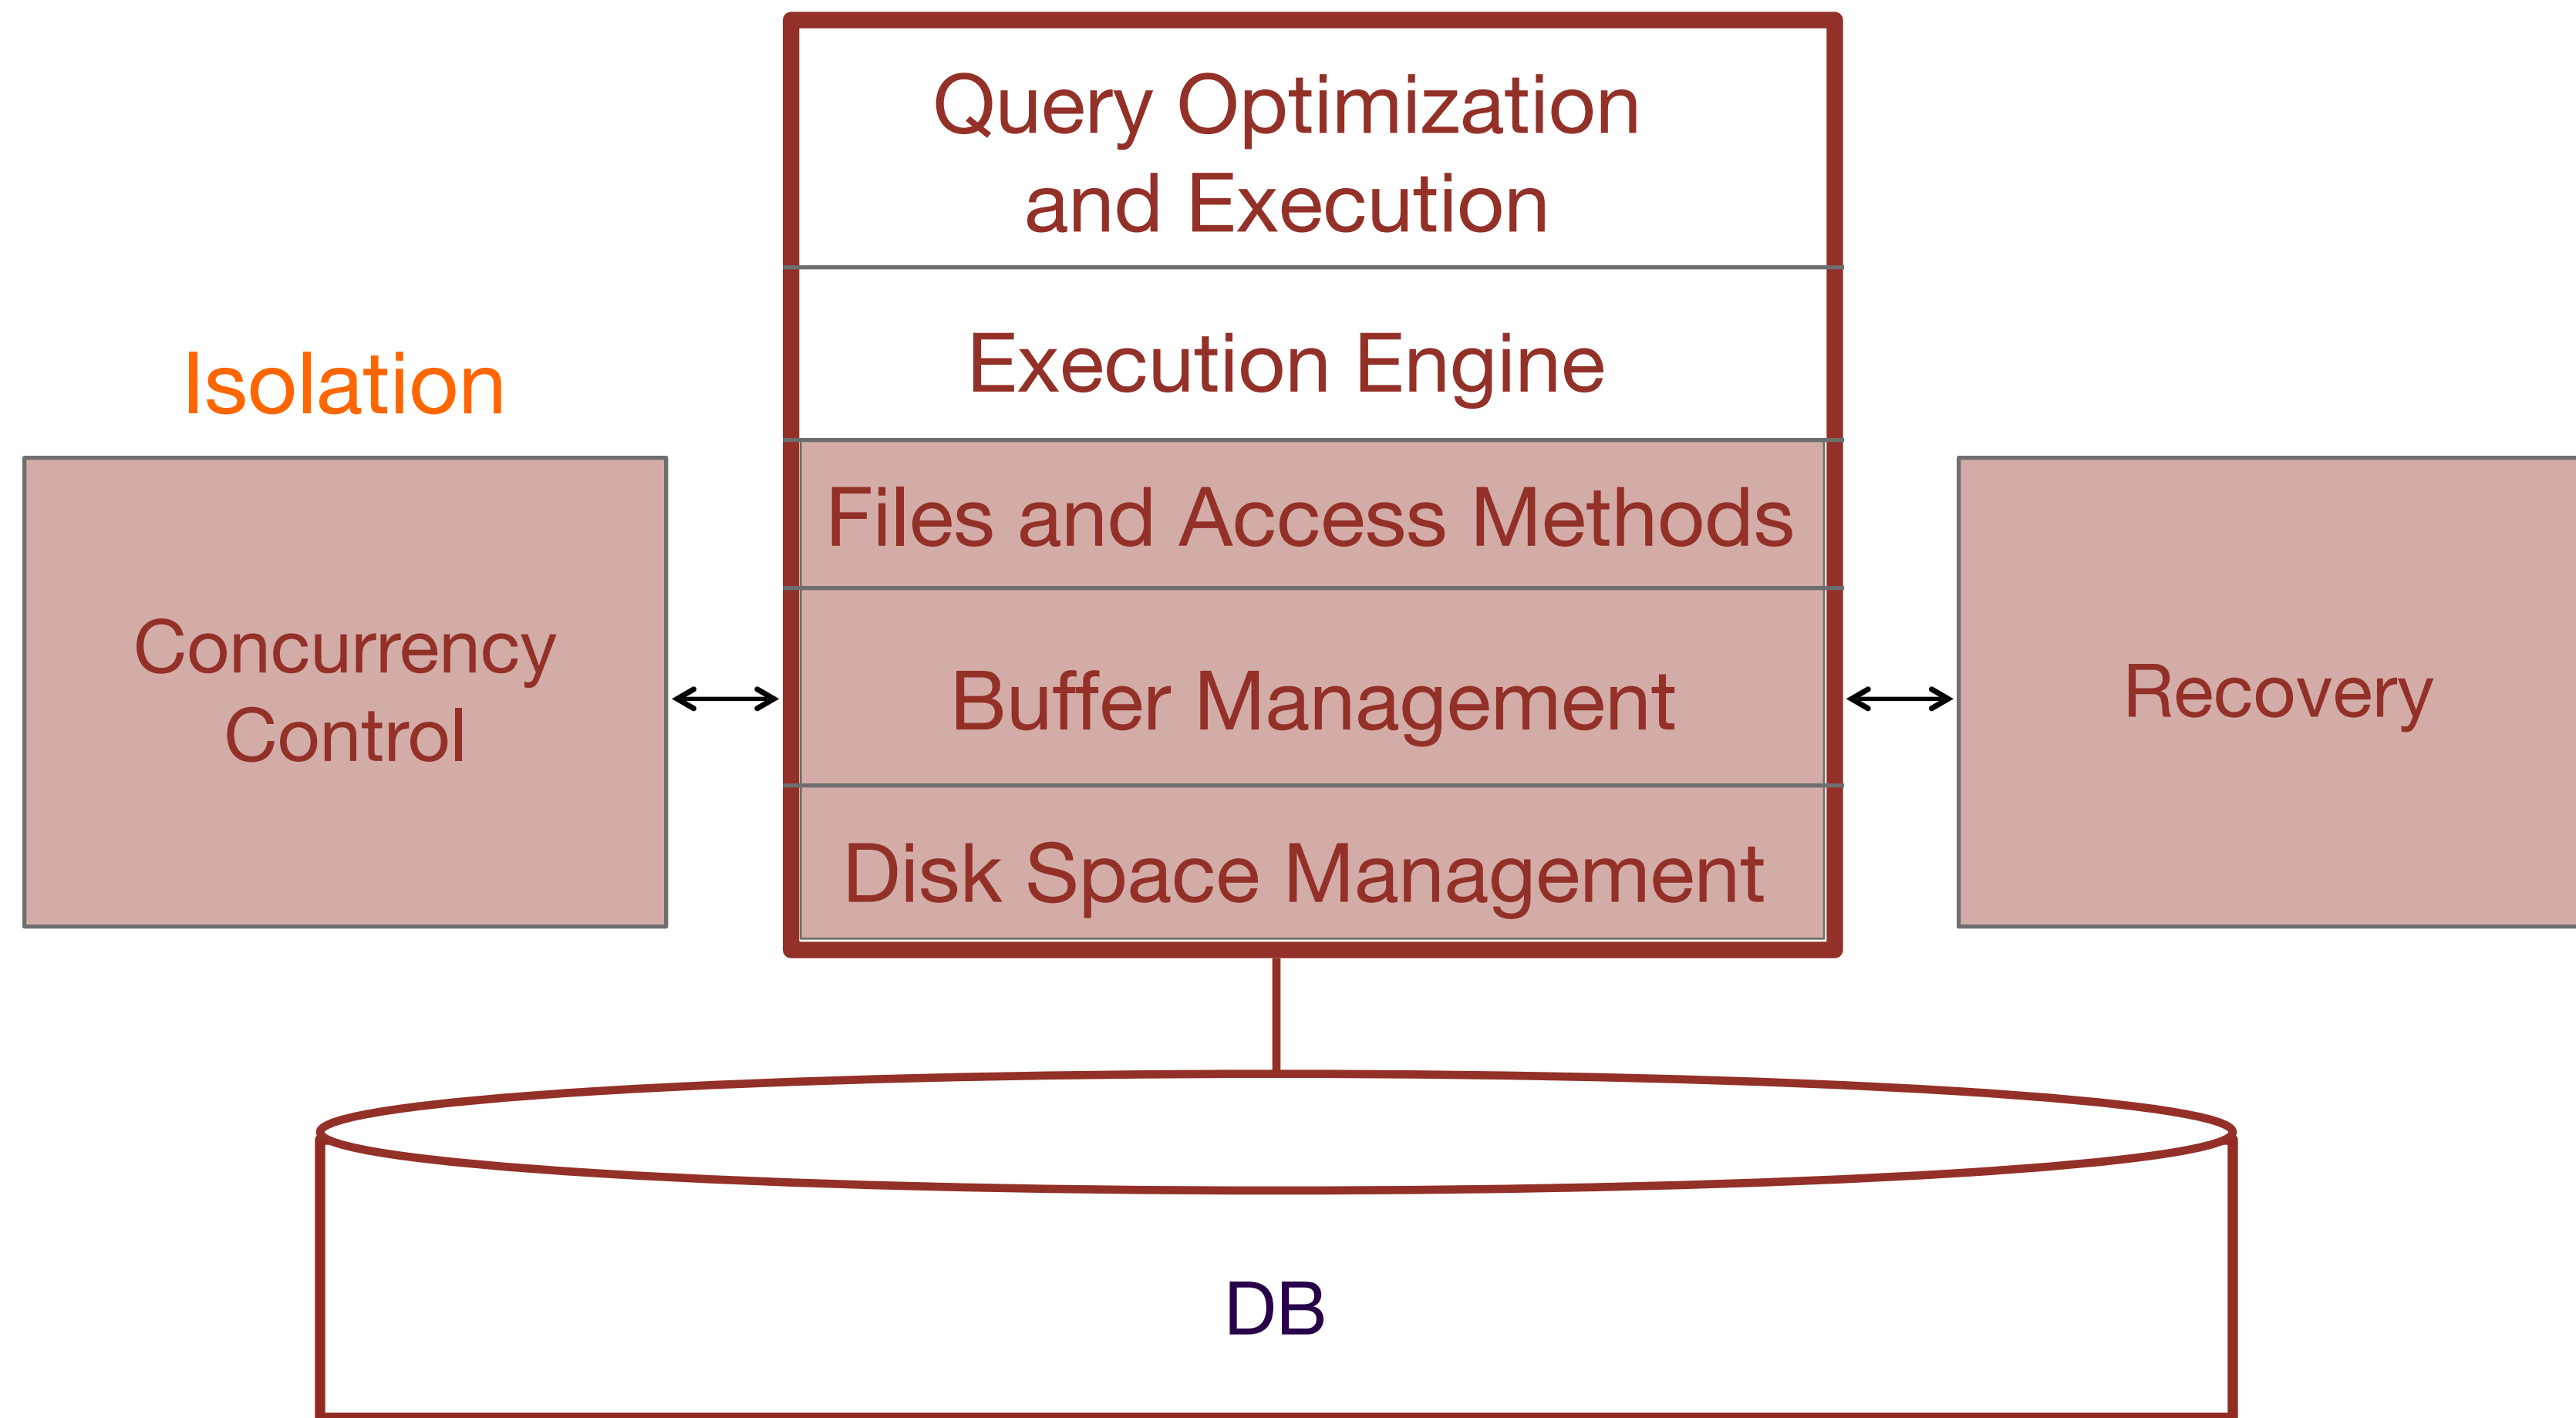
- Previous example could be input directly in psql
- Transaction often run from client program through a **driver** or **database adapter**
 - e.g., Psycopg (<http://initd.org/psycopg/>)



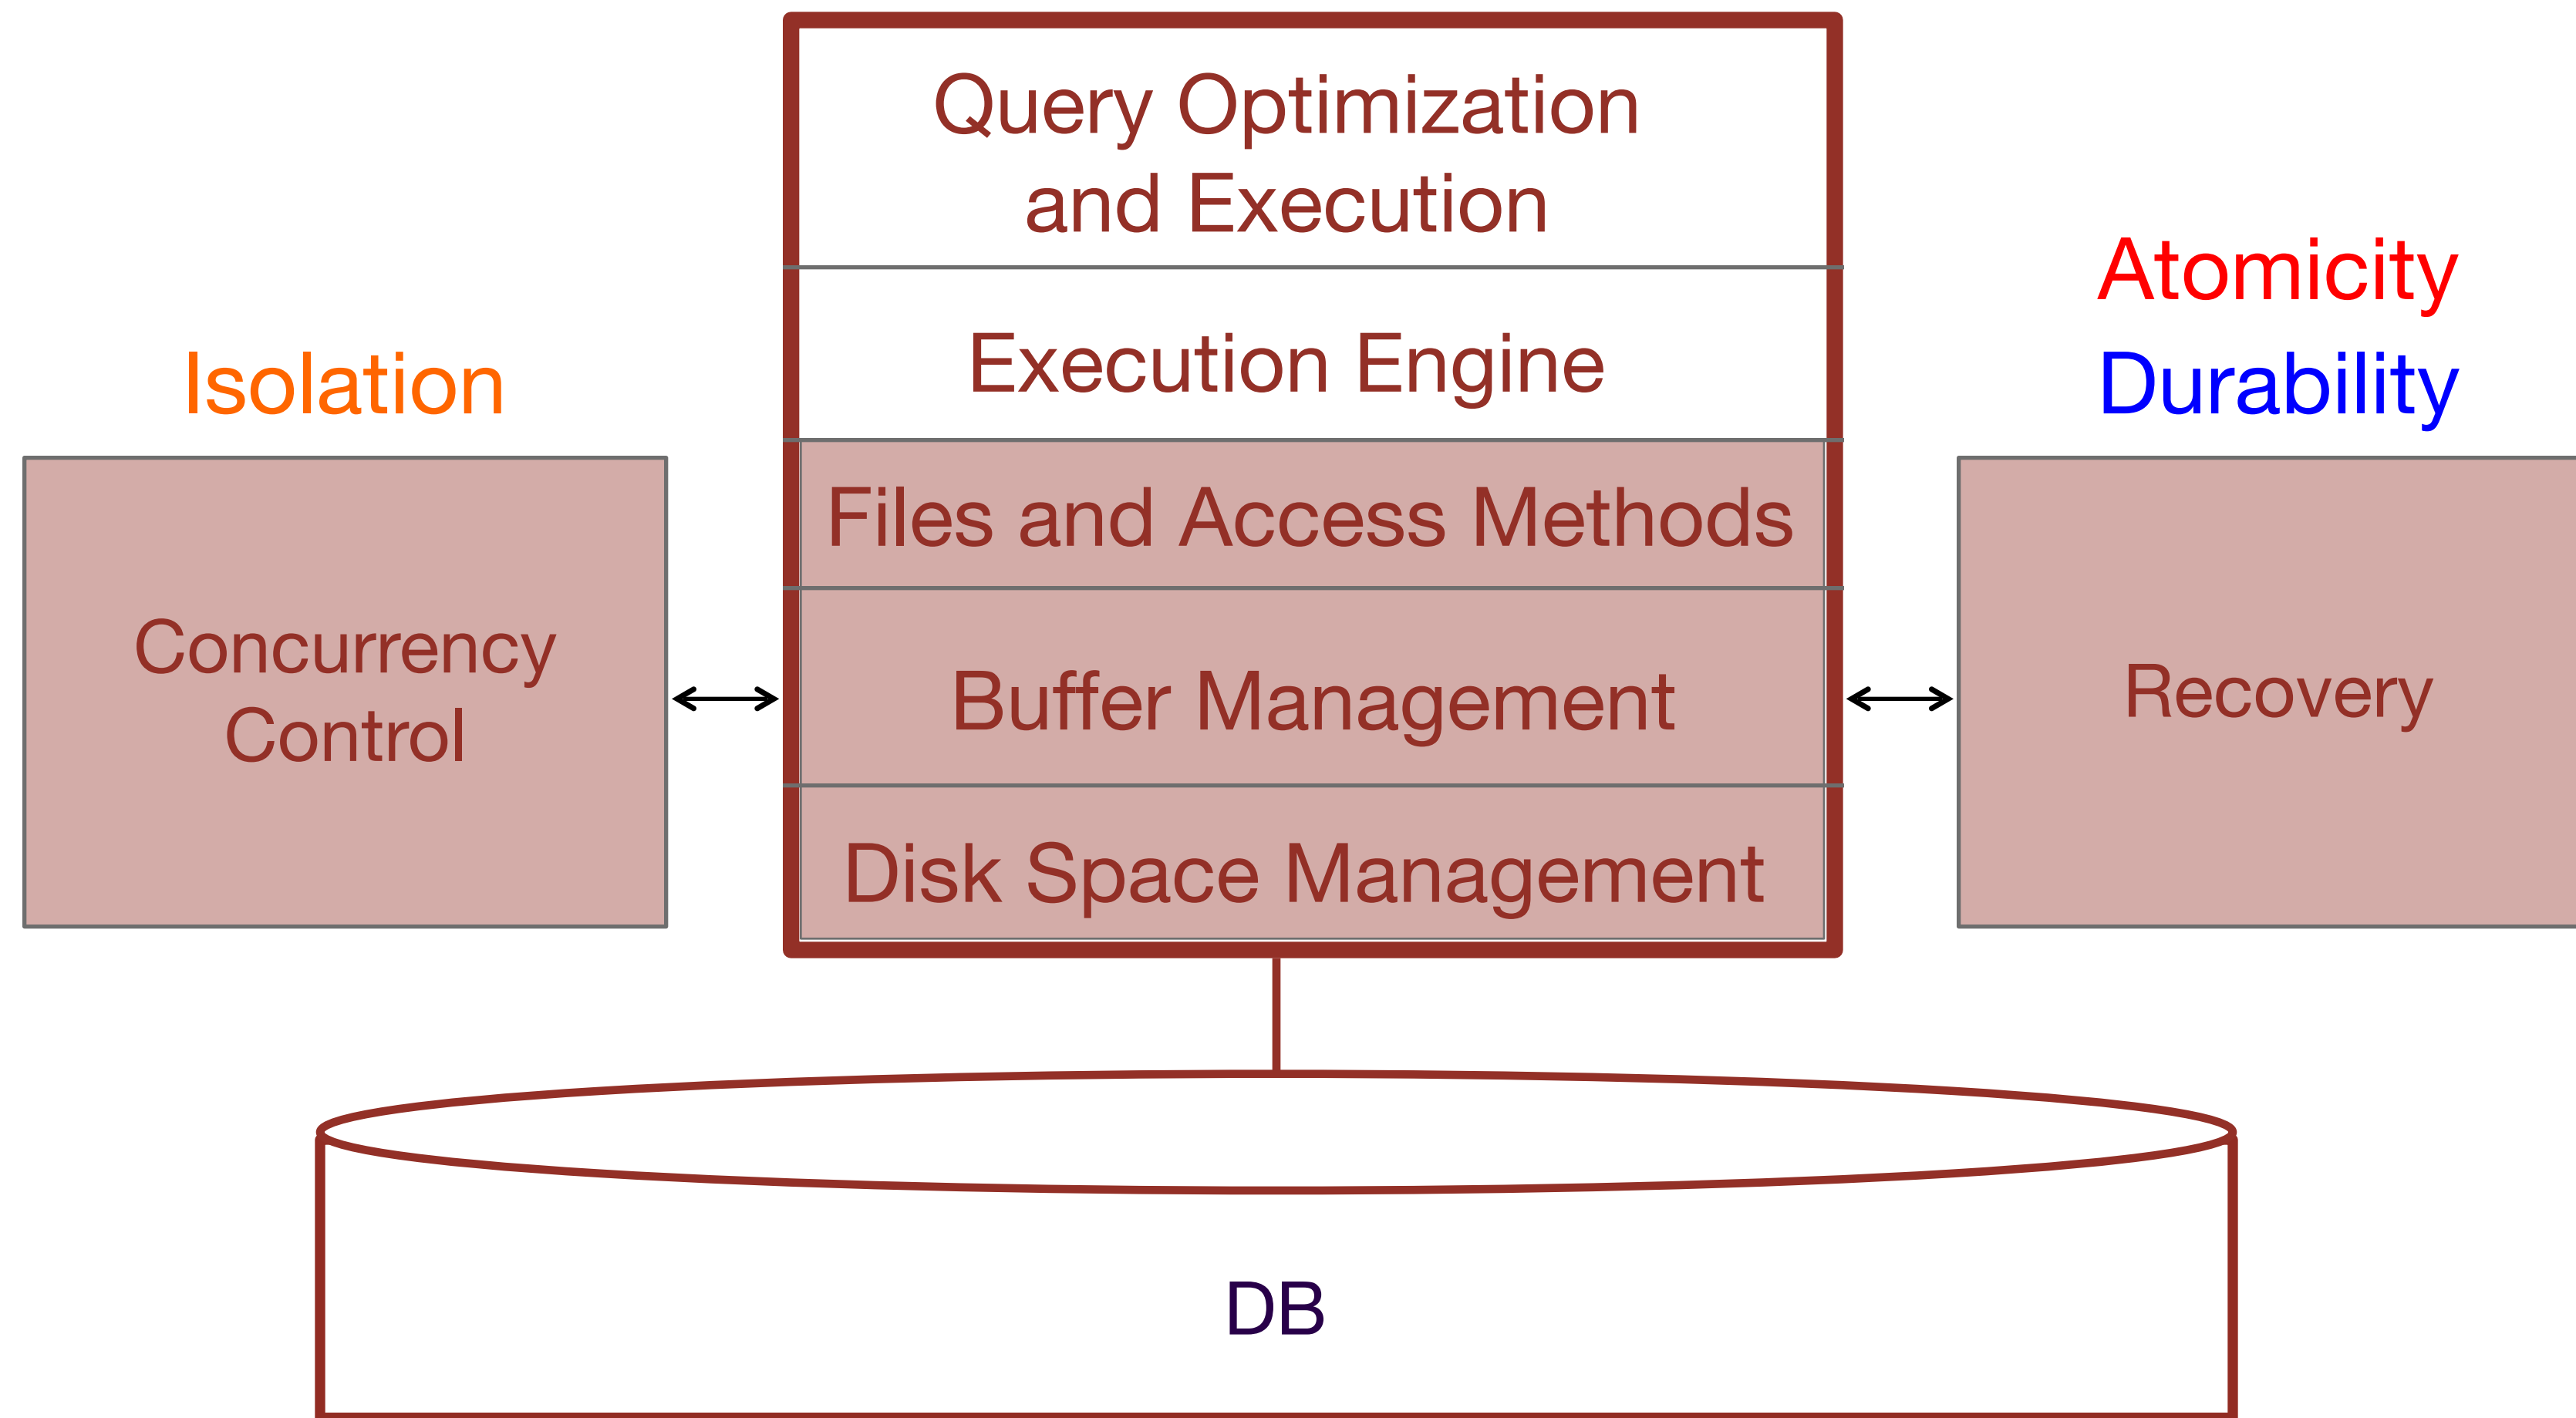
Classic Architecture of a DBMS



Classic Architecture of a DBMS



Classic Architecture of a DBMS



Discussion: Transactions in SQL

- Take script `txn_example.sql` from Absalon and open two `psql` terminals
- Run first this code fragment:

```
BEGIN;
UPDATE accounts
SET balance = balance * 1.01;
-- do not commit yet
```
- In a separate terminal, then run:

```
BEGIN;
UPDATE accounts
SET balance = balance + 10
WHERE name = 'Dmitriy';
...
```
- Discussion: If you run the scripts above in PostgreSQL via two `psql` terminals, what are the outcomes? Why?
- Can you now commit or rollback and see what happens to the balances?

Questions so far?

A Query

```
SELECT DISTINCT X.location, X.date, count(Y.date)
FROM covid X JOIN covid Y
ON X.location = Y.location AND X.new_cases > Y.new_cases
GROUP BY X.location, X.date
```

A Query

```
SELECT DISTINCT X.location, X.date, count(Y.date)
FROM covid X JOIN covid Y
ON X.location = Y.location AND X.new_cases > Y.new_cases
GROUP BY X.location, X.date
```

Wait for 25 seconds!

A Query

```
SELECT DISTINCT X.location, X.date, count(Y.date)
FROM covid X JOIN covid Y
ON X.location = Y.location AND X.new_cases > Y.new_cases
GROUP BY X.location, X.date
```

Wait for 25 seconds!

```
...
Zimbabwe | 2021-05-16 | 99
Zimbabwe | 2021-05-17 | 125
Zimbabwe | 2021-05-18 | 191
Zimbabwe | 2021-05-19 | 160
Zimbabwe | 2021-05-20 | 191
Zimbabwe | 2021-05-21 | 222
Zimbabwe | 2021-05-22 | 147
(72976 rows)
```

A(nother?) Query

```
CREATE INDEX covid_i ON covid(location, new_cases, date);  
  
SELECT DISTINCT X.location, X.date, count(Y.date)  
FROM covid X JOIN covid Y  
ON X.location = Y.location AND X.new_cases > Y.new_cases  
GROUP BY X.location, X.date
```

A(nother?) Query

```
CREATE INDEX covid1 ON covid(location, new_cases, date);

SELECT DISTINCT X.location, X.date, count(Y.date)
FROM covid X JOIN covid Y
ON X.location = Y.location AND X.new_cases > Y.new_cases
GROUP BY X.location, X.date
```

```
...
Zimbabwe | 2021-05-16 | 99
Zimbabwe | 2021-05-17 | 125
Zimbabwe | 2021-05-18 | 191
Zimbabwe | 2021-05-19 | 160
Zimbabwe | 2021-05-20 | 191
Zimbabwe | 2021-05-21 | 222
Zimbabwe | 2021-05-22 | 147
(72976 rows)
```

A(nother?) Query

```
CREATE INDEX covidi ON covid(location, new_cases, date);  
  
SELECT DISTINCT X.location, X.date, count(Y.date)  
FROM covid X JOIN covid Y  
ON X.location = Y.location AND X.new_cases > Y.new_cases  
GROUP BY X.location, X.date
```

Wait for 7.5 seconds!

```
...  
Zimbabwe | 2021-05-16 | 99  
Zimbabwe | 2021-05-17 | 125  
Zimbabwe | 2021-05-18 | 191  
Zimbabwe | 2021-05-19 | 160  
Zimbabwe | 2021-05-20 | 191  
Zimbabwe | 2021-05-21 | 222  
Zimbabwe | 2021-05-22 | 147  
(72976 rows)
```

A(nother?) Query

```
CREATE INDEX covid1 ON covid(location, new_cases, date);
```

0.2 s

```
SELECT DISTINCT X.location, X.date, count(Y.date)
```

```
FROM covid X JOIN covid Y
```

```
ON X.location = Y.location AND X.new_cases > Y.new_cases
```

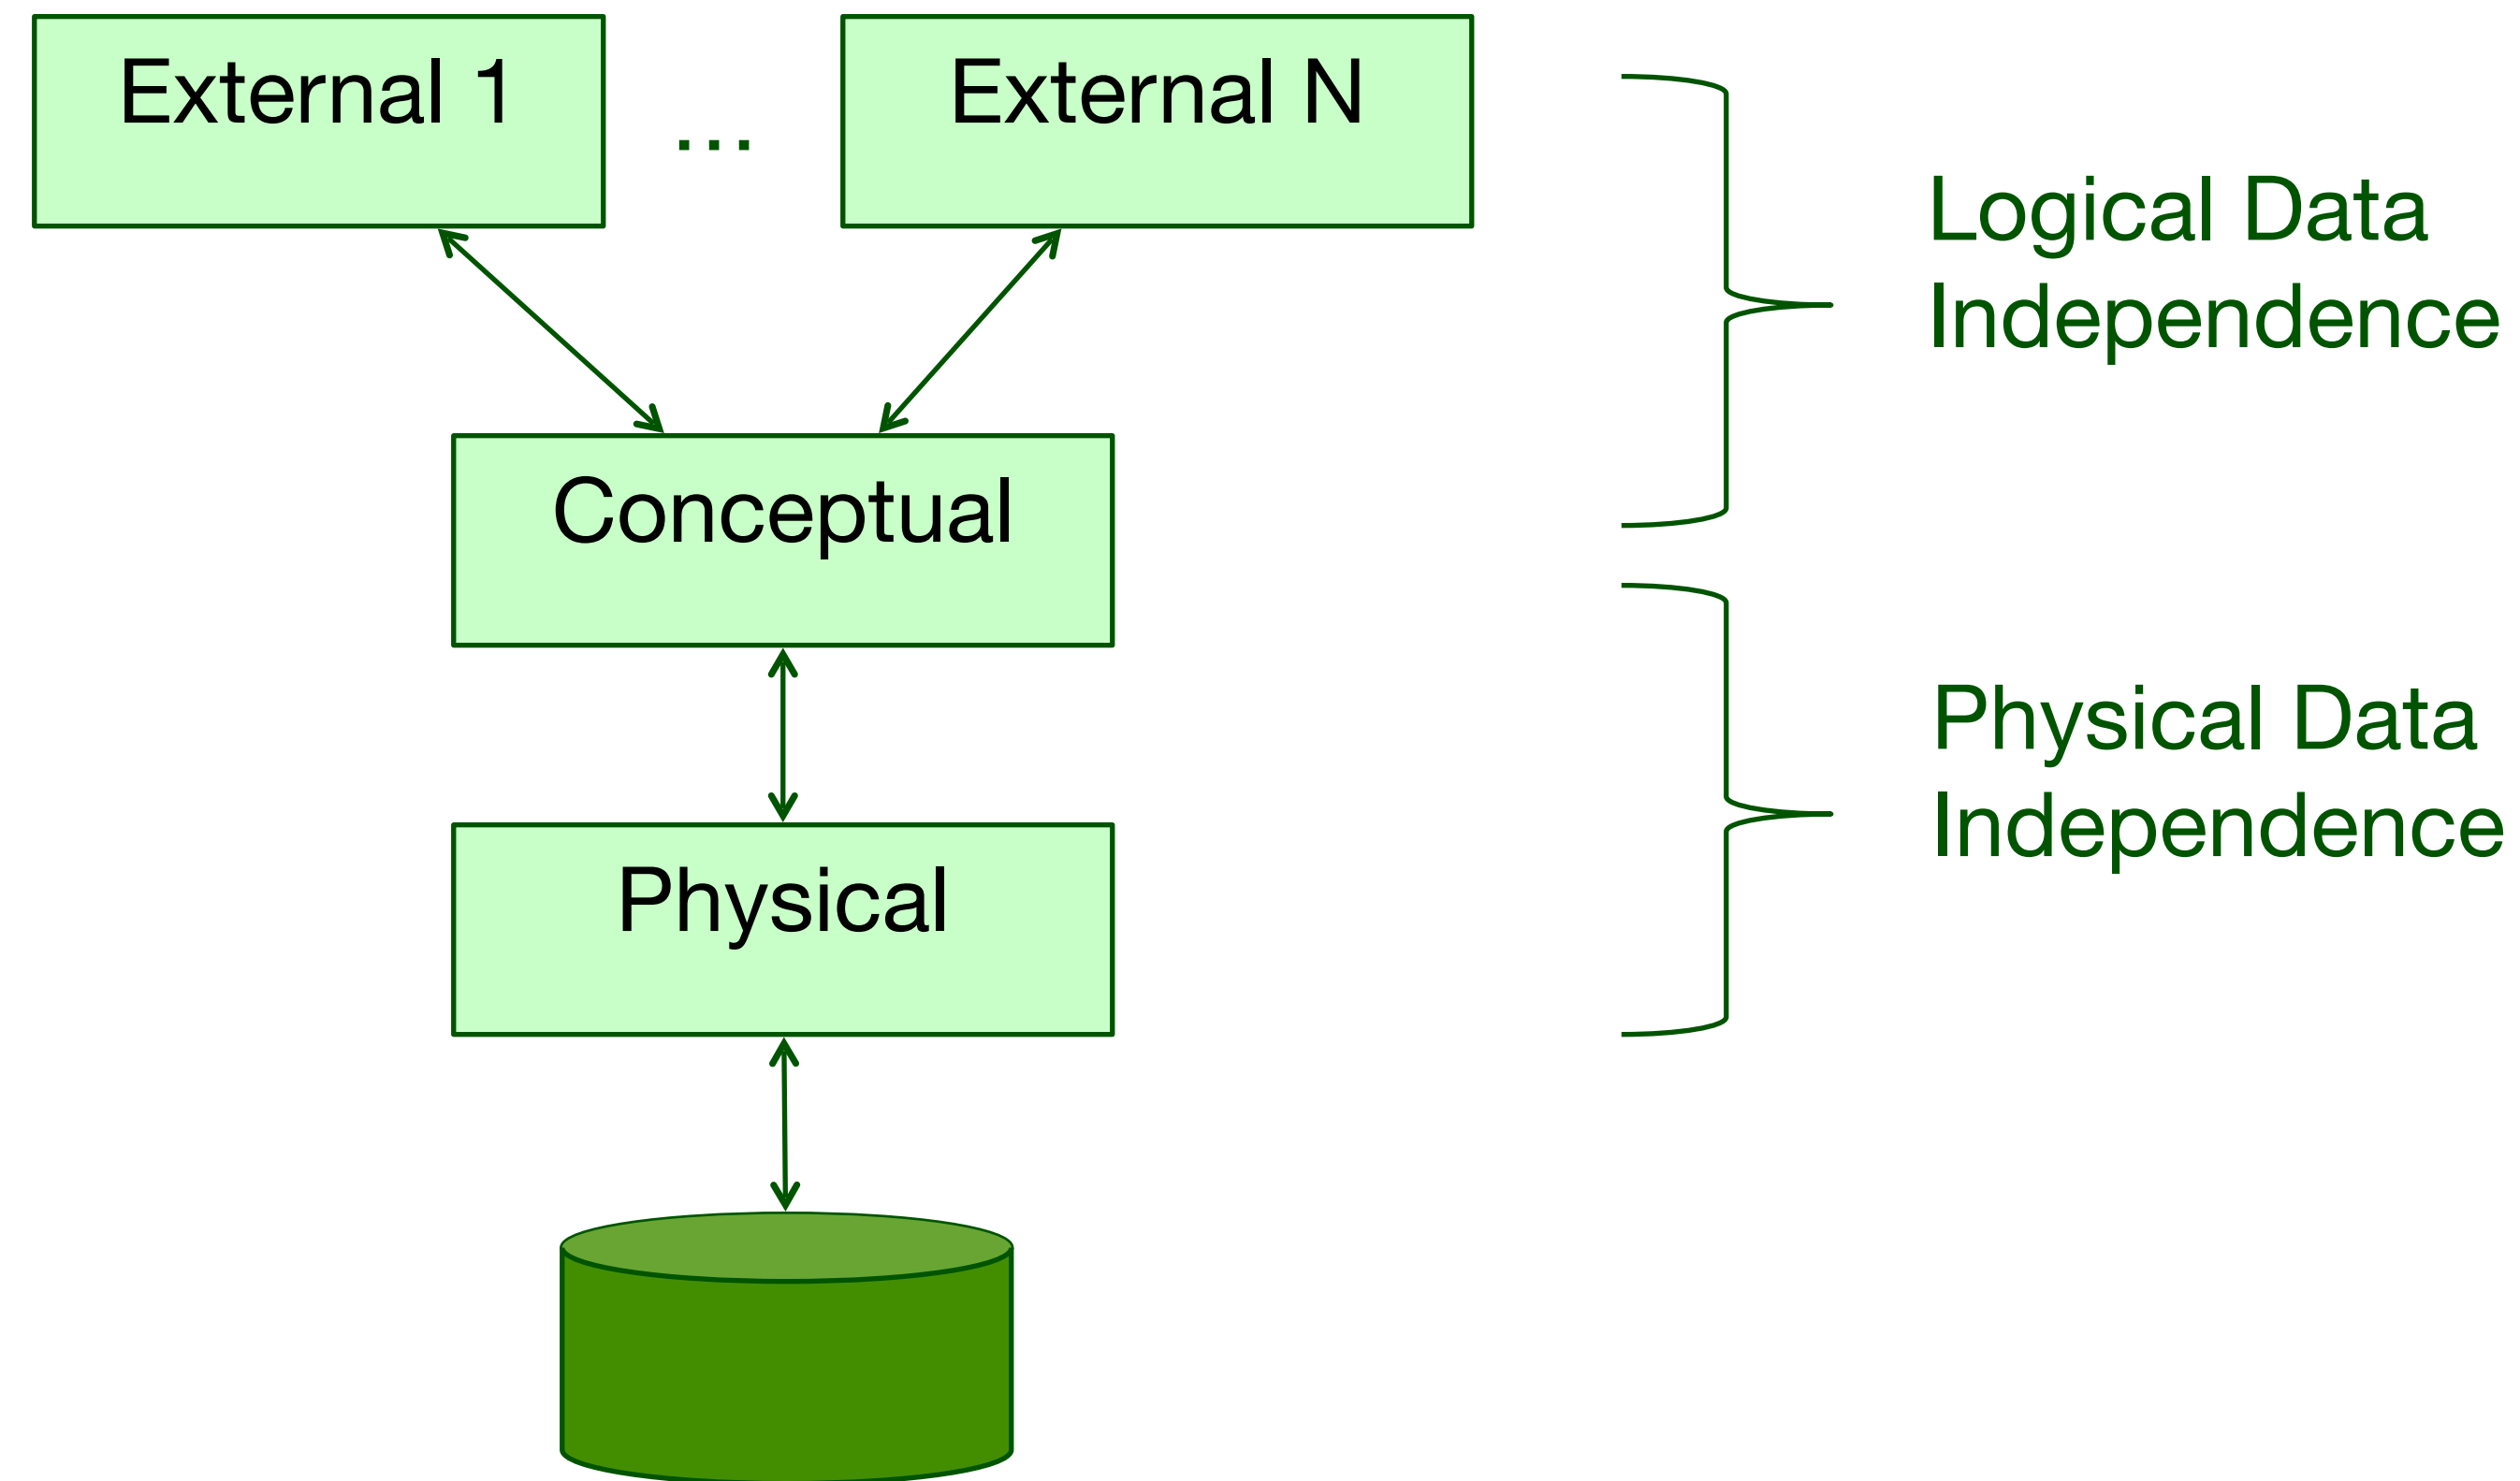
```
GROUP BY X.location, X.date
```

7.3 s

Wait for 7.5 seconds!

```
...
Zimbabwe | 2021-05-16 | 99
Zimbabwe | 2021-05-17 | 125
Zimbabwe | 2021-05-18 | 191
Zimbabwe | 2021-05-19 | 160
Zimbabwe | 2021-05-20 | 191
Zimbabwe | 2021-05-21 | 222
Zimbabwe | 2021-05-22 | 147
(72976 rows)
```


Physical Data Independence



How can we physically represent relations?

Employees

<u>ssn</u>	name	lot
0983763423	John	10
9384392483	Jane	10
3743923483	Jill	20

Departments

<u>did</u>	dname	budget
101	Sales	10K
105	Purchasing	20K
108	Databases	1000K

Works_In

<u>ssn</u>	<u>did</u>	since
0983763423	101	1 Jan 2003
0983763423	108	2 Jan 2003
9384392483	108	1 Jun 2002


Files of Records

- DBMS software operates on records, and files of records.
- File: A collection of pages (or blocks), each containing a collection of records.
Must support:
 - insert/delete/modify record
 - read a particular record (specified using record id)
 - scan all records (possibly with some conditions on the records to be retrieved)

Unordered (Heap) Files

- Simplest file structure contains records in no particular order.
- As file grows and shrinks, disk pages are allocated and de-allocated.
- To support record level operations, we must:
 - keep track of the **pages** in a file
 - keep track of **free space** on pages
 - keep track of the **records** on a page
- There are many alternatives for keeping track of this.

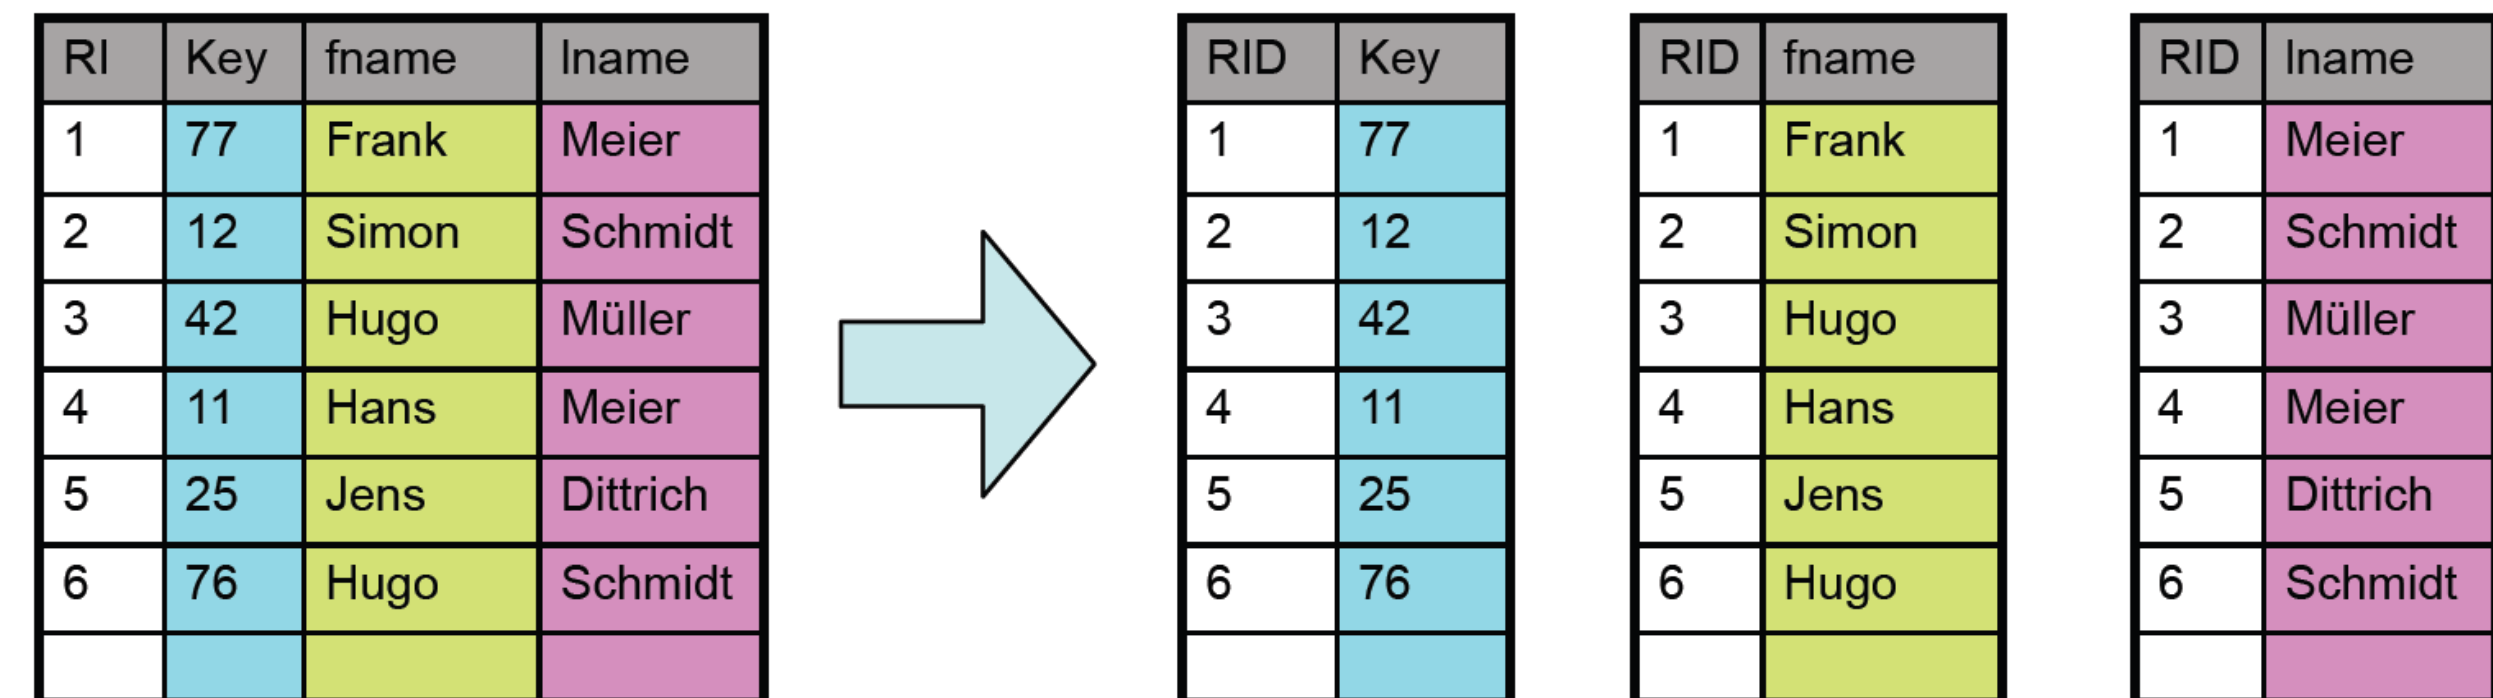
Heap File Using a Page Directory



- The entry for a page can include the number of free bytes on the page.
- The directory is a collection of pages; e.g., implemented via linked list
 - Much smaller than linked list of all heap file pages!

Alternative Storage Organizations

- Classically, heap file follow a **row-store** organization
 - n-ary Storage Model (NSM)
 - “array of tuples”
- Alternatives:
 - Column-store**: Decomposition Storage Model (DSM)
 - “tuple of arrays”
 - Hybrid**: Partition Attributes Across (PAX)
 - “array of tuples of arrays”



How to access individual records when you do not have the record ID?

```
SELECT *  
FROM   departments  
WHERE  did >= 100  
       AND did <= 200
```

```
SELECT *  
FROM   employees  
WHERE  name = 'John'
```

- **Heap file: scan the whole file!**
- Is there a way to map attribute values directly to records in heap file?

How to access individual records when you do not have the record ID?

```
SELECT *  
FROM   departments  
WHERE  did >= 100  
       AND did <= 200
```


```
SELECT *  
FROM   employees  
WHERE  name = 'John'
```

- **Heap file: scan the whole file!**
- Is there a way to map attribute values directly to records in heap file?

Indices are data structures to do so!

Range Searches


- “Find all students with $\text{gpa} > 3.0$ ”
 - If data **entries** are sorted, do binary search to find first such student, then scan to find others.
- Problem?



Data (Entries) File

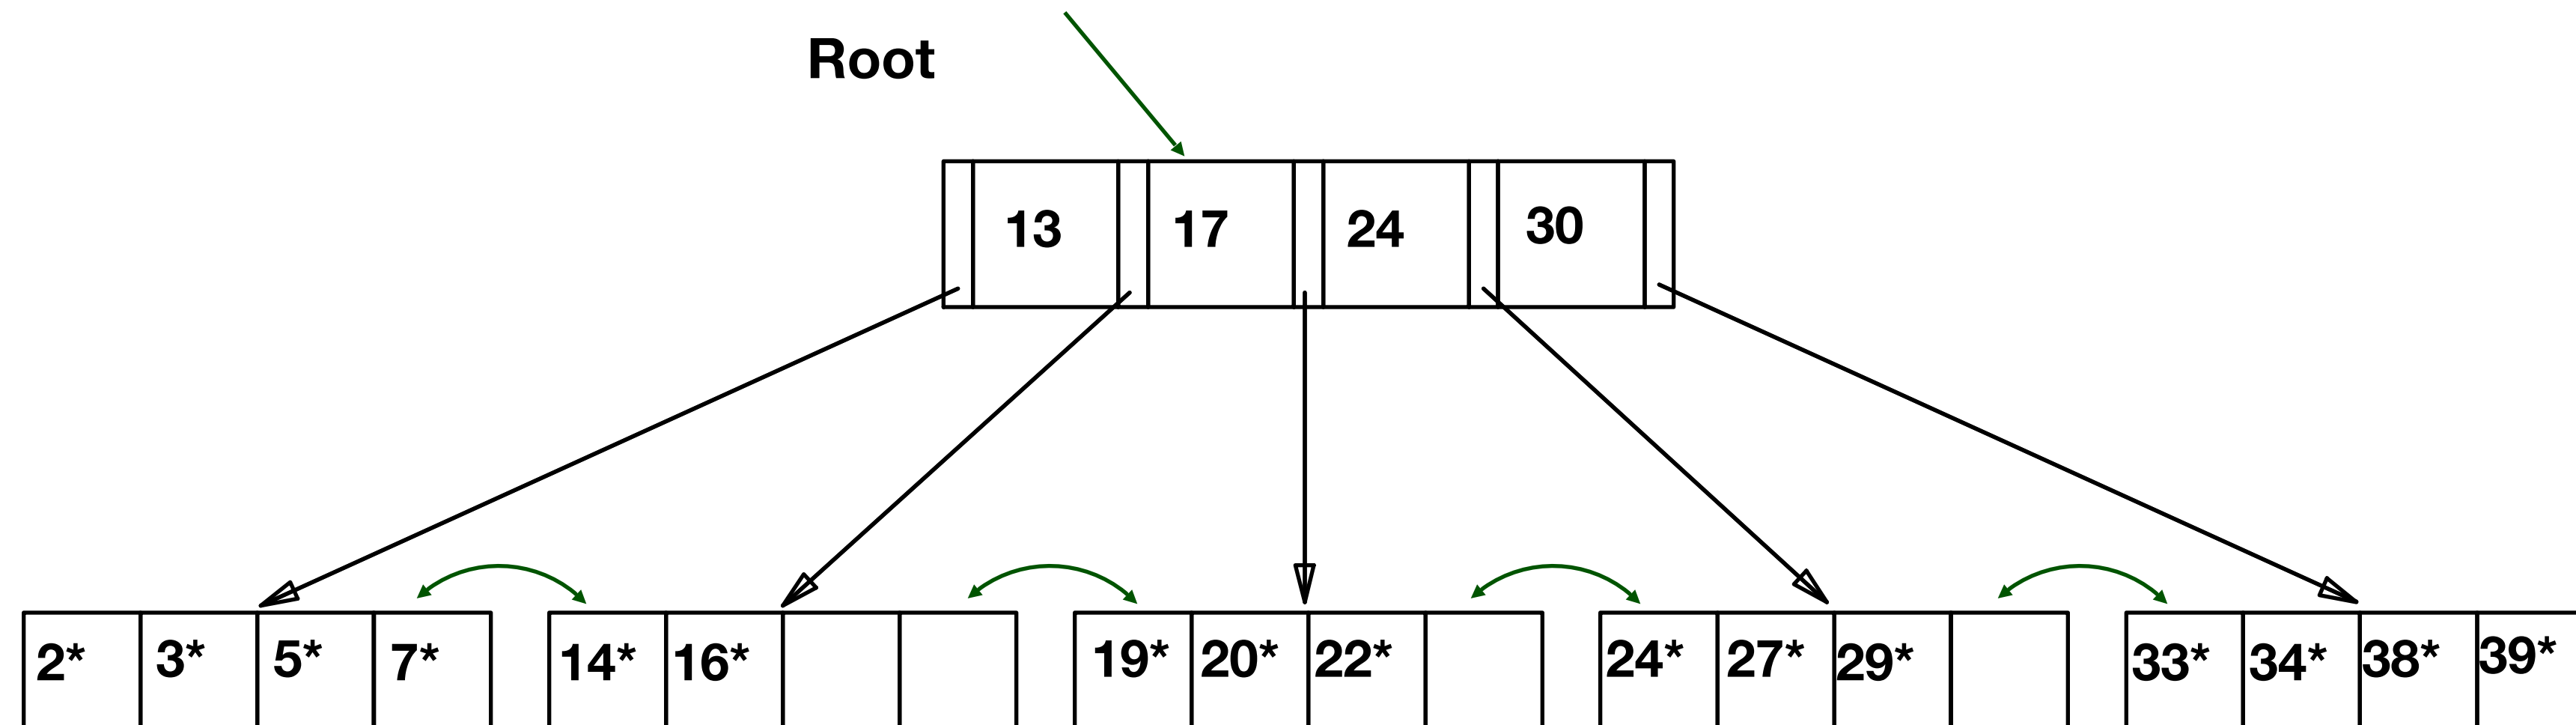
B+ Tree: The Most Widely Used Index

- Insert/delete at $\log_{\text{fanout}}\left(\frac{\text{number of entries}}{\text{leaf capacity}}\right)$ cost; keep tree *height-balanced*.
- Minimum 50% occupancy (except for root). Each node contains $d \leq m \leq 2d$ entries. The parameter d is called the *order* of the tree.
- Supports equality and range-searches efficiently.



Example B+ Tree

- Search begins at root, and key comparisons direct it to a leaf (as in ISAM).
- Search for 5*, 15*, all data entries $\geq 24^*$...




Based on the search for 15, we know it is not in the tree!*


B+ Trees in Practice

- Typical order: 100. Typical fill-factor: 67%.
 - average fanout = 133
- Typical capacities:
 - Height 4: $133^4 = 312,900,721$ records
 - Height 3: $133^3 = 2,352,637$ records
- Can often hold top levels in main memory:
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 MBytes


Inserting 23* ...



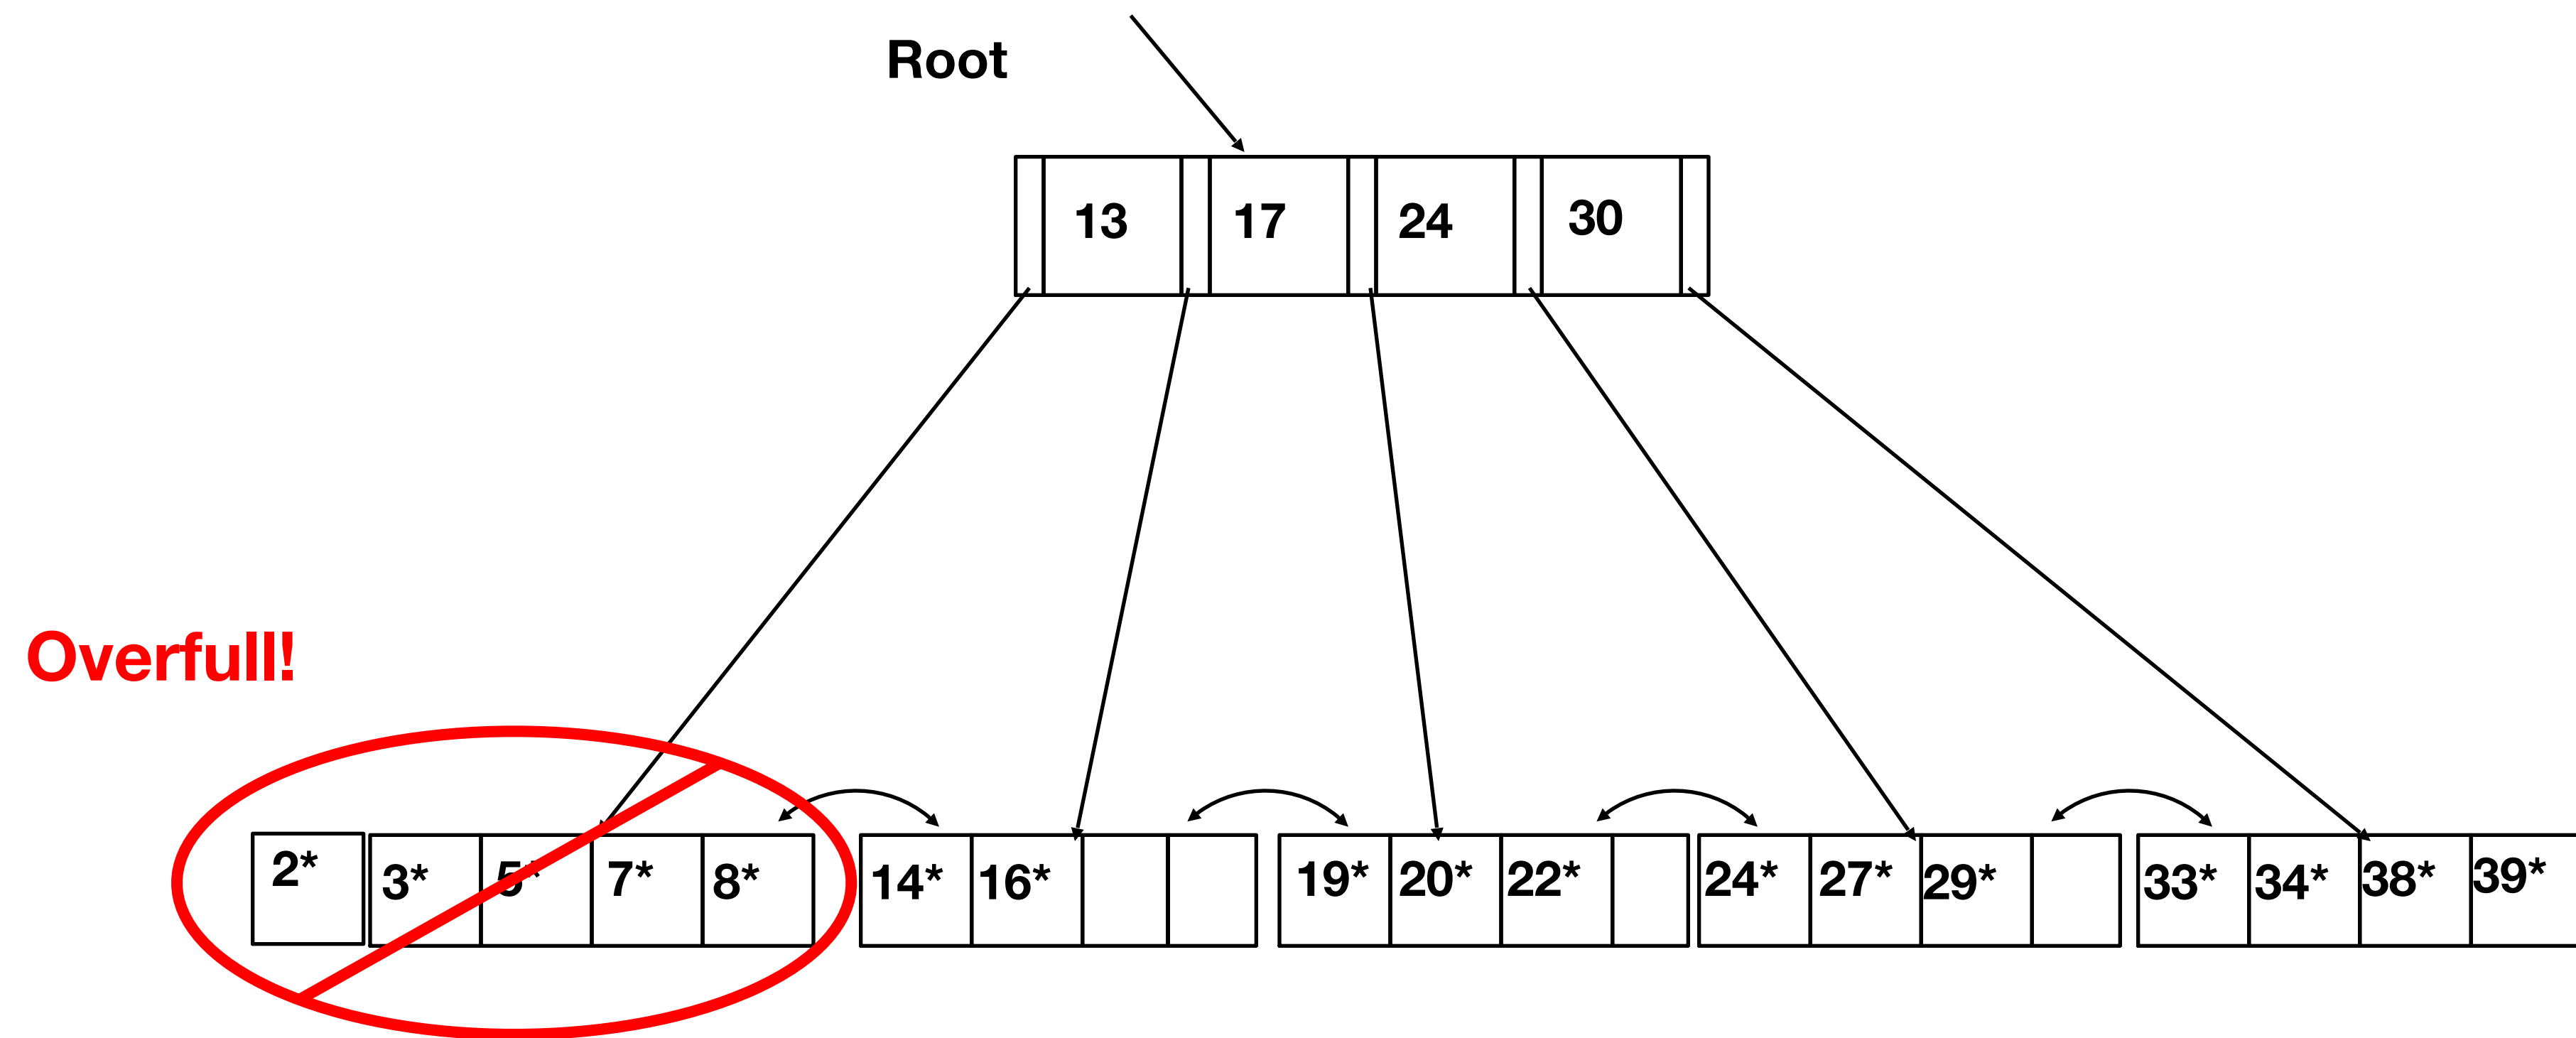
After Inserting 23*




Inserting 8* ...



Inserting 8* ...



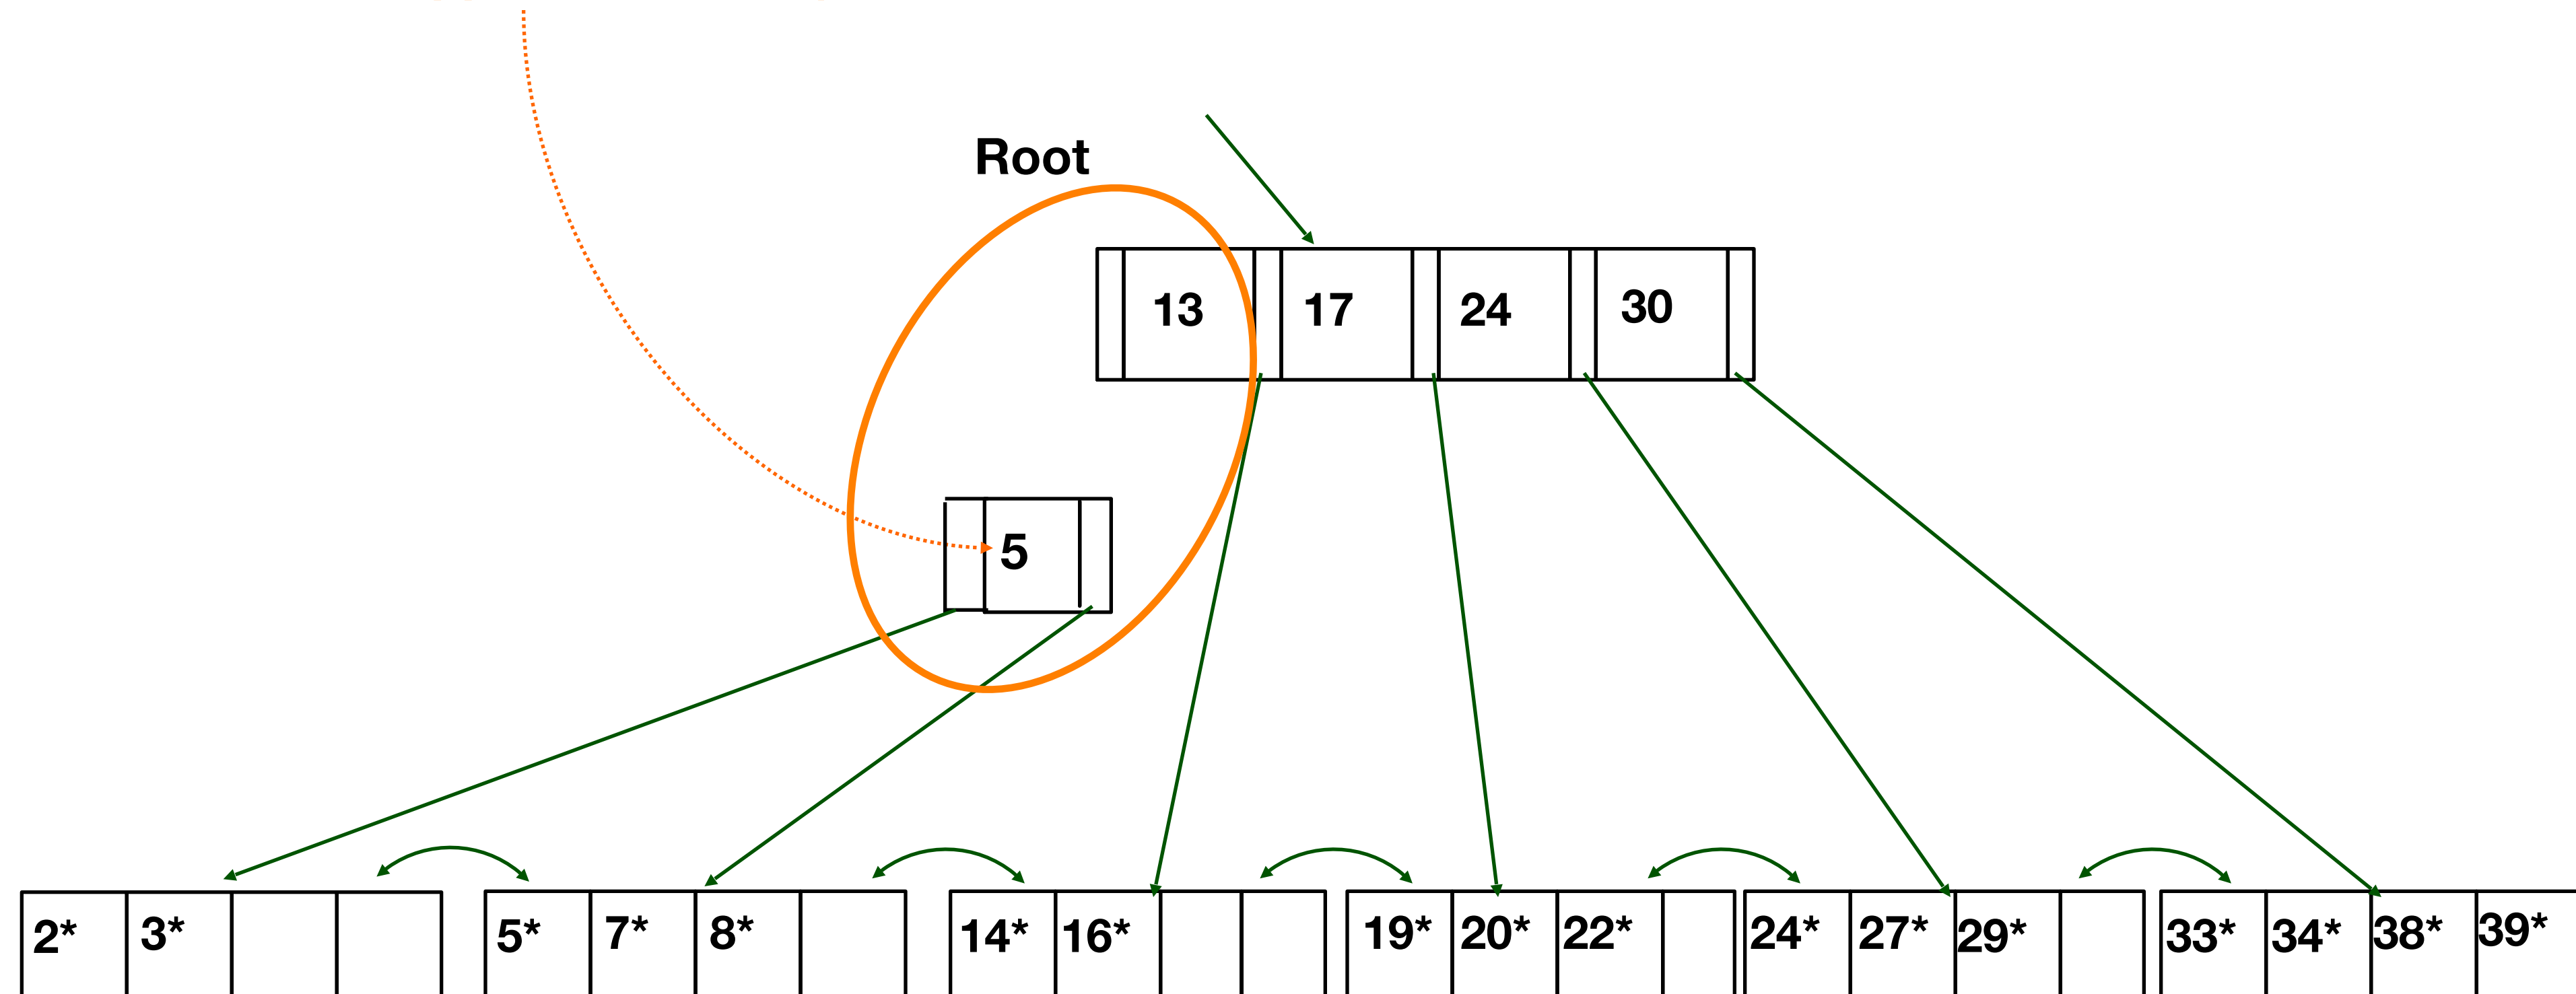
Splitting an Overfull Leaf



- New nodes not underfull: $2d+1 \rightarrow d$ and $d + 1$
- Single pointer in parent replaced by two pointers with key between
- Key **copied** from right node of new pair.


Inserting 8* ...

Entry to be inserted in parent node
(Note that 5 is **copied** up and
continues to appear in the leaf)




Inserting 8* ...

(Note that 5 is **copied up** and continues to appear in the leaf)

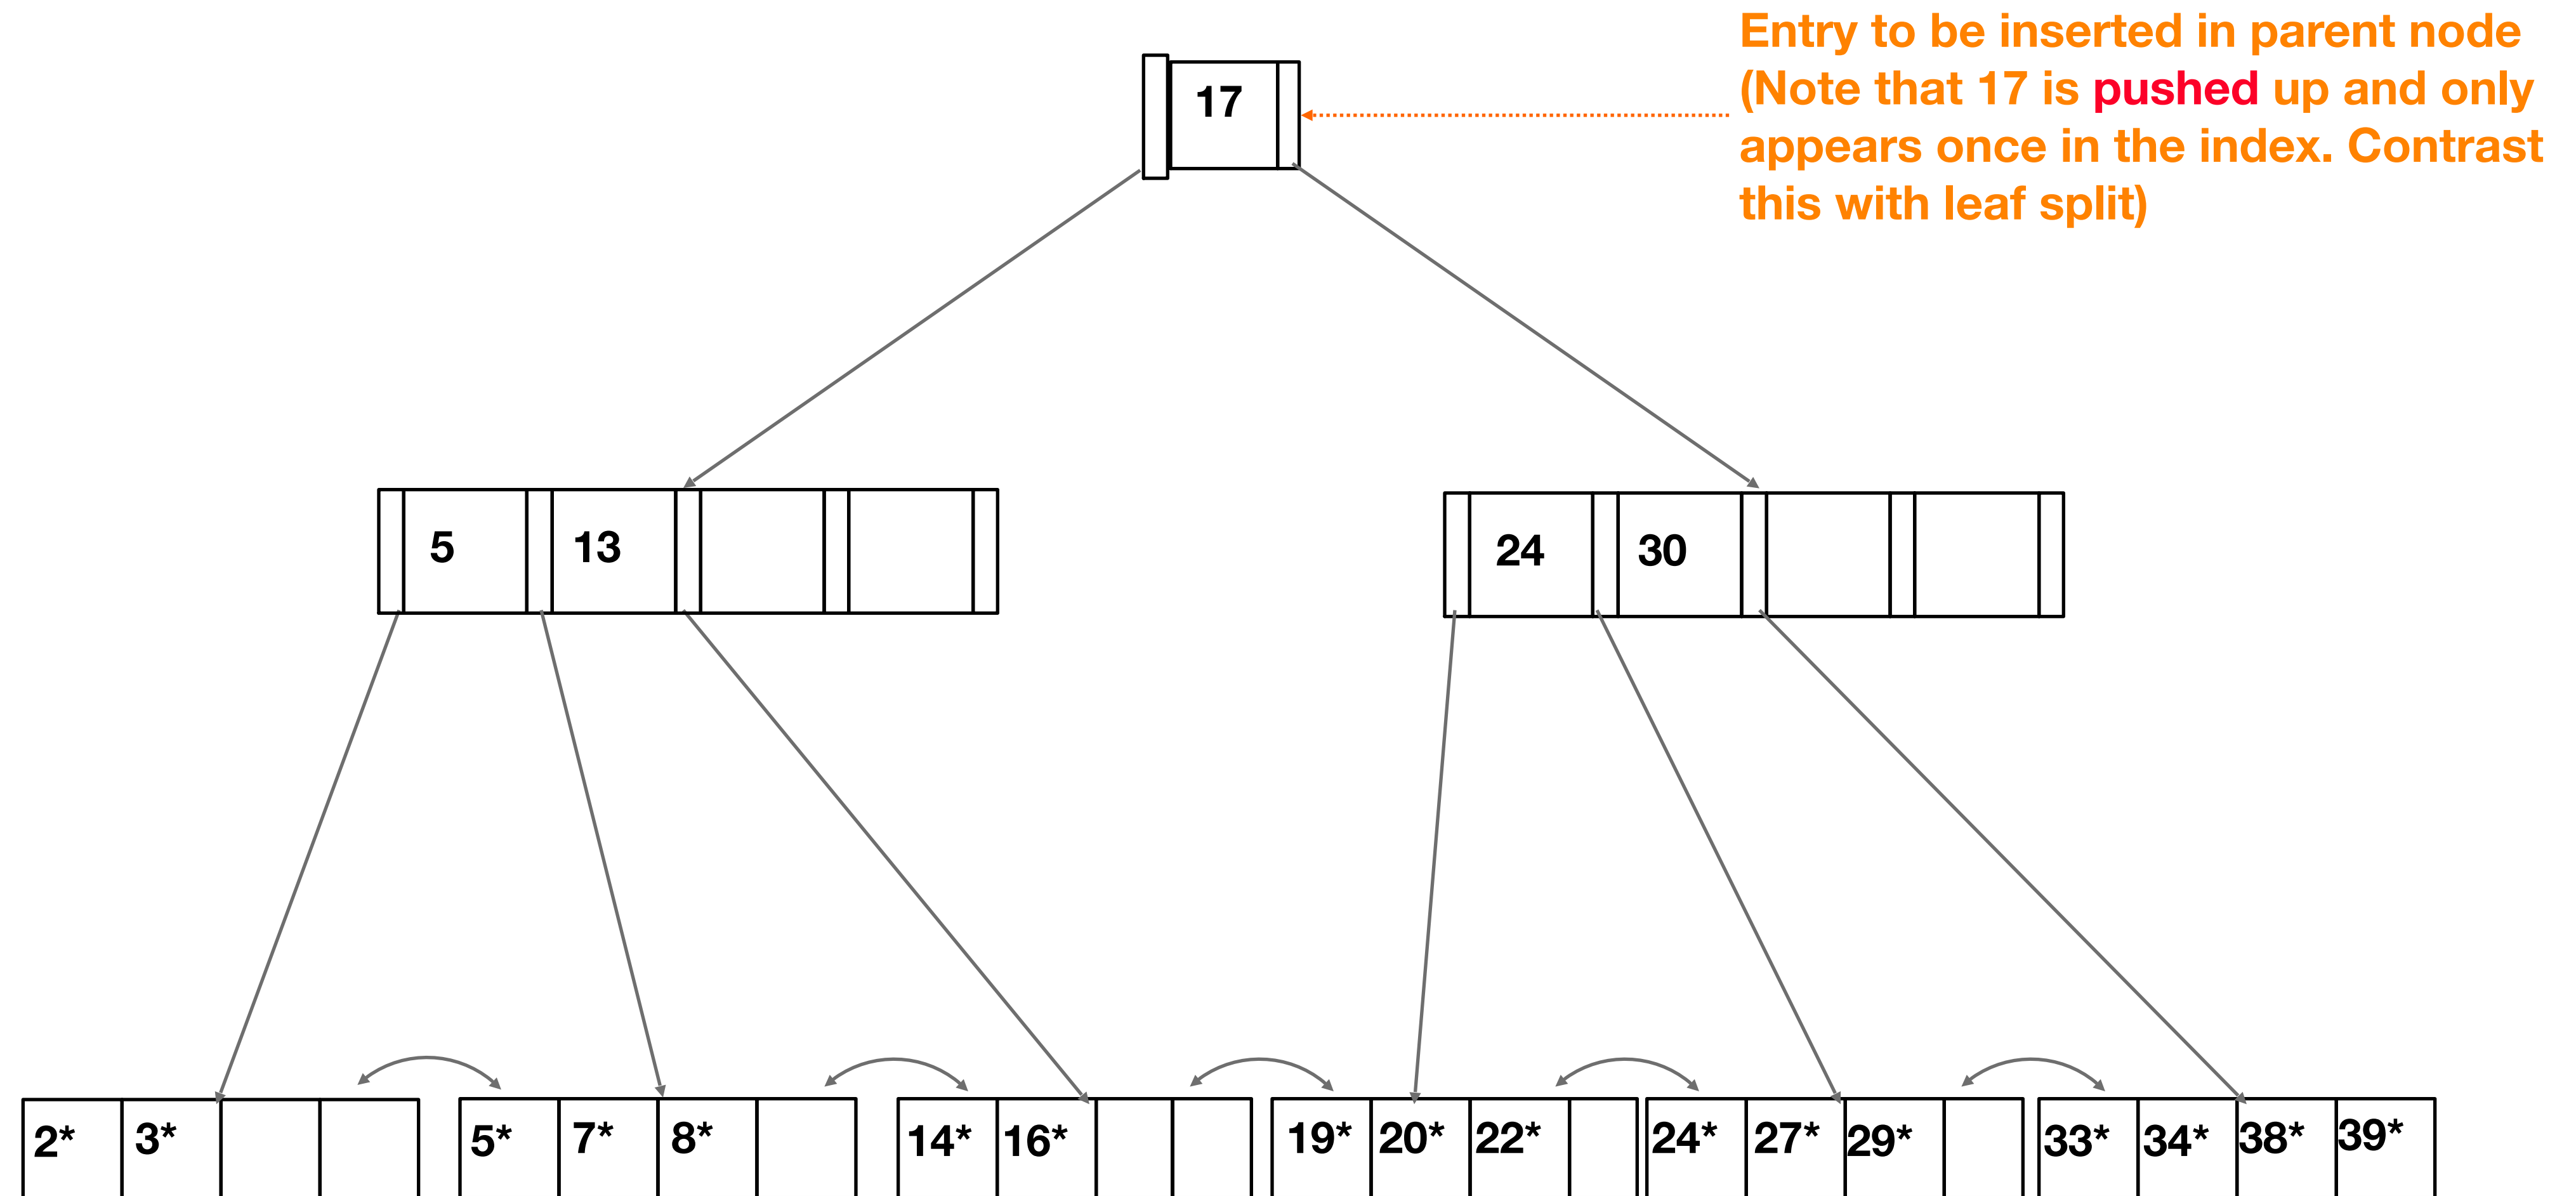


Splitting an Overfull Nonleaf ...

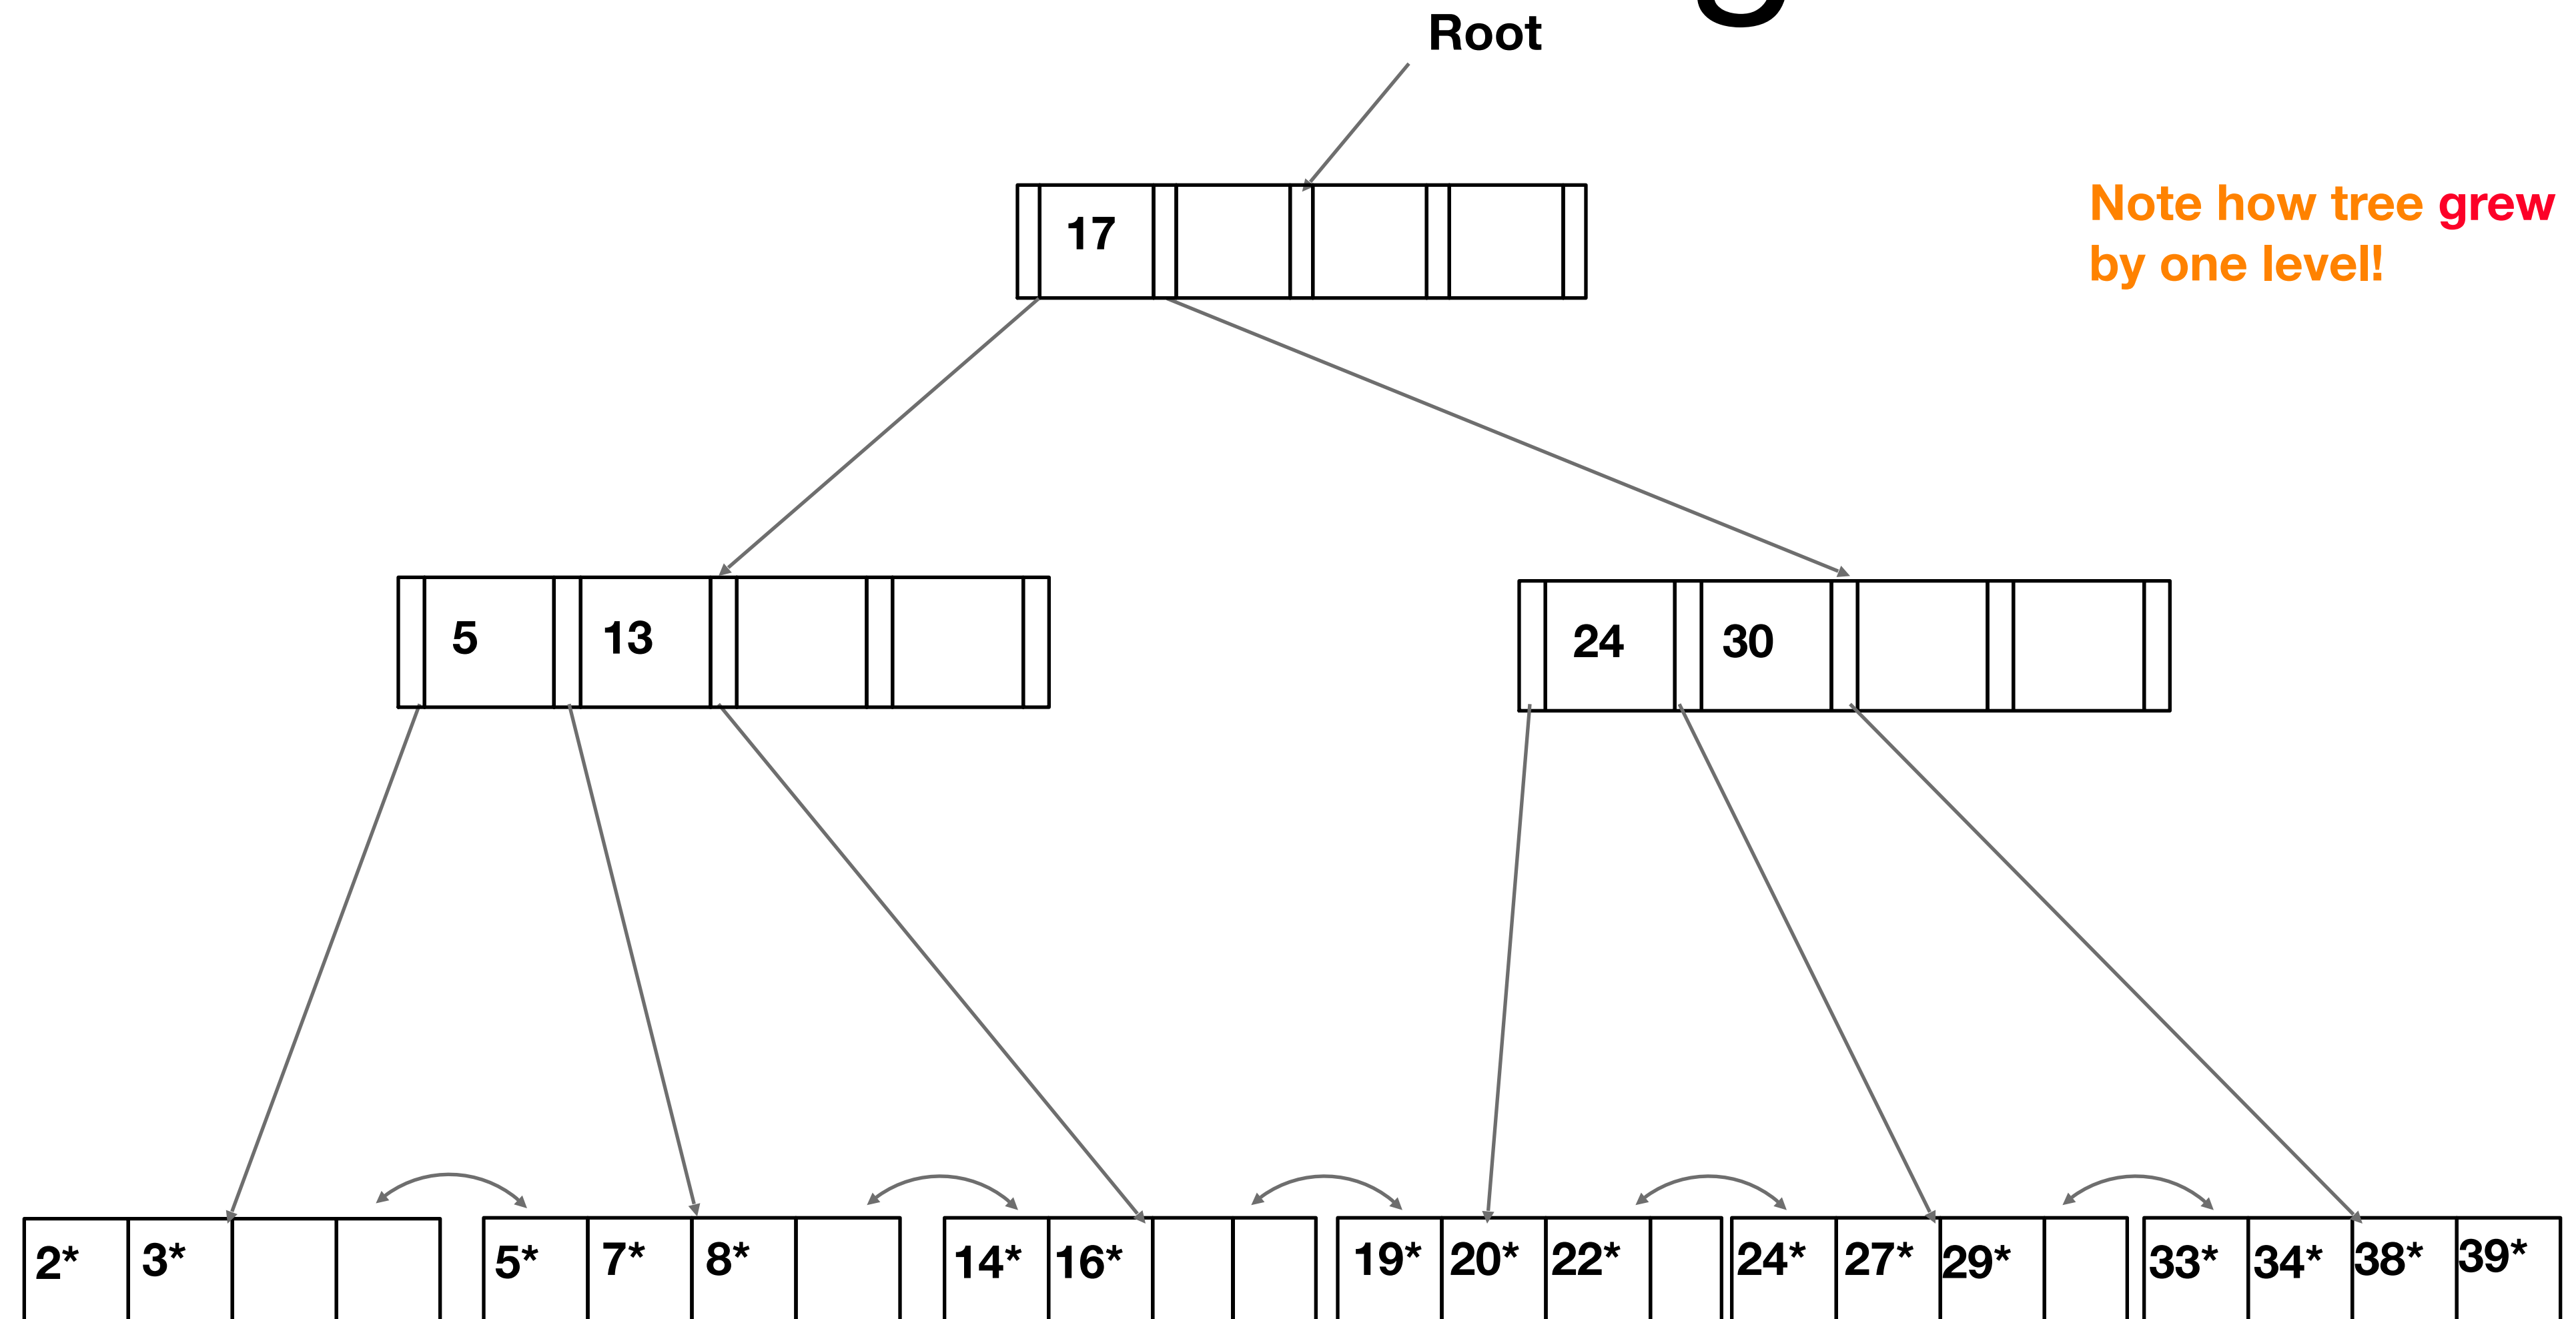


- Like leaf split except
- $2d+1 \rightarrow d$ and d (and **1**)
- **Key moved** from right node of new pair, not copied.

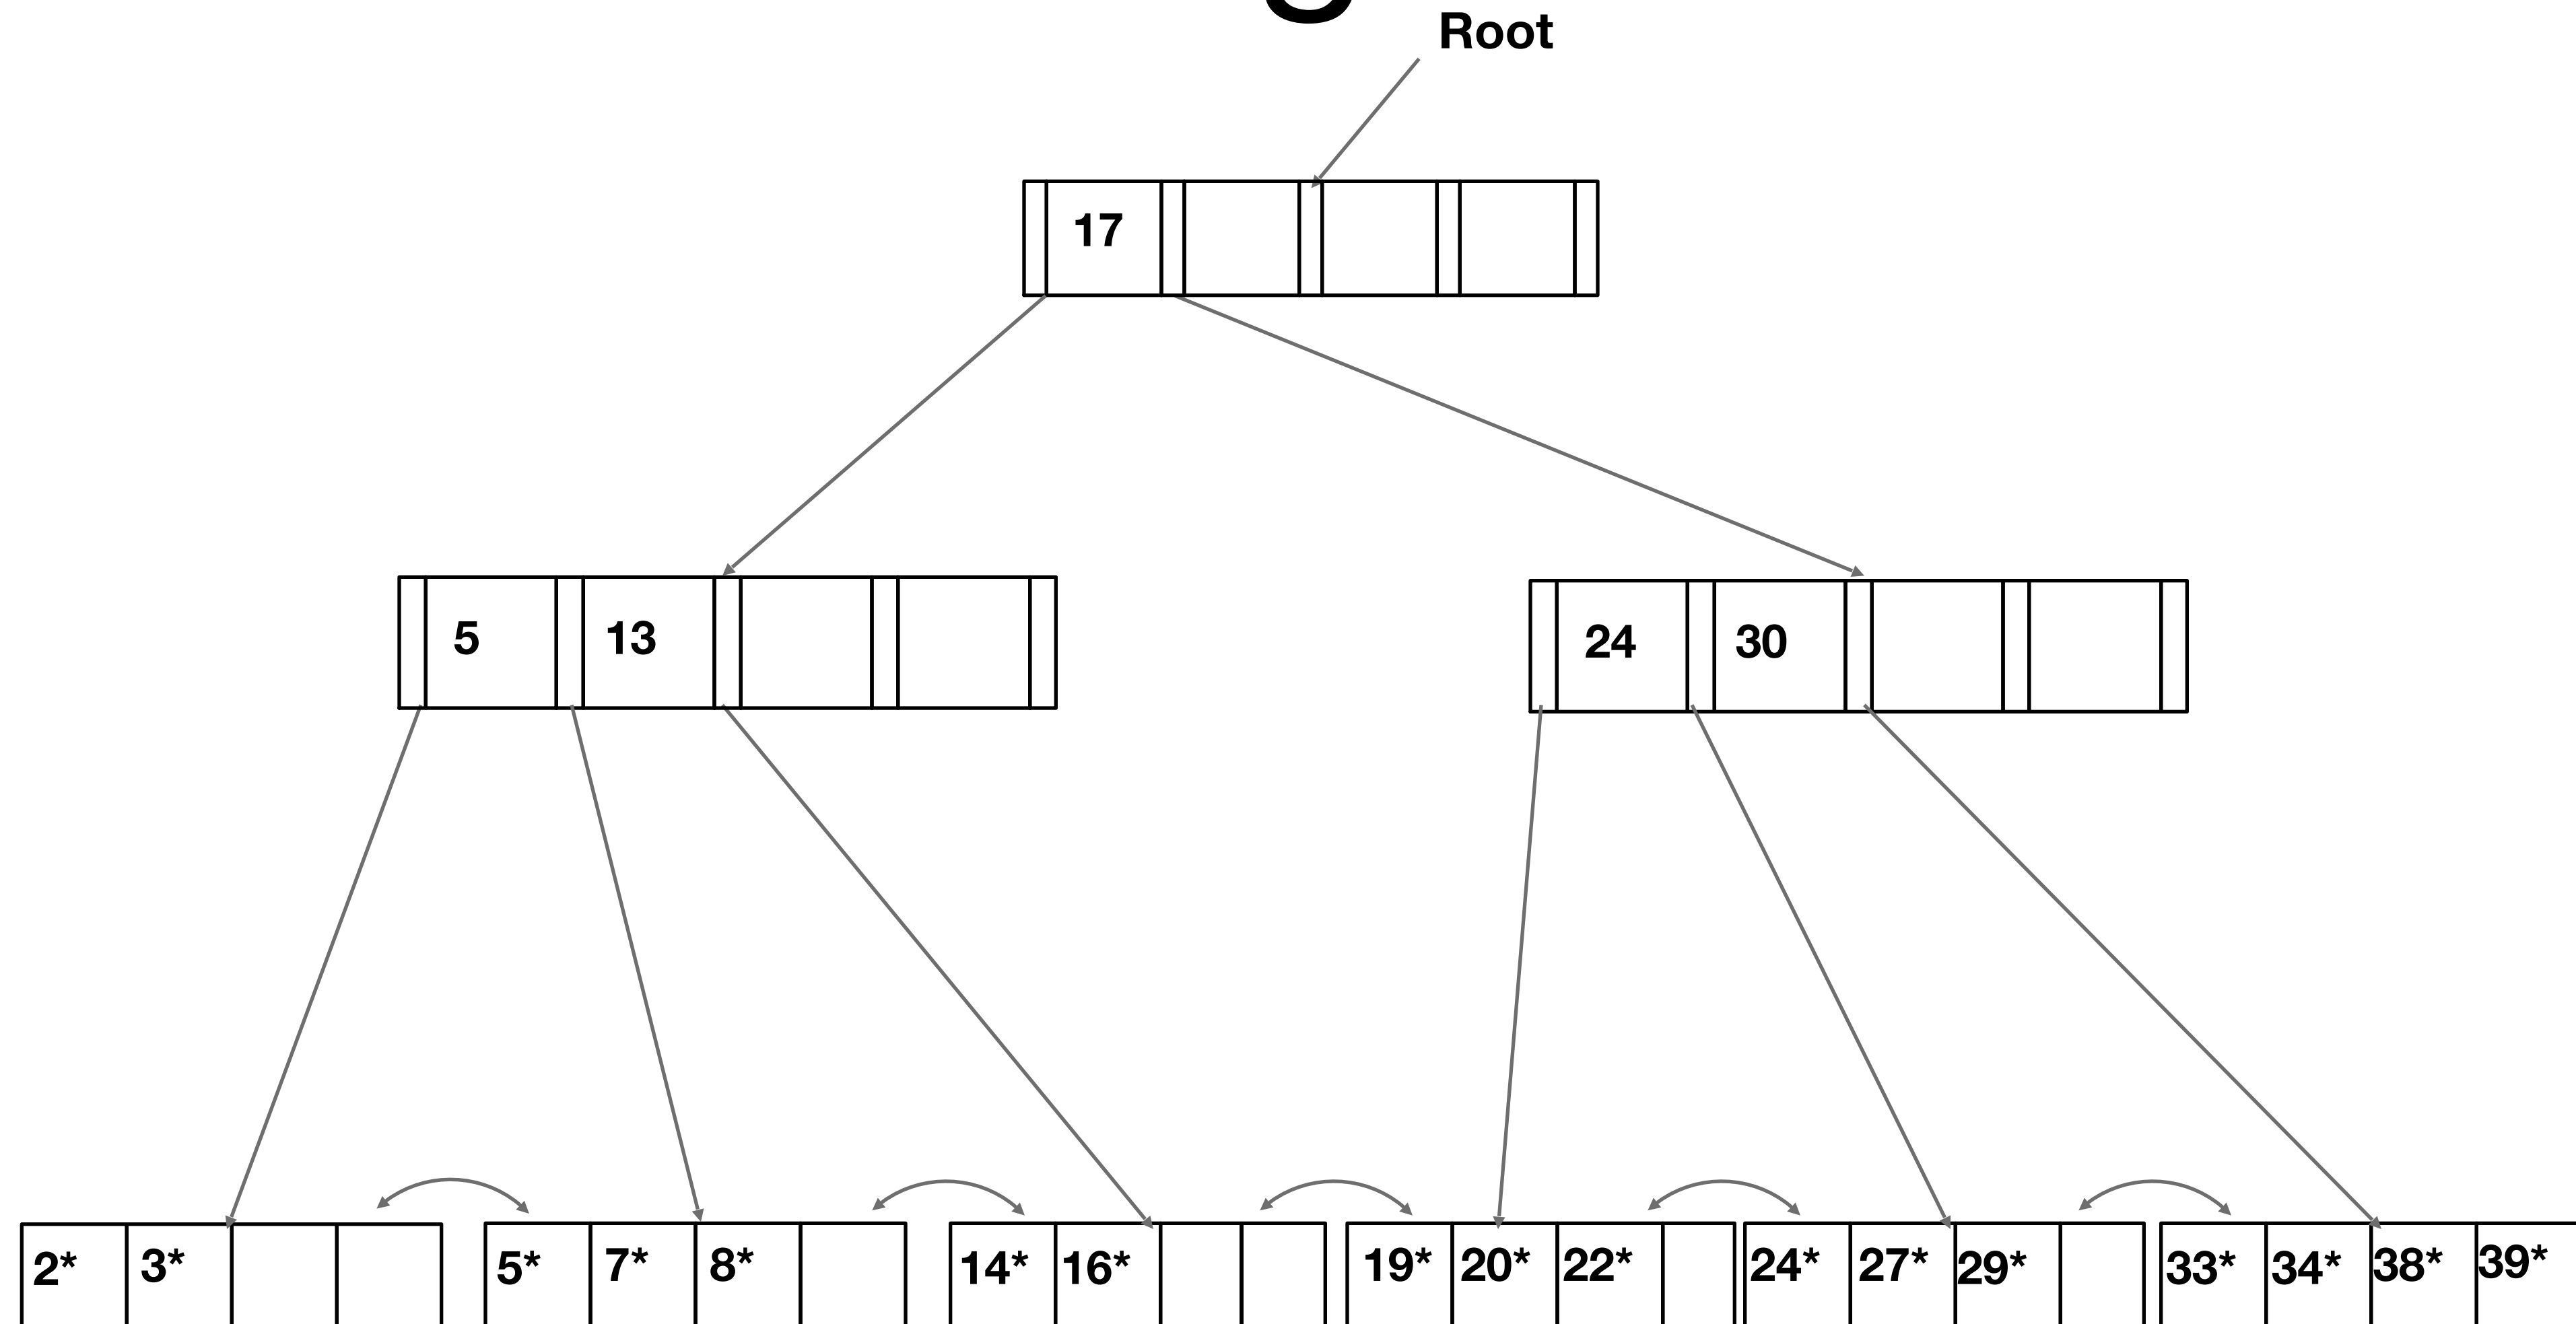
Inserting 8* ...



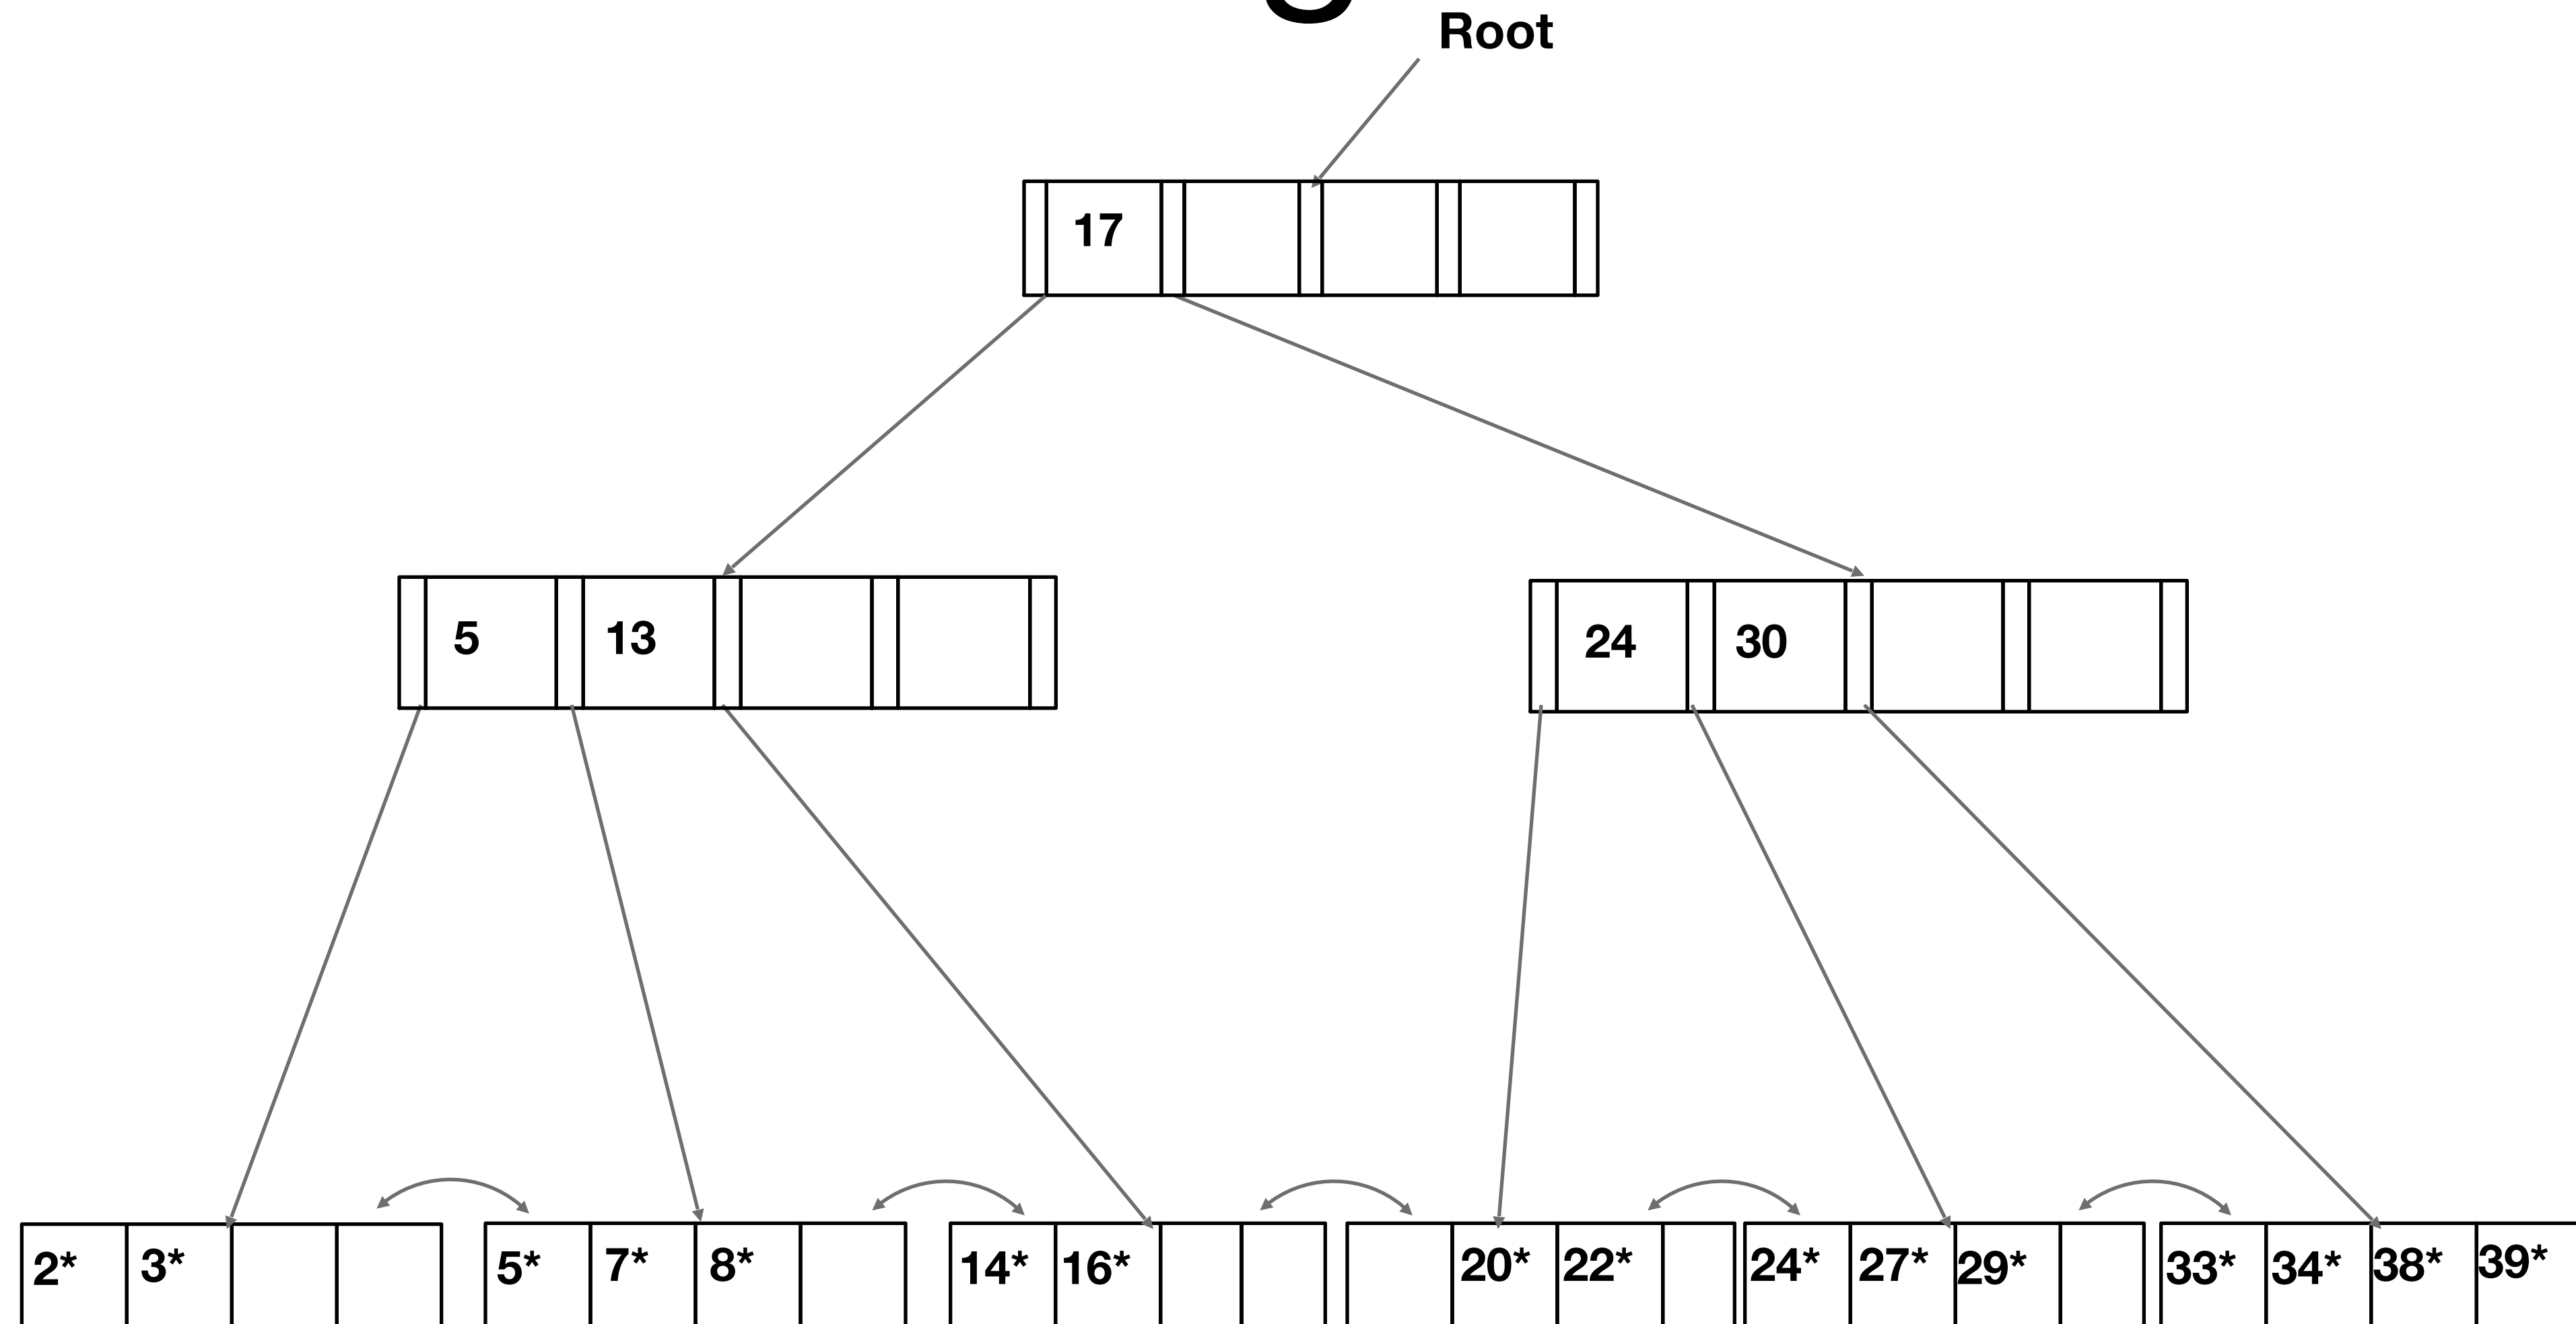
After Inserting 8*



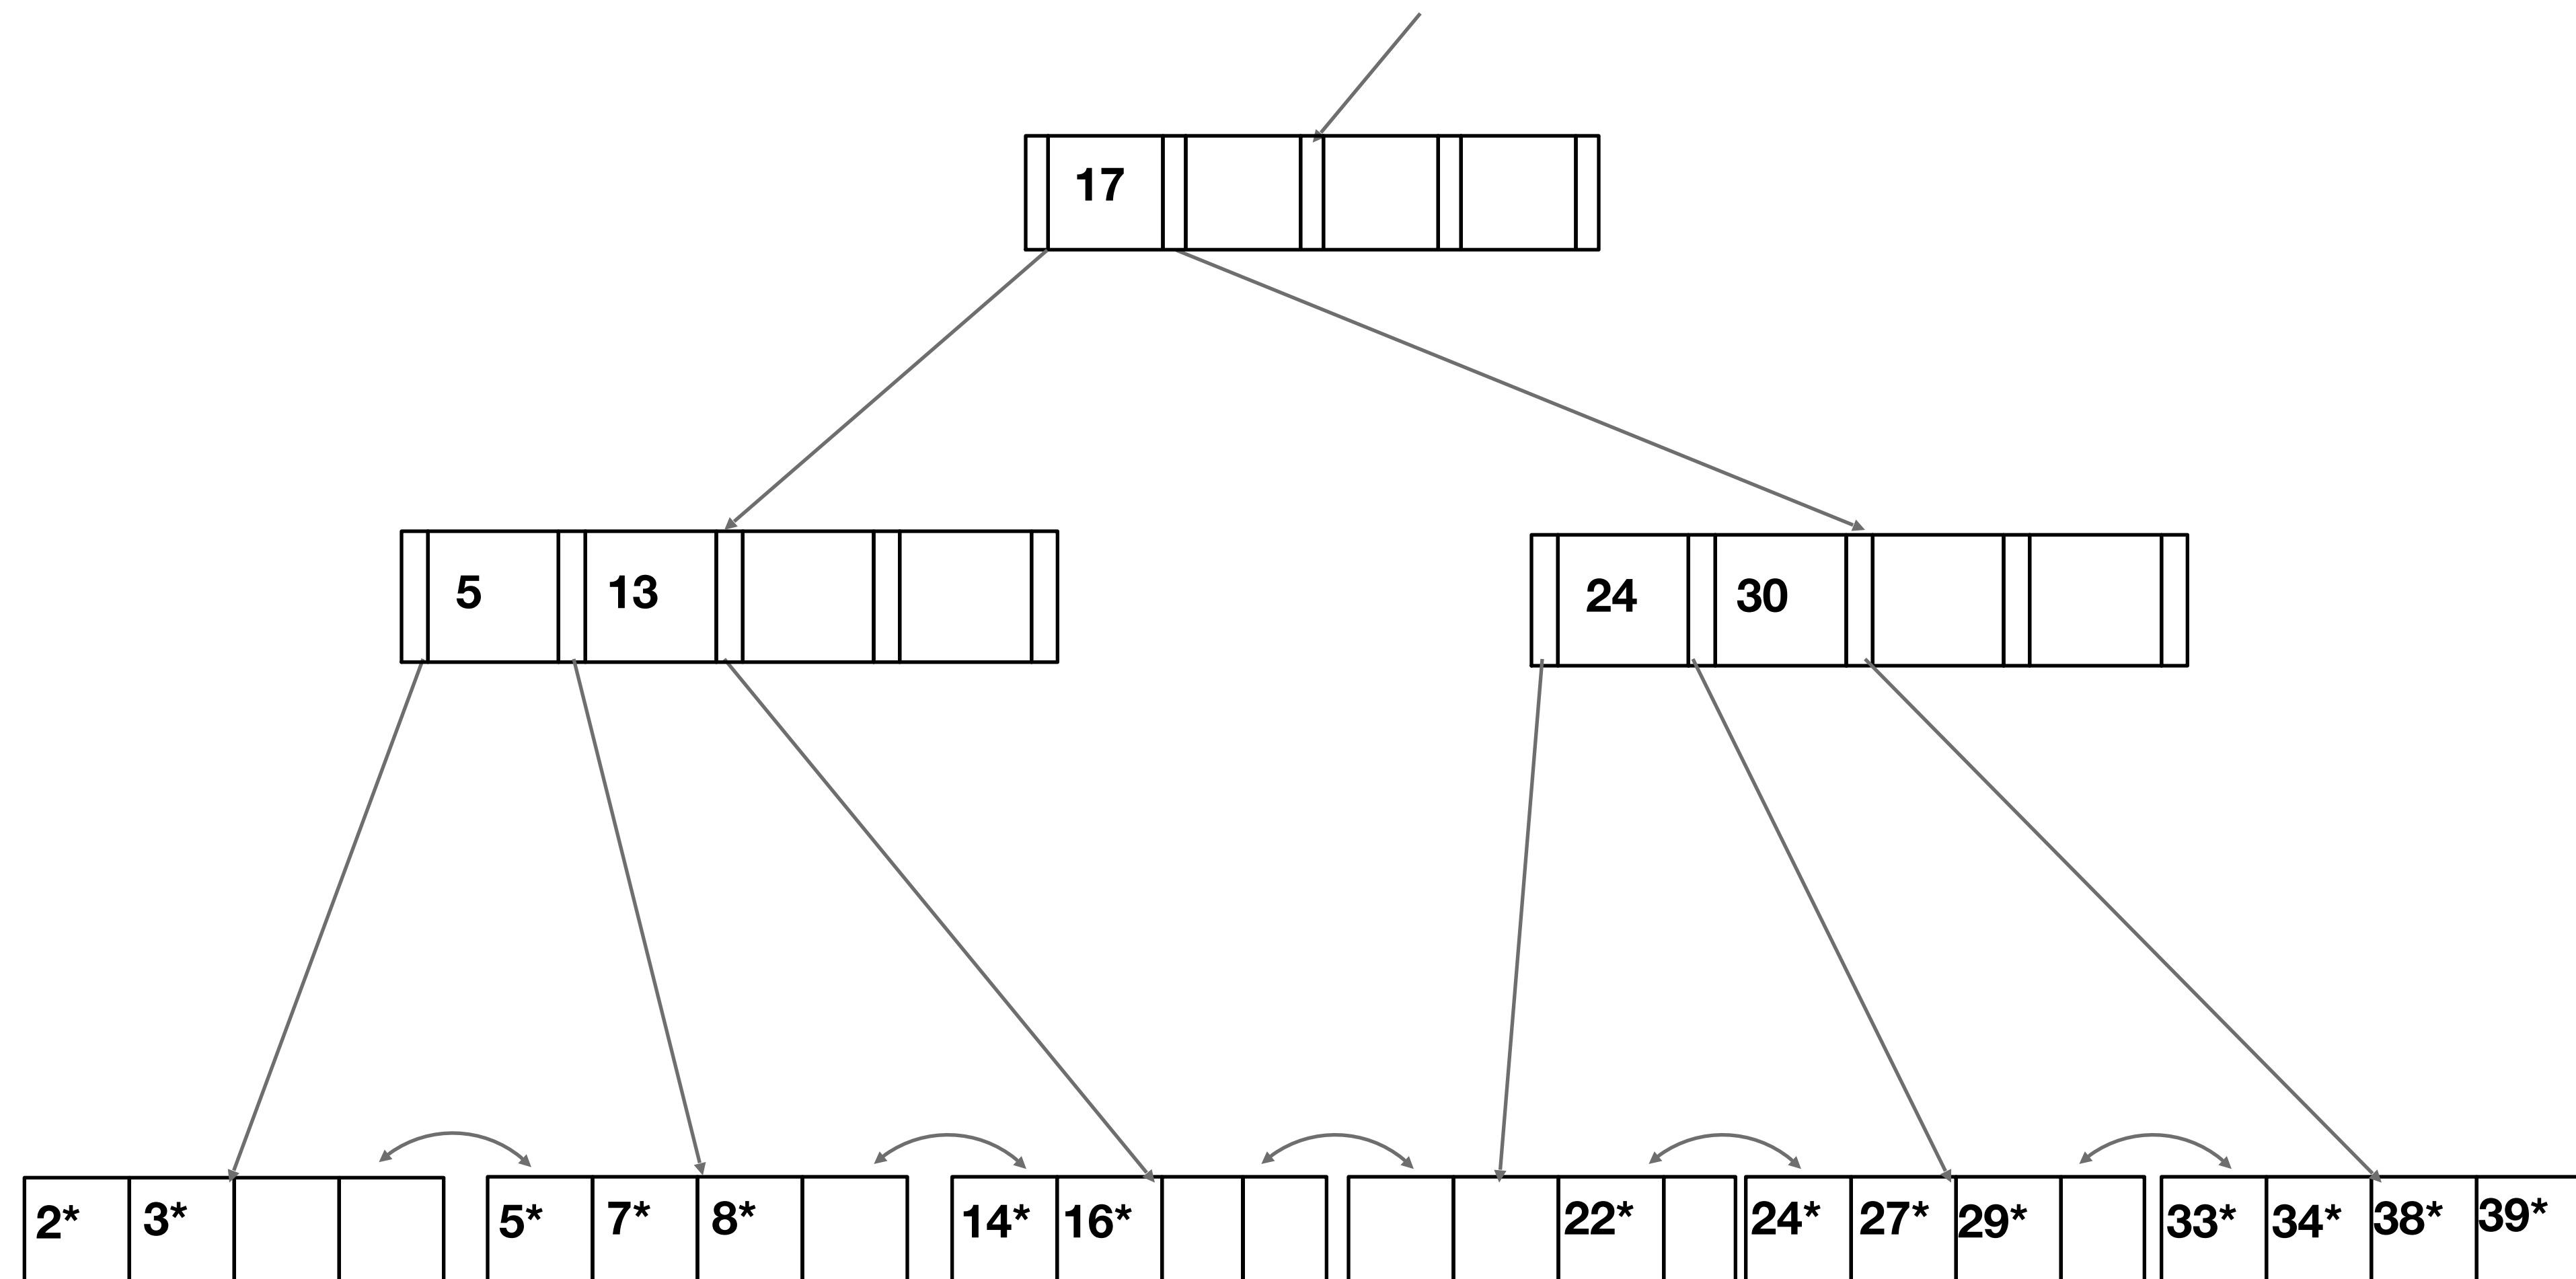
Deleting 19* ...



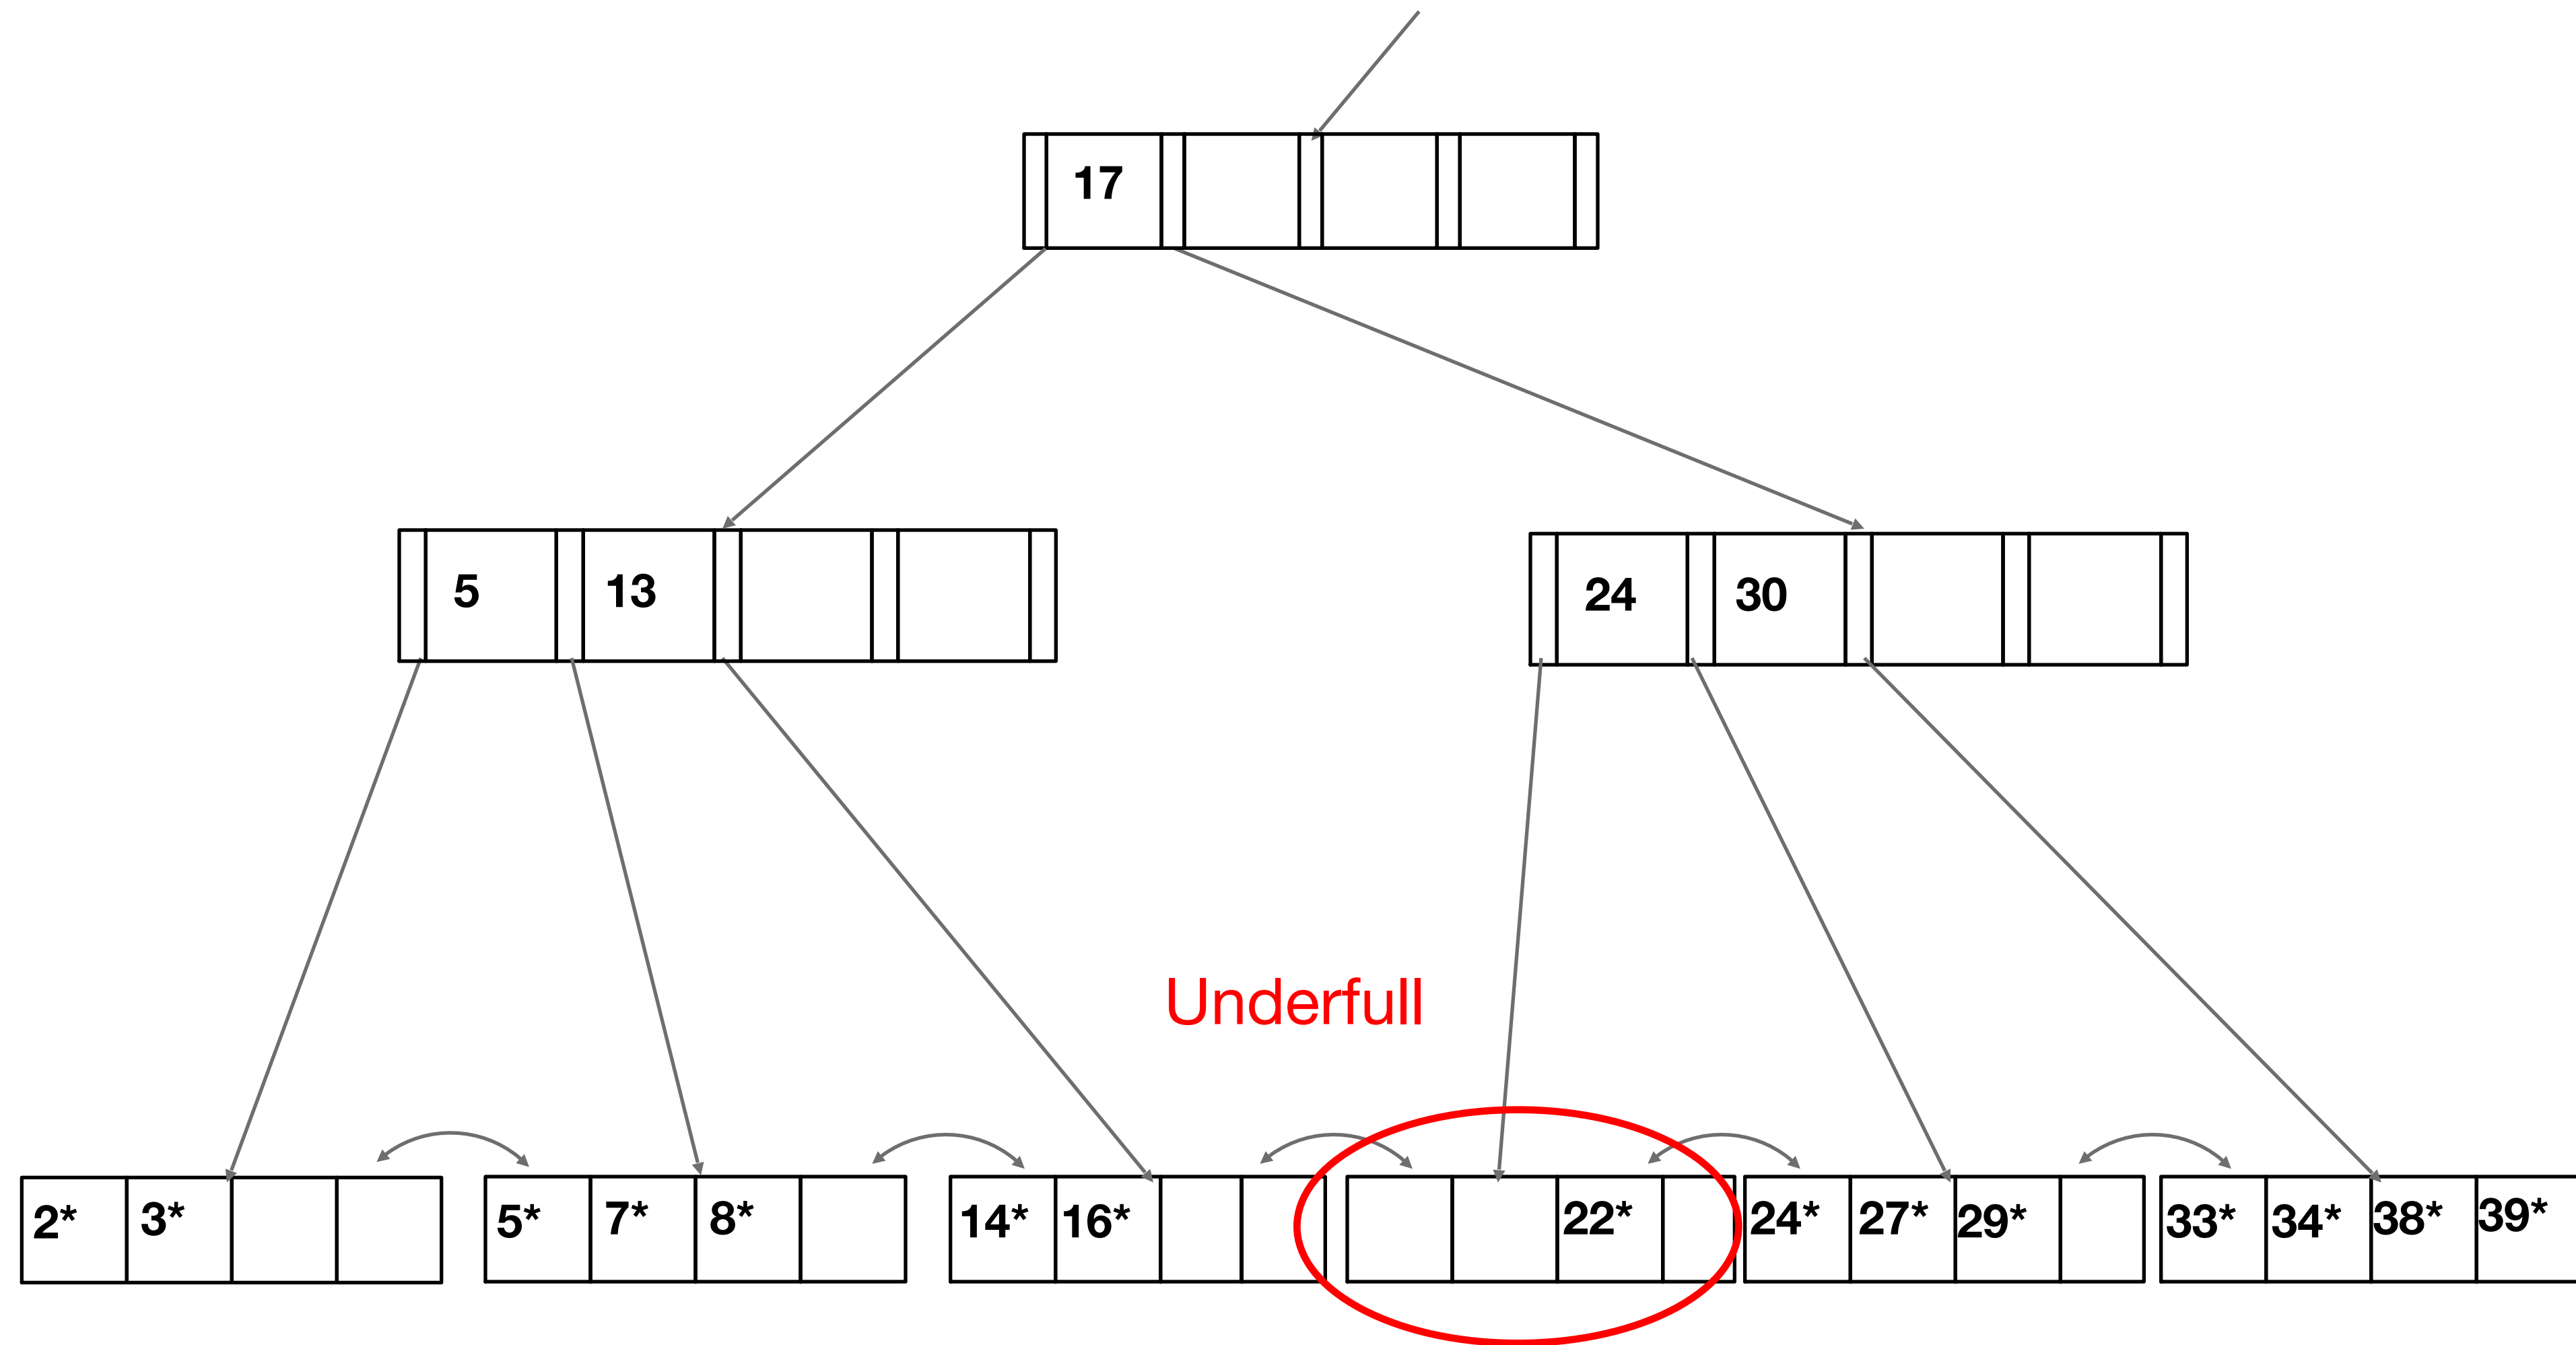
Deleting 20* ...




Deleting 20* ...



Deleting 20* ...

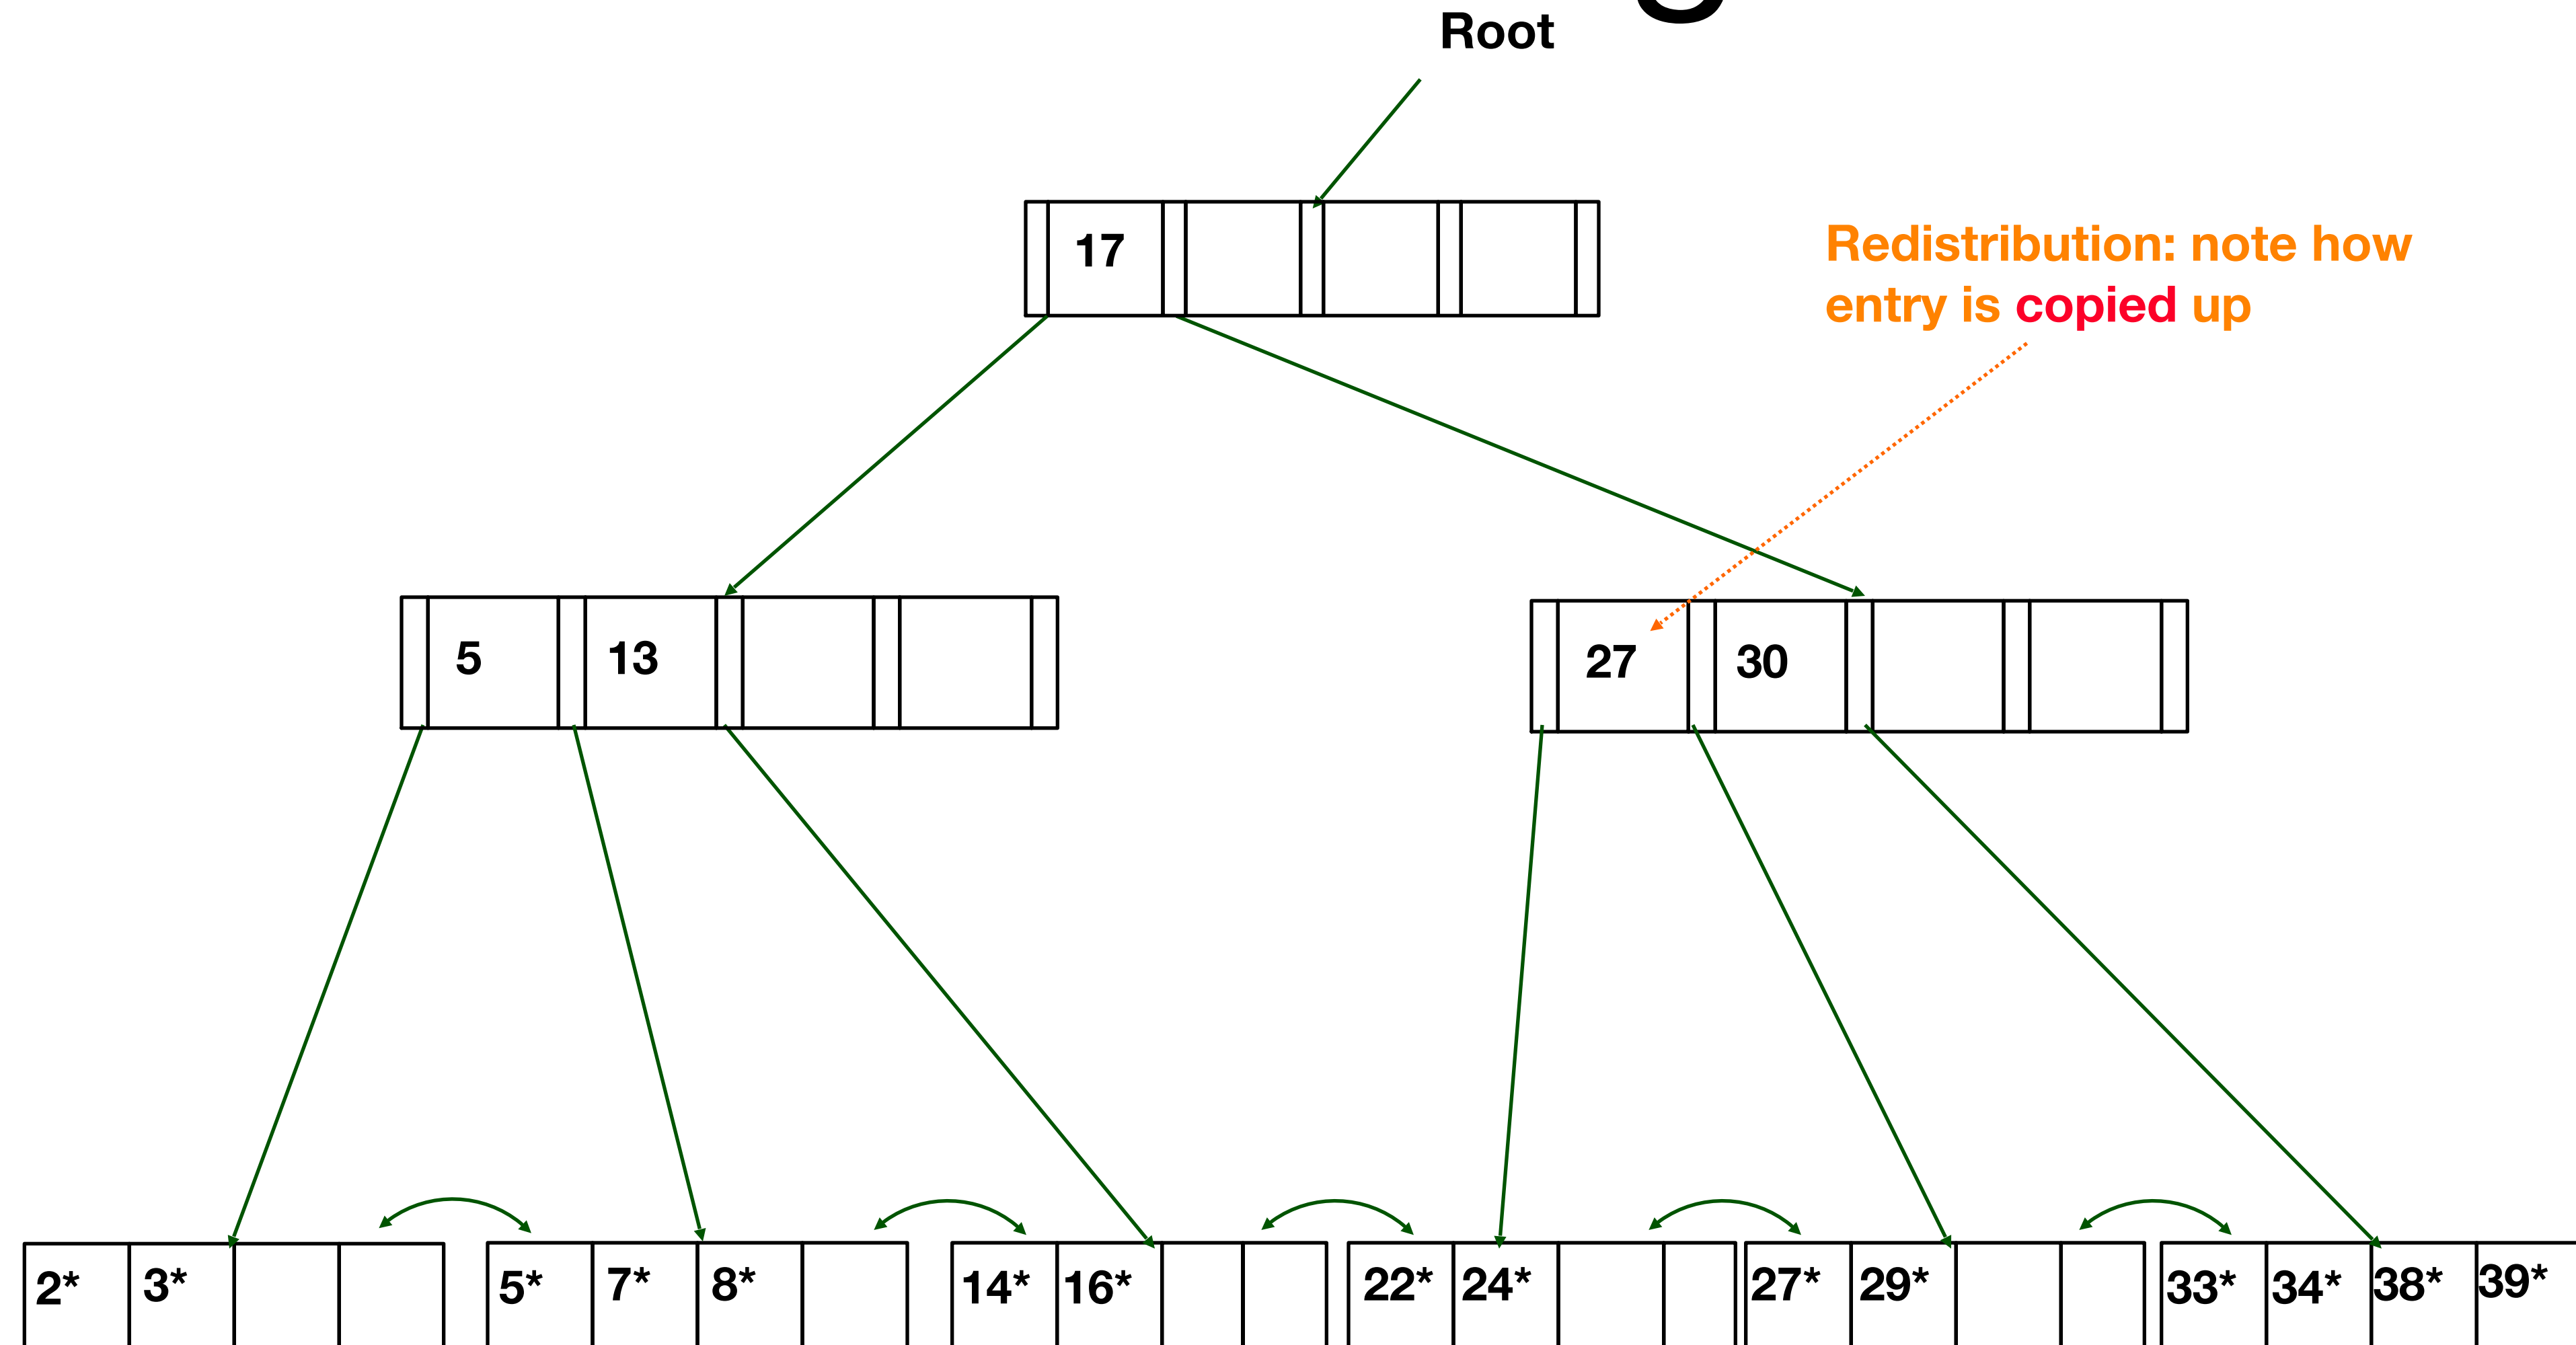


Redistributing Underfull Leaves ...

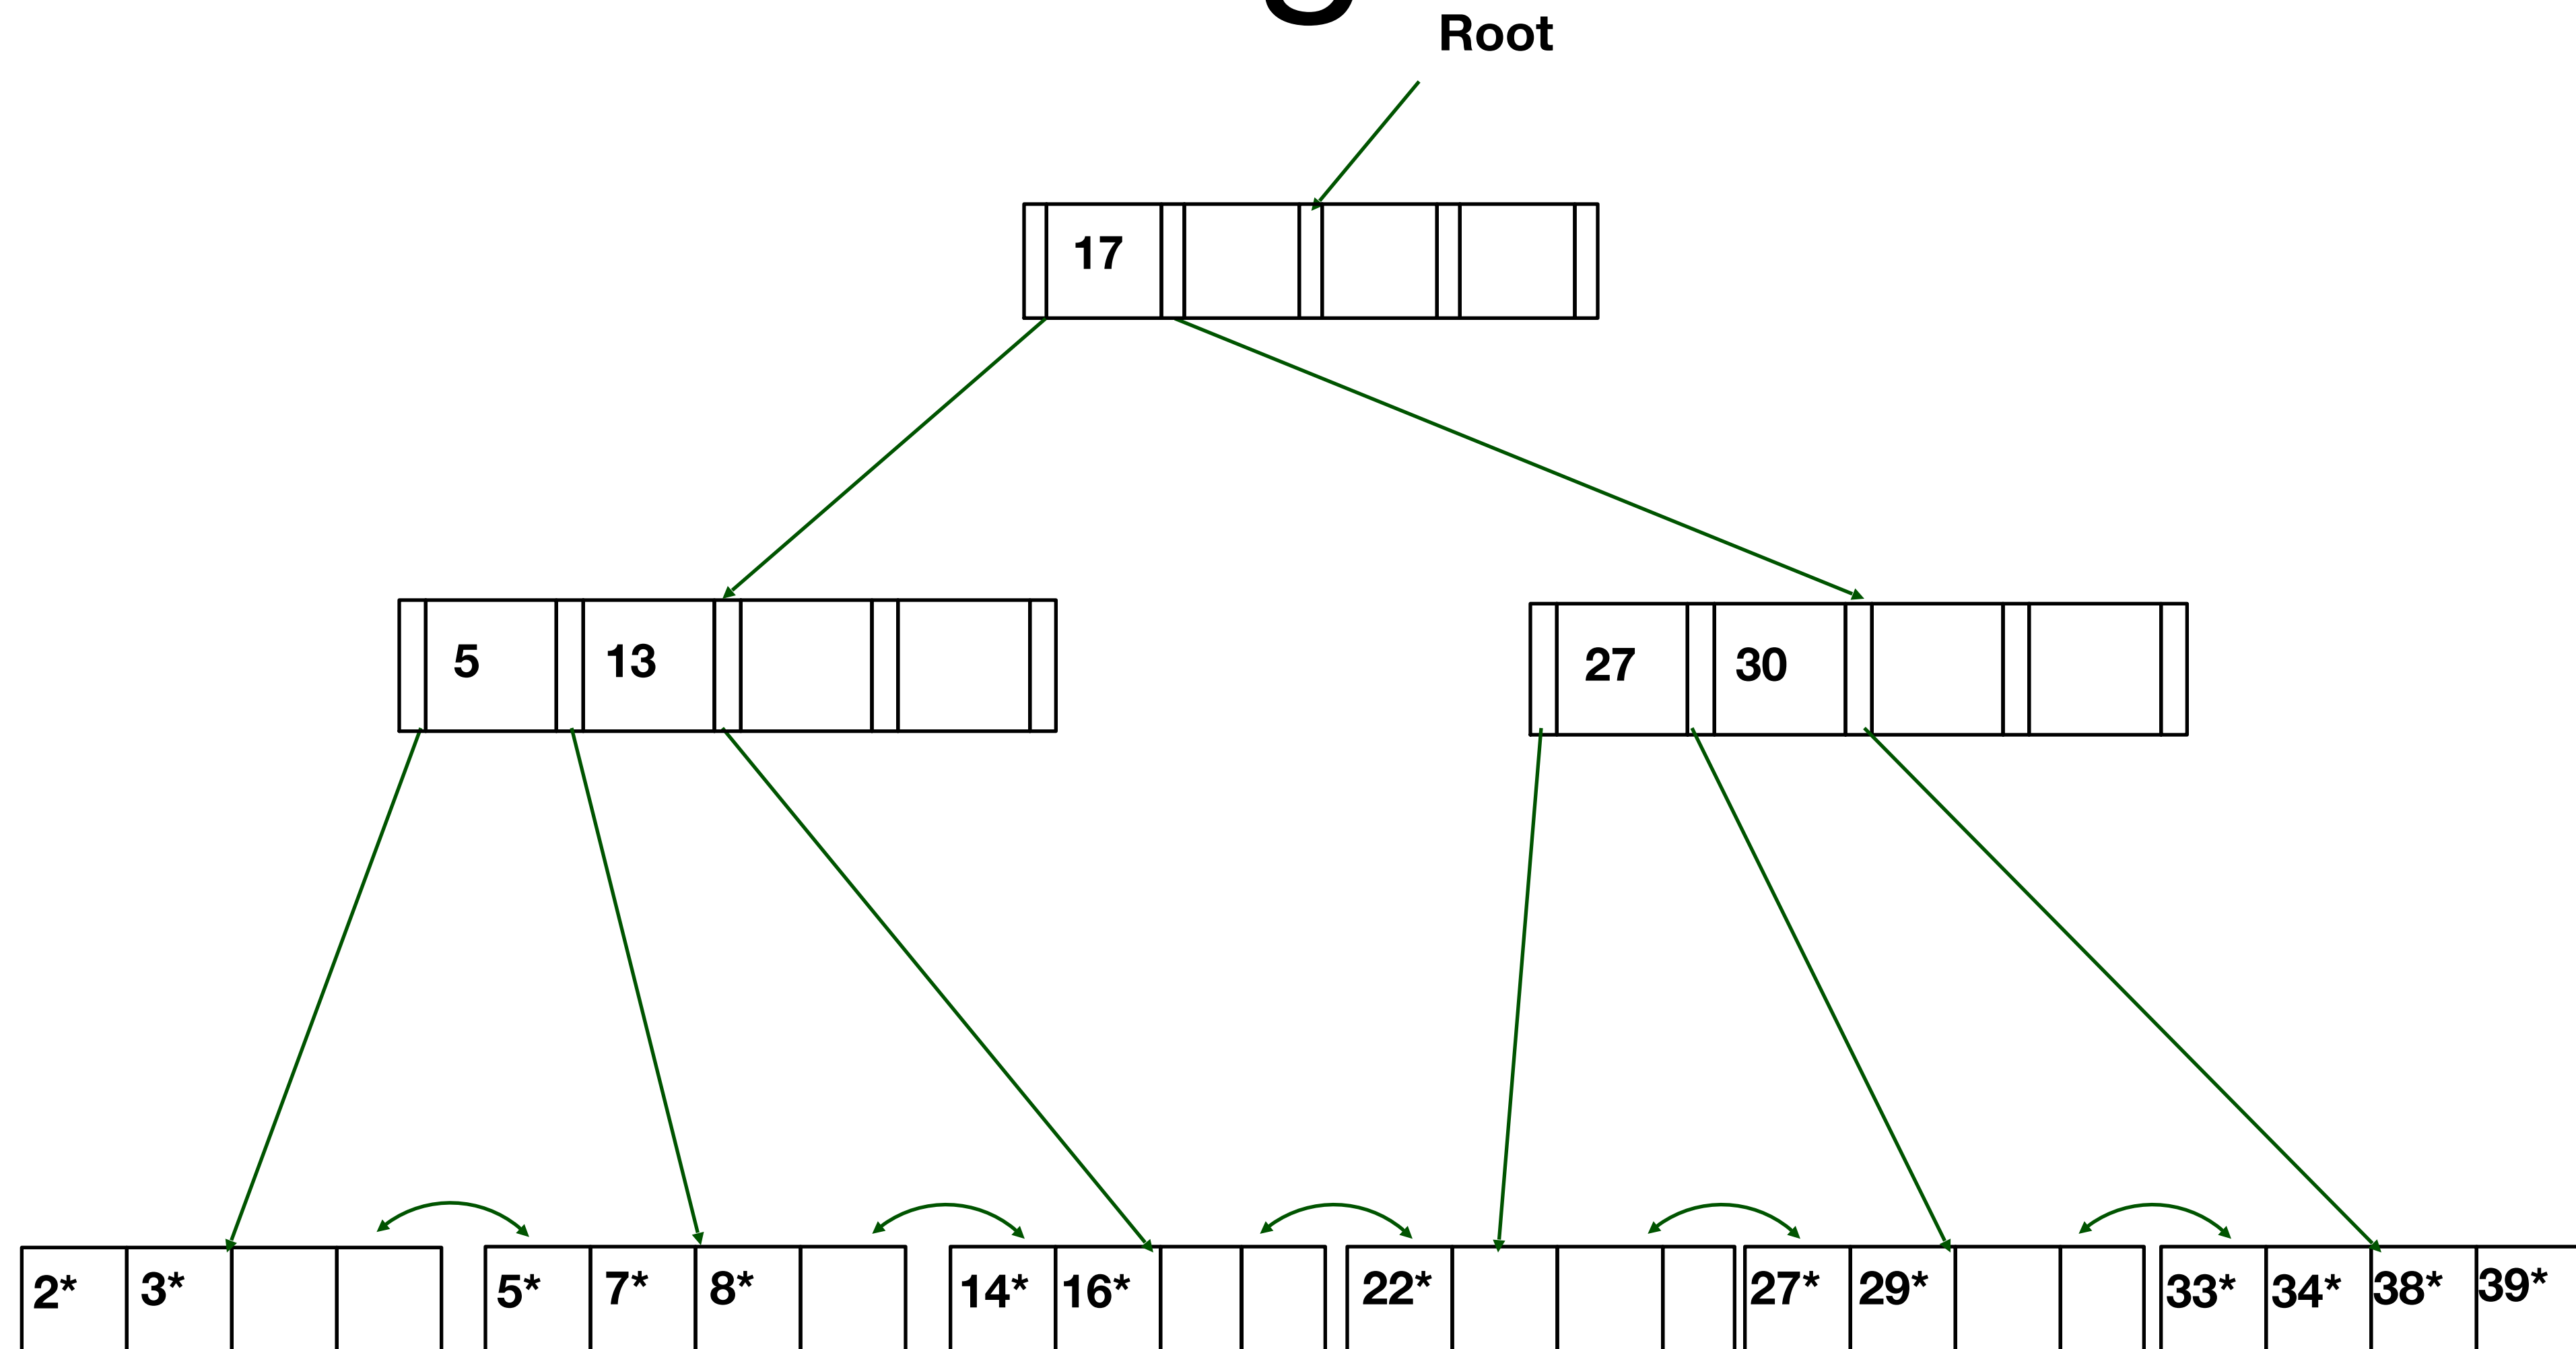


- Key in parent is modified
- Degree of parent unchanged
- Note there are 3 disk writes

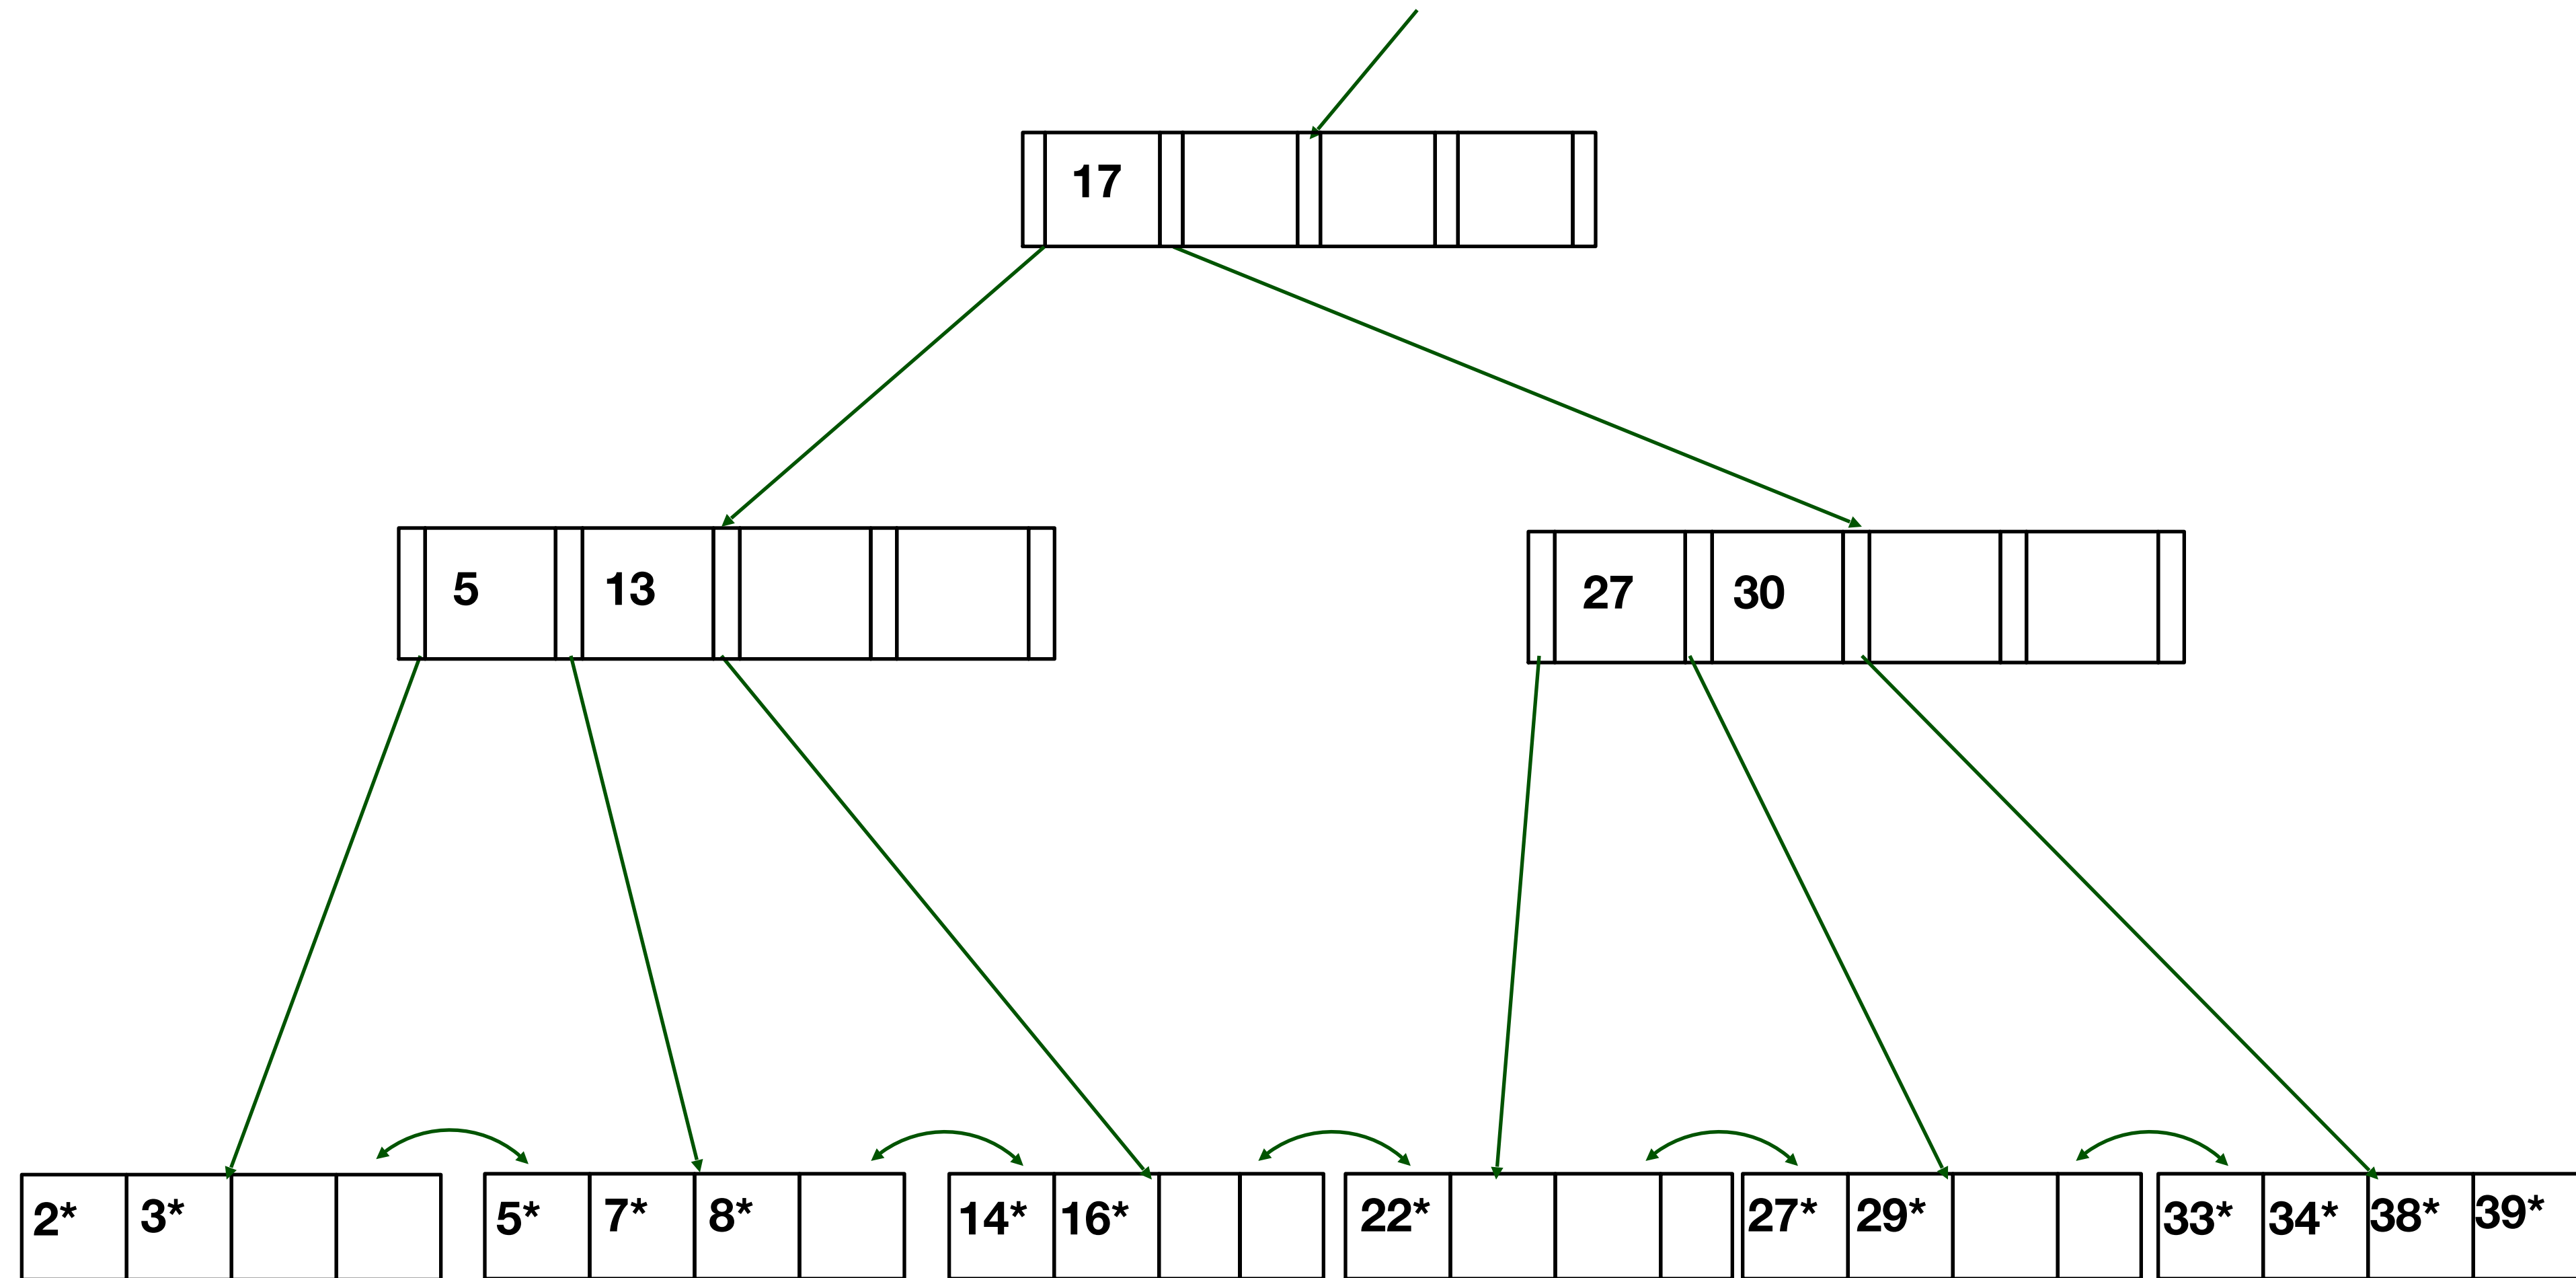
After Deleting 20*



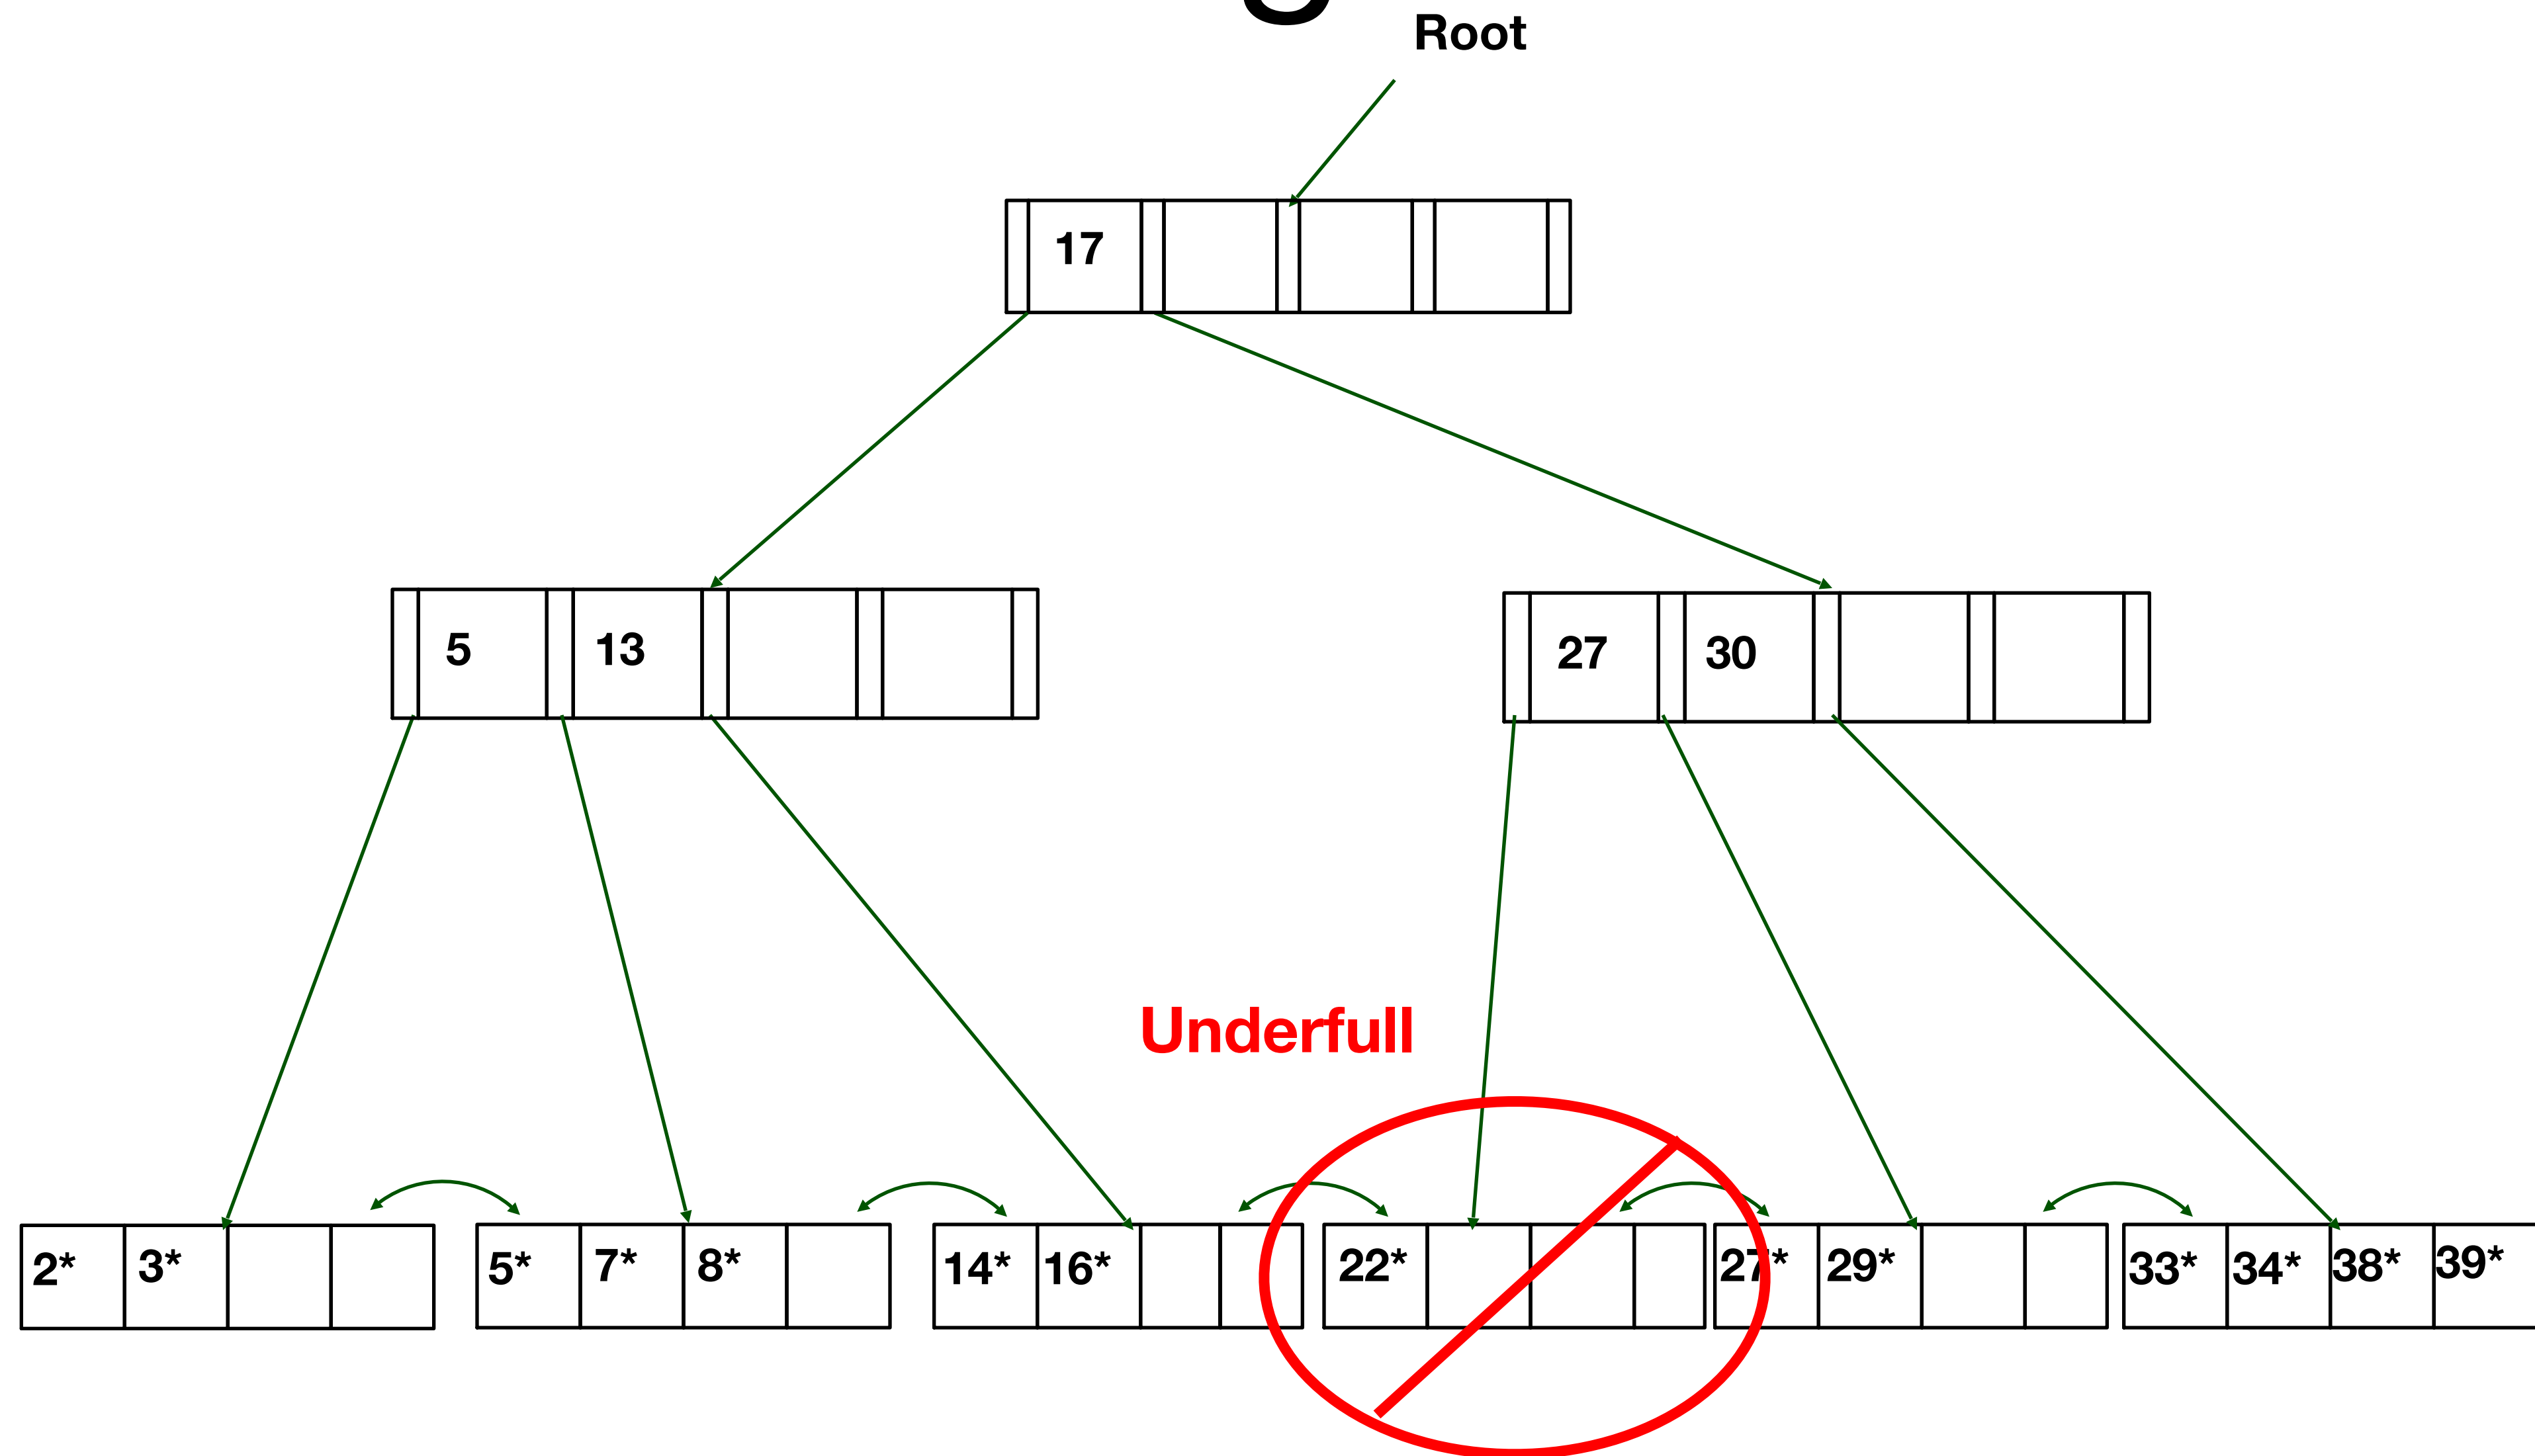
Deleting 24* ...




Deleting 24* ...



Deleting 24* ...

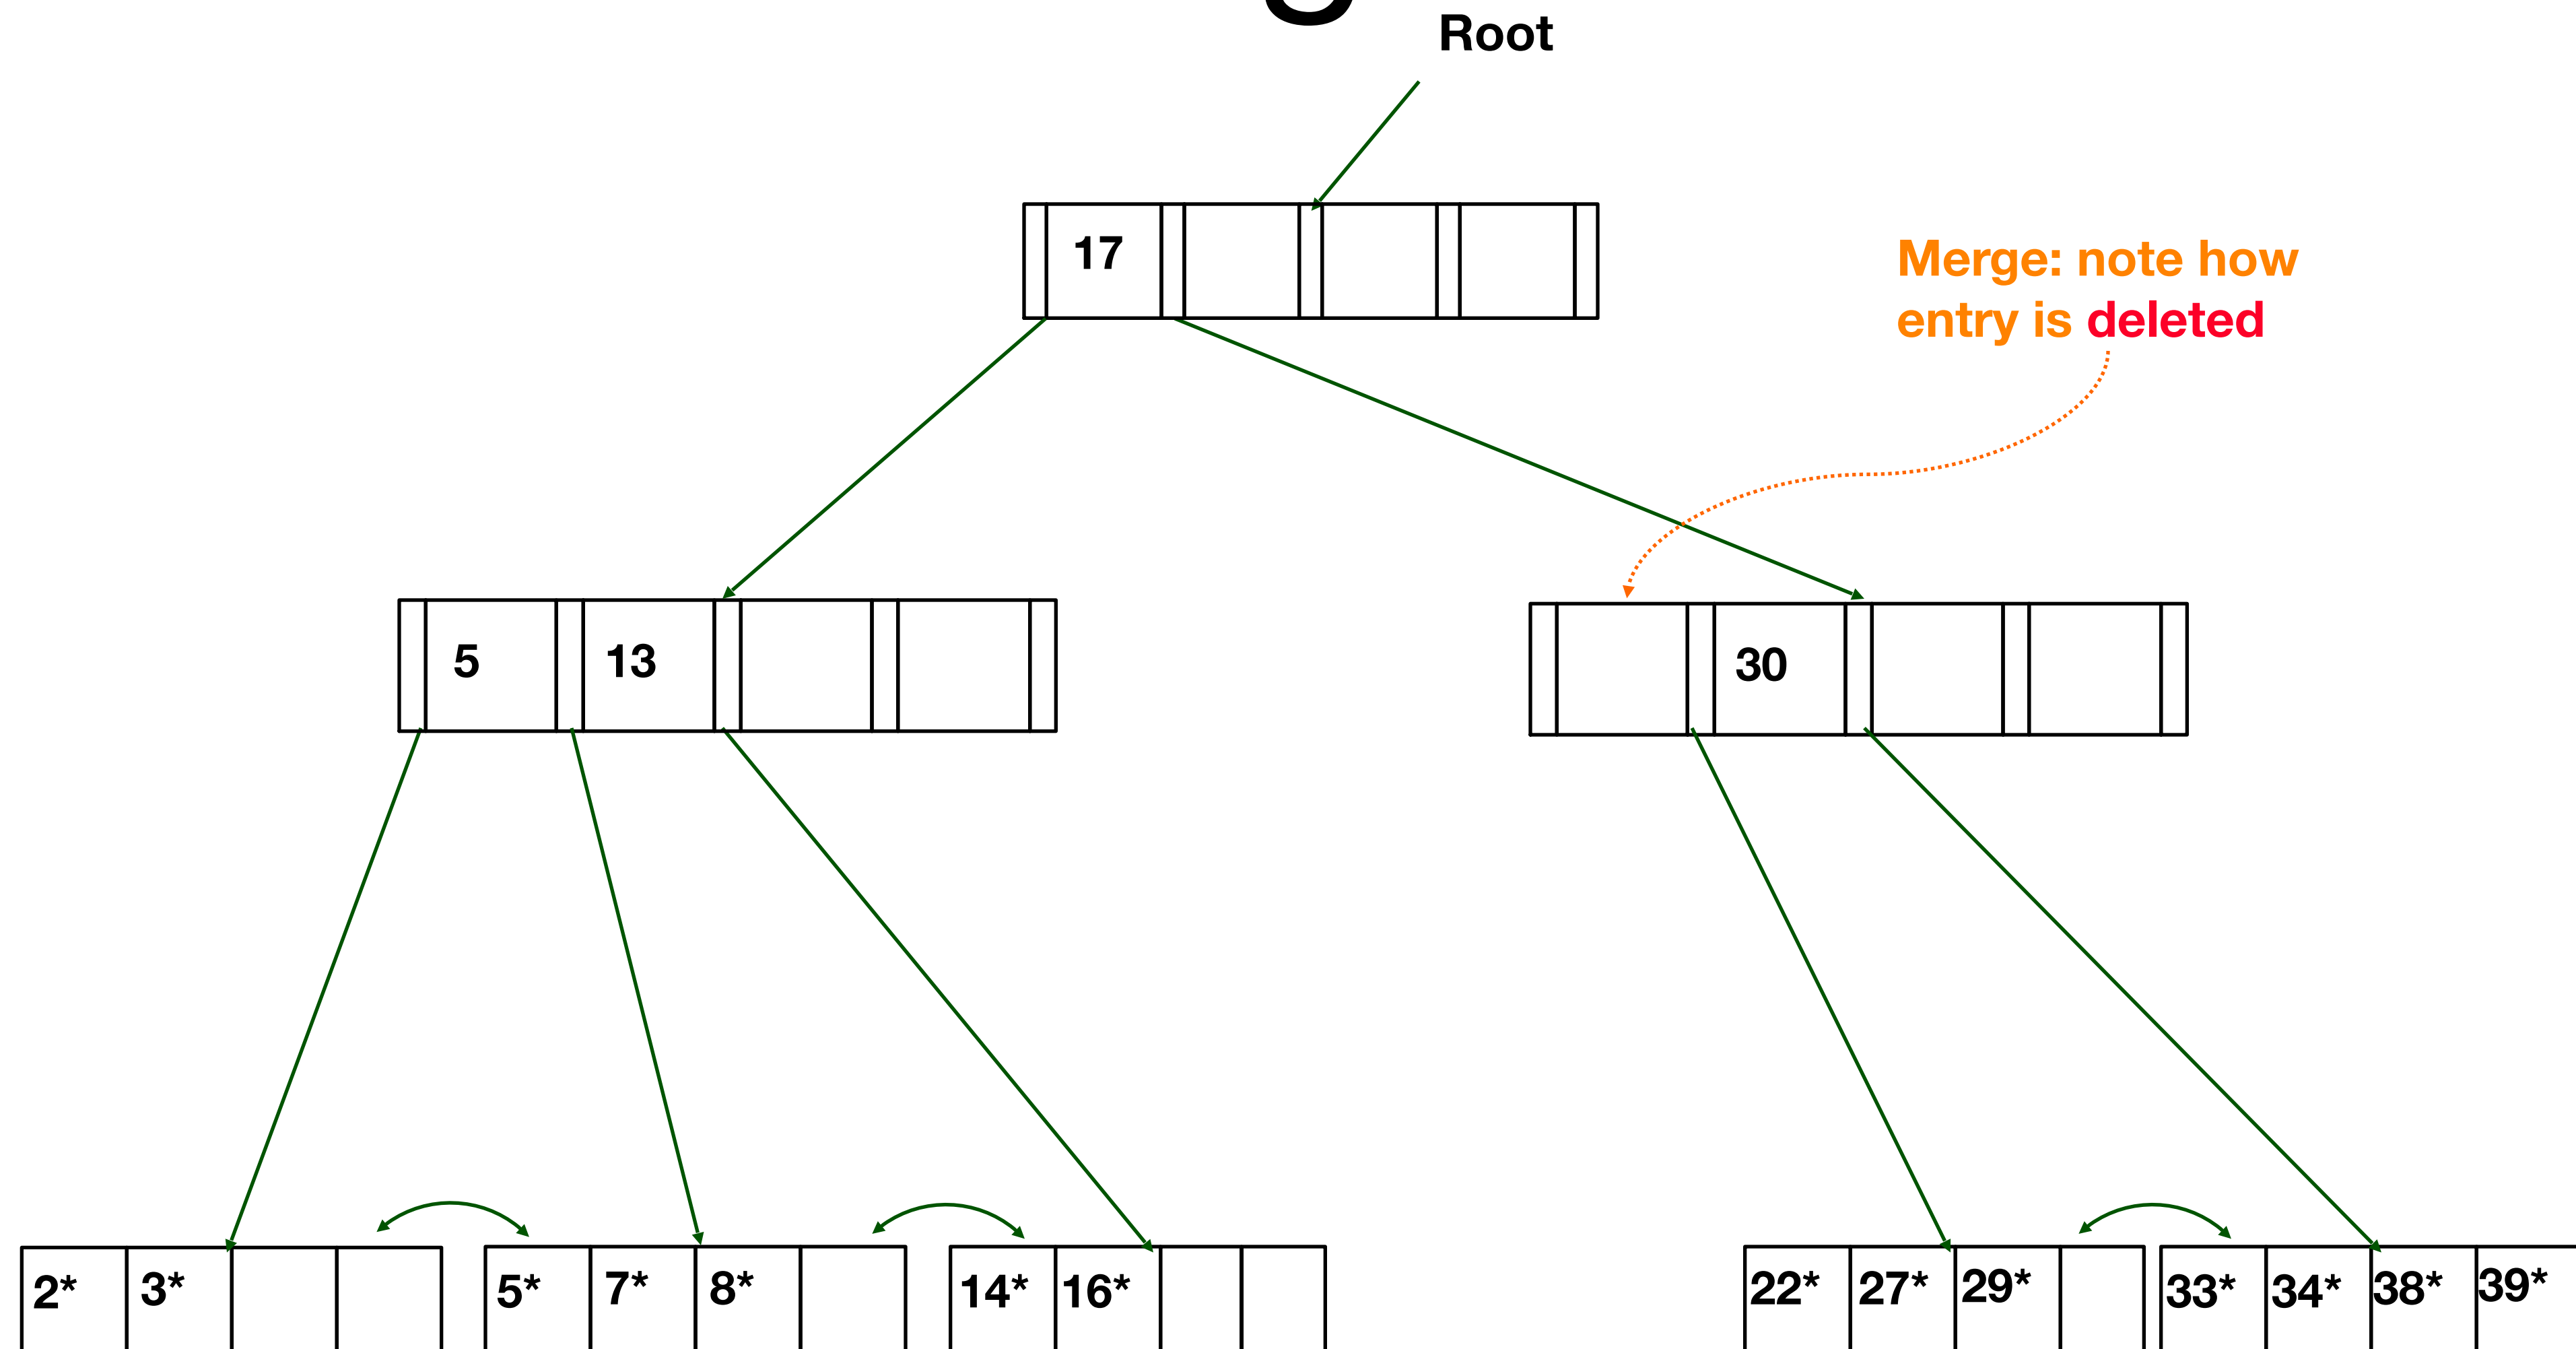


Merging/Deleting Underfull Leaves ...

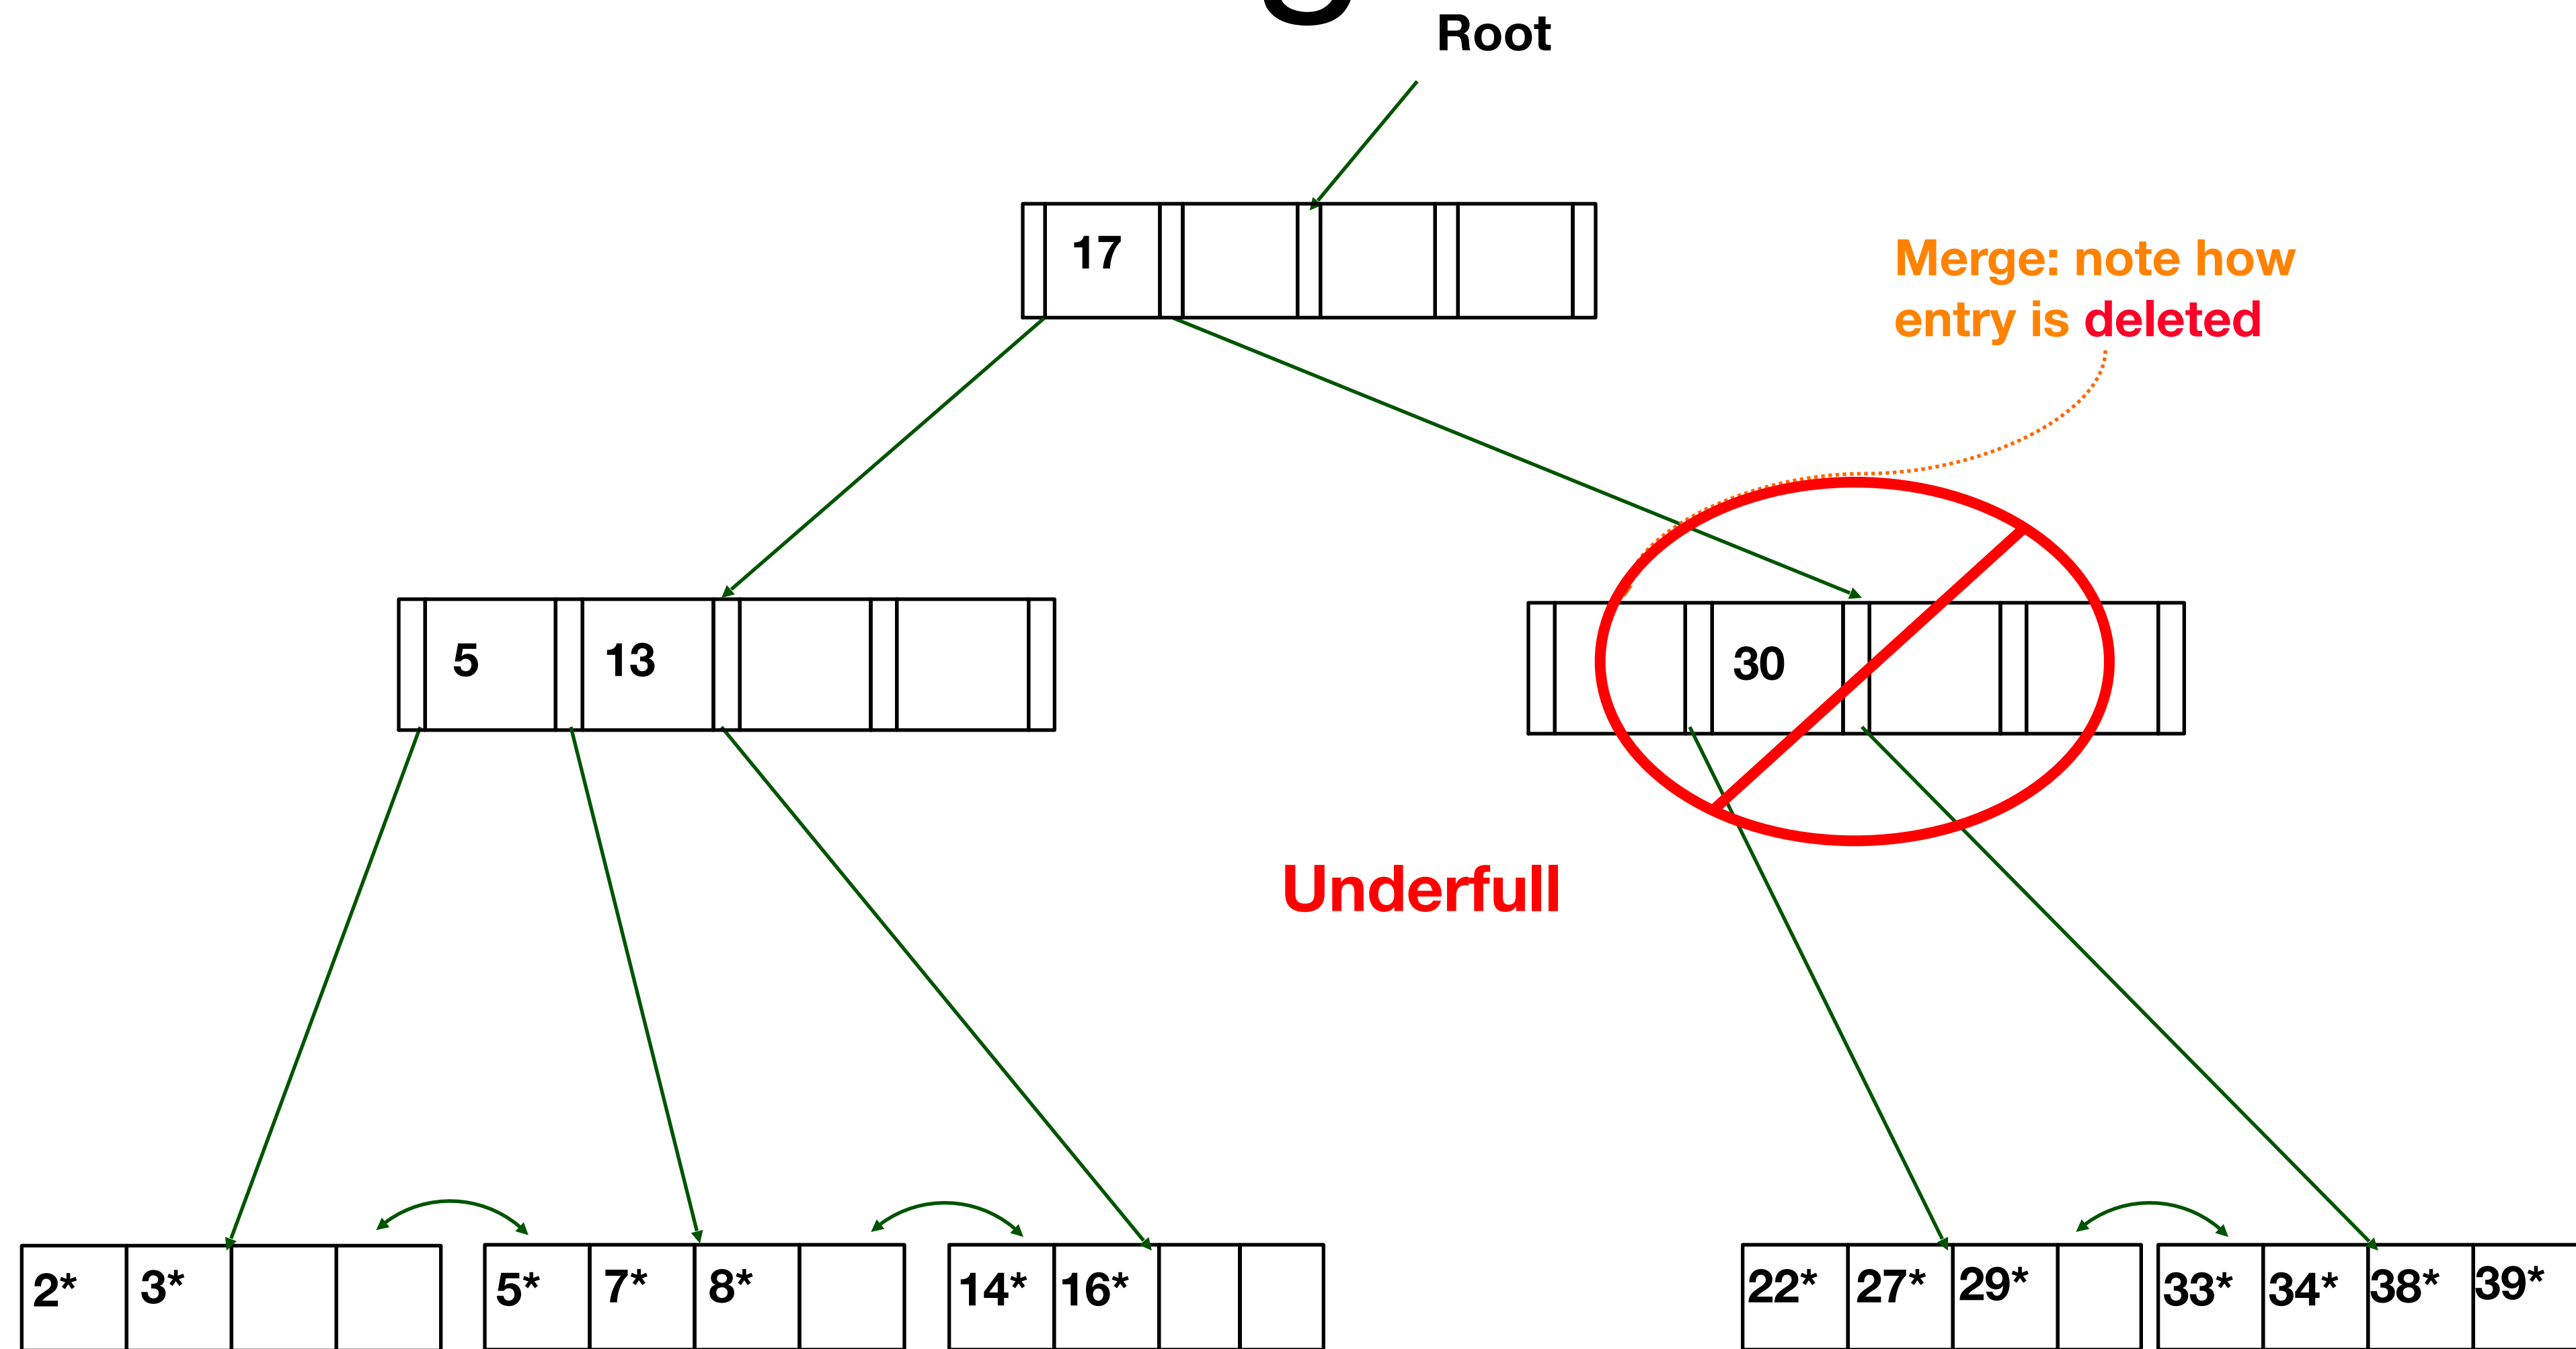


- If the node fits in its sibling, delete node ...
- This **deletes** a key from parent
- Parent may become underfull, so repeat


Deleting 24* ...



Deleting 24* ...

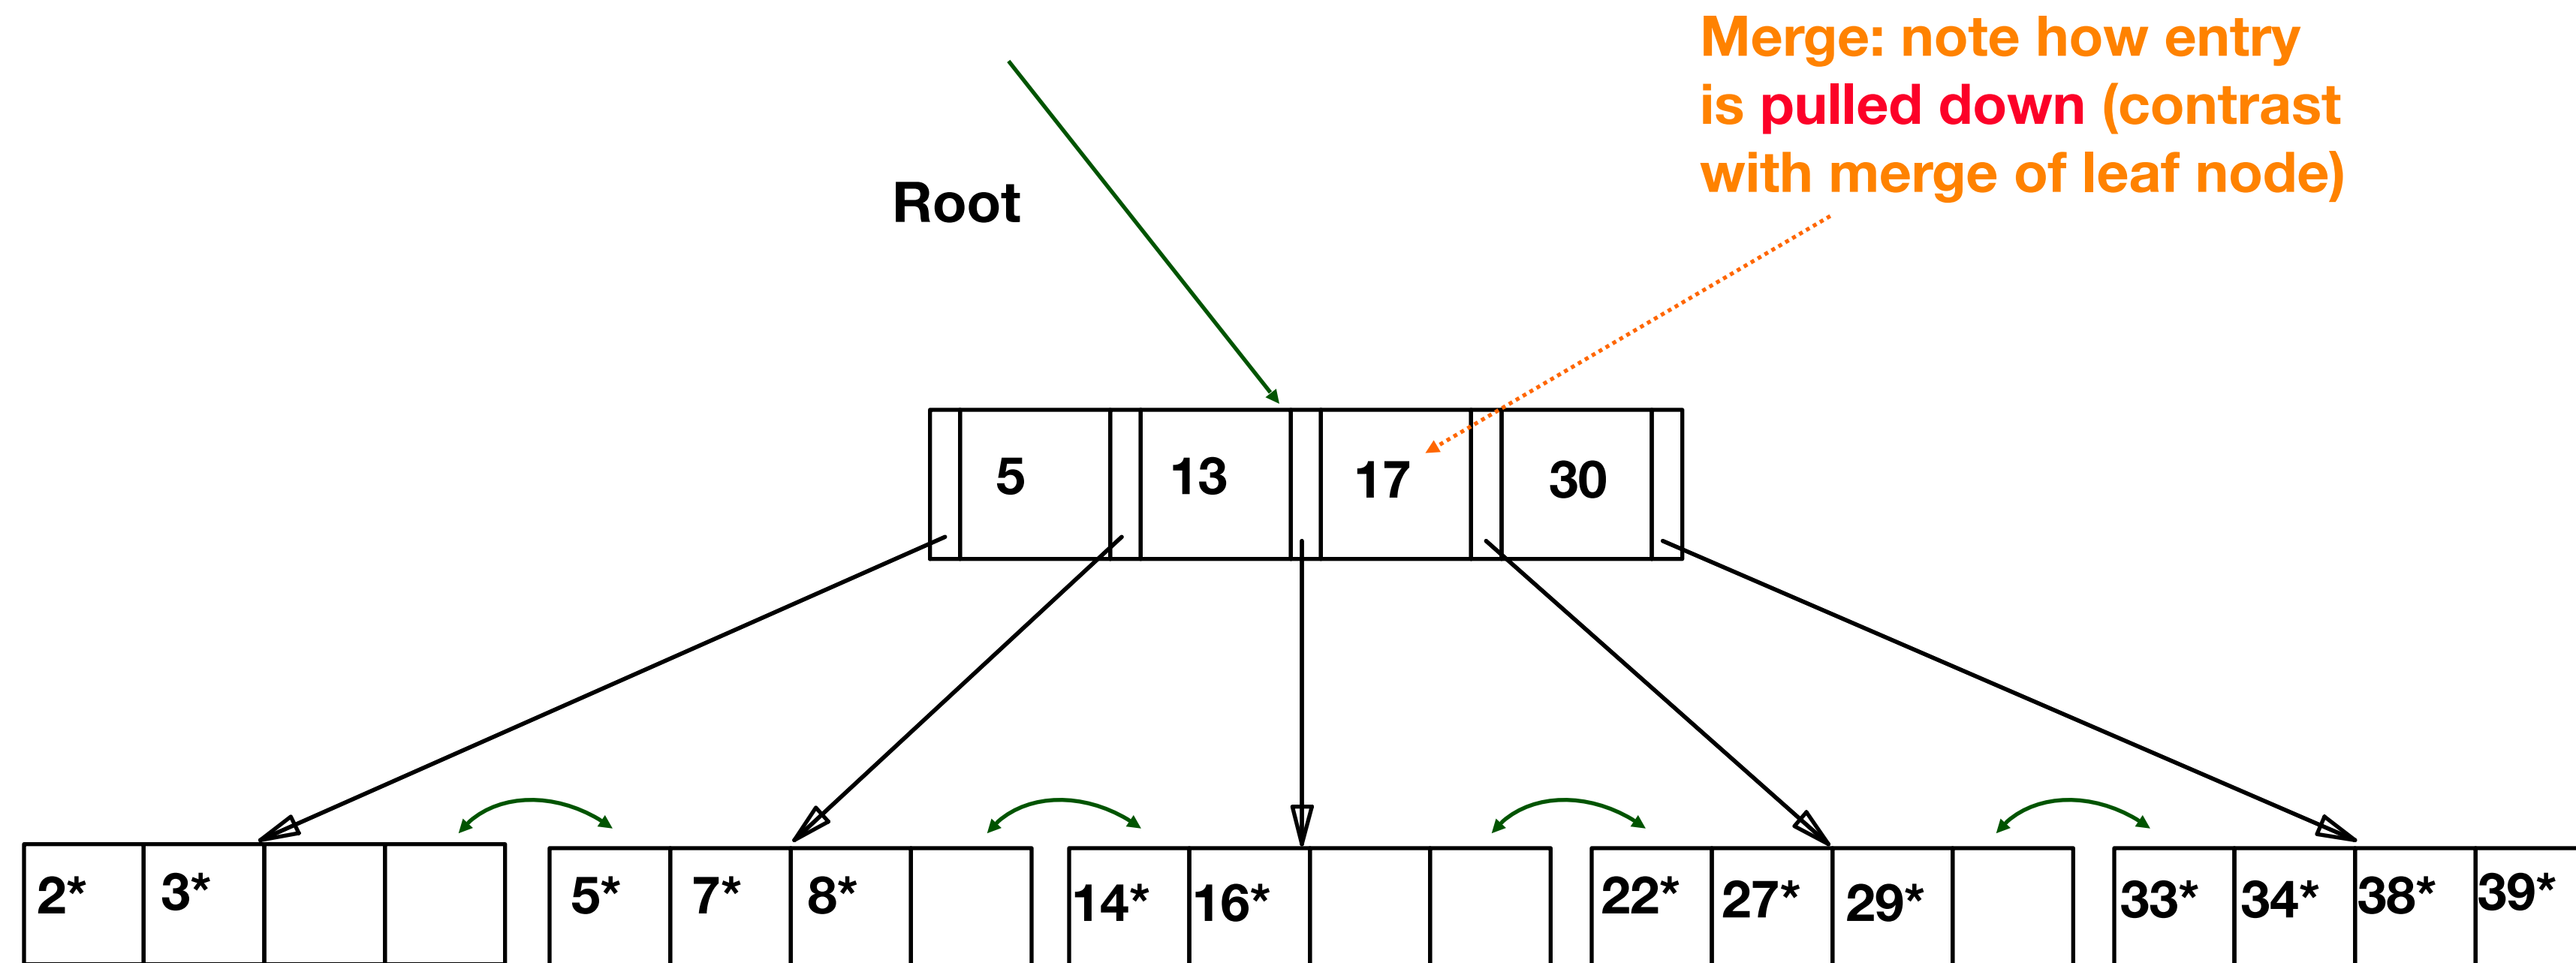


Merging Underfull Nonleaves ...




- If the node fits in its sibling, delete ...
- This **moves** a key from parent to merged child
- Parent may become underfull, so repeat

Deleting 24* ...




Redistributing Underfull Nonleaves ...



- Must redistribute if sibling node is full
- Keys are moved rather than copied
- Degree of parent unchanged

Composite Search Keys



What should we learn today?

- Explain the concept of **(materialized) views** and create them in SQL
- Explain the concept of **triggers** and create them in SQL
- Argue for when to use **(materialized) views** vs. **triggers**
- Explain the **ACID** properties of transactions
- Formulate **transaction programs** in SQL
- Identify the main types of representations and access methods for relations in DBMS, namely **heap files** and **indexes**
- Explain the core algorithms for insertion, deletion, and search in a **B⁺-tree**

